



# Guiding Deep Probabilistic Models

Timur Garipov

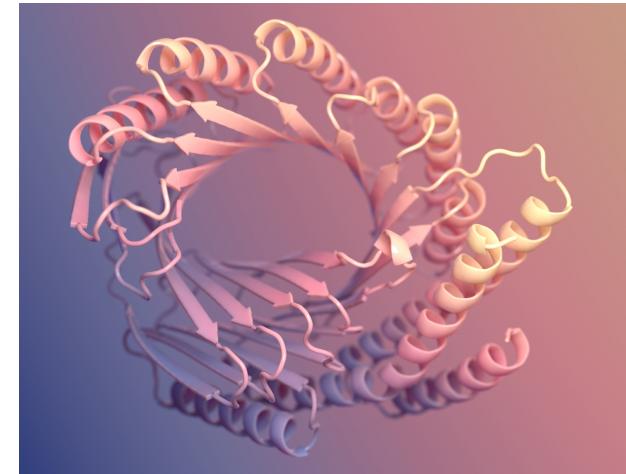
06/13/2024



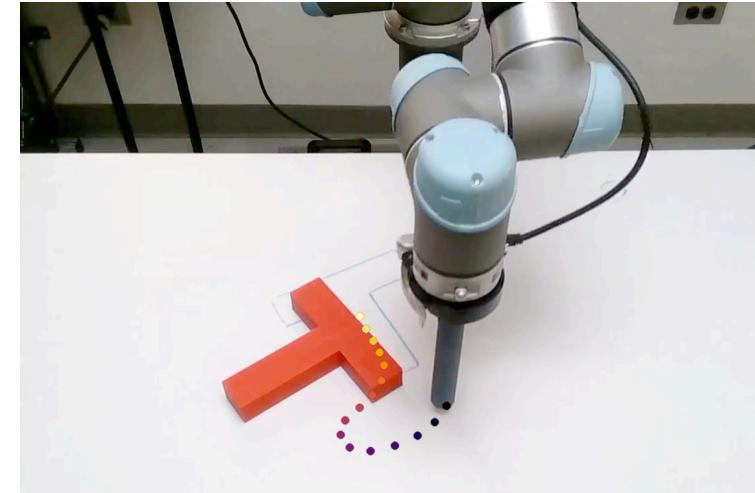
**GPT-4**  
[OpenAI, 2023]



**DALL-E 3**  
[OpenAI, 2023]



**RFdiffusion**  
[Watson et al., 2023]



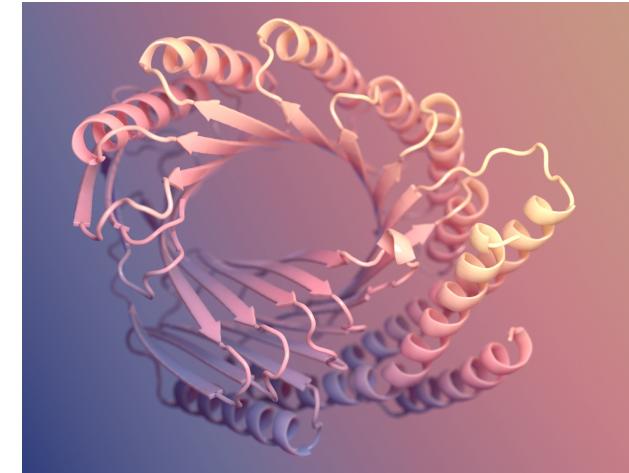
**Diffusion Policy**  
[Chi et al., 2023]



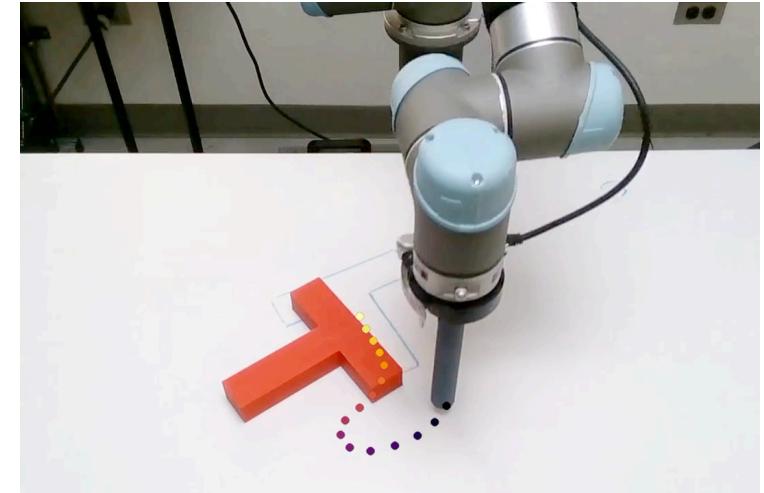
**GPT-4**  
**[OpenAI, 2023]**



**DALL-E 3**  
**[OpenAI, 2023]**



**RFdiffusion**  
**[Watson et al., 2023]**



**Diffusion Policy**  
**[Chi et al., 2023]**

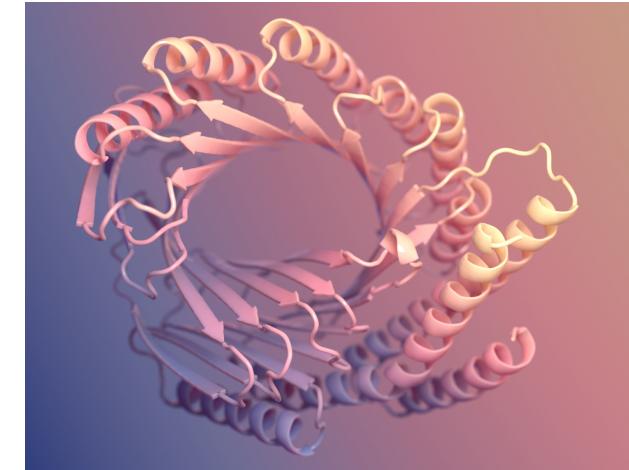
Ground-breaking impact in language modeling, image generation, sciences, robotics



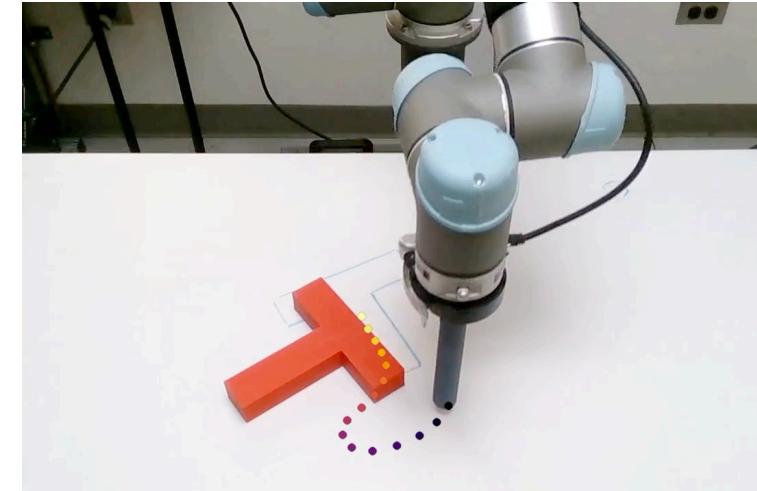
**GPT-4**  
**[OpenAI, 2023]**



**DALL-E 3**  
**[OpenAI, 2023]**



**RFdiffusion**  
**[Watson et al., 2023]**



**Diffusion Policy**  
**[Chi et al., 2023]**

Ground-breaking impact in language modeling, image generation, sciences, robotics

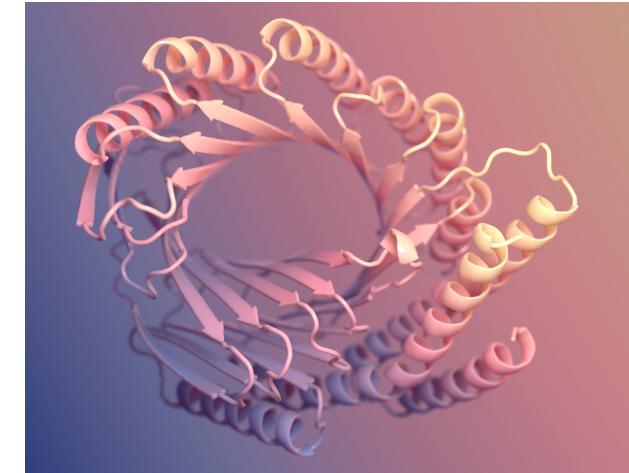
How to make progress in areas where direct supervision signals are limited?



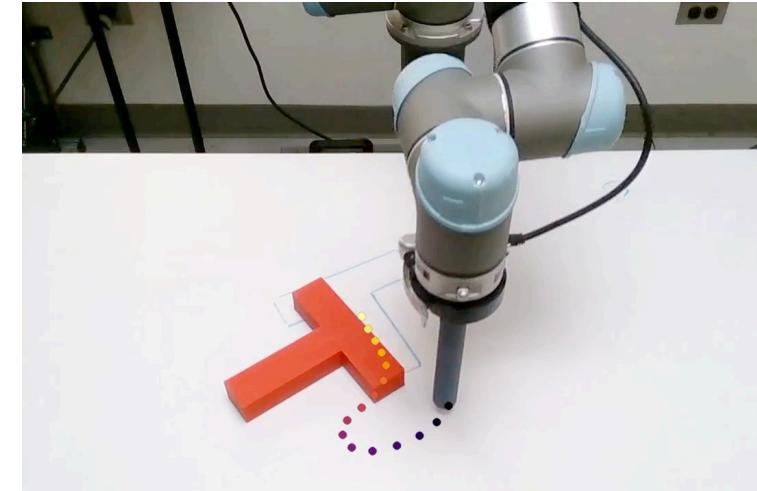
**GPT-4**  
**[OpenAI, 2023]**



**DALL-E 3**  
**[OpenAI, 2023]**



**RFdiffusion**  
**[Watson et al., 2023]**



**Diffusion Policy**  
**[Chi et al., 2023]**

Ground-breaking impact in language modeling, image generation, sciences, robotics

How to make progress in areas where direct supervision signals are limited?

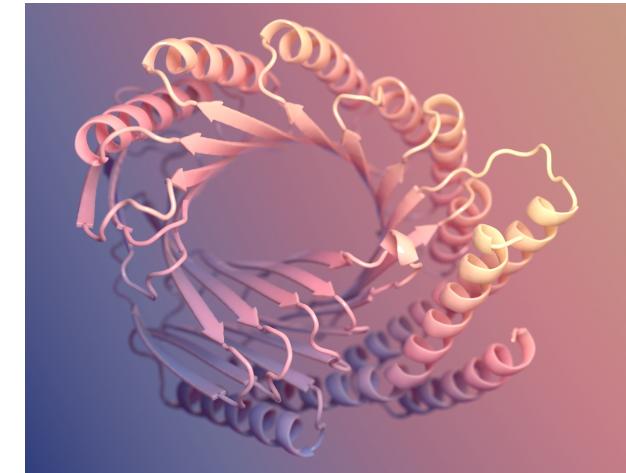
**Reasoning and planning**



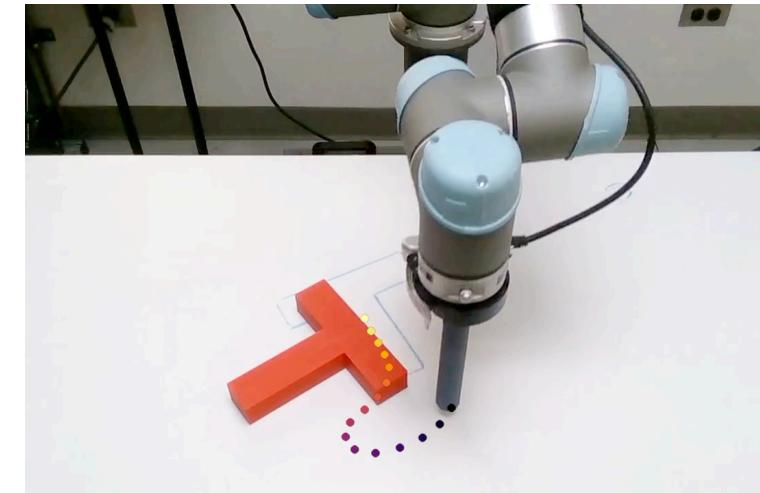
**GPT-4**  
**[OpenAI, 2023]**



**DALL-E 3**  
**[OpenAI, 2023]**



**RFdiffusion**  
**[Watson et al., 2023]**



**Diffusion Policy**  
**[Chi et al., 2023]**

Ground-breaking impact in language modeling, image generation, sciences, robotics

How to make progress in areas where direct supervision signals are limited?

**Reasoning and planning**

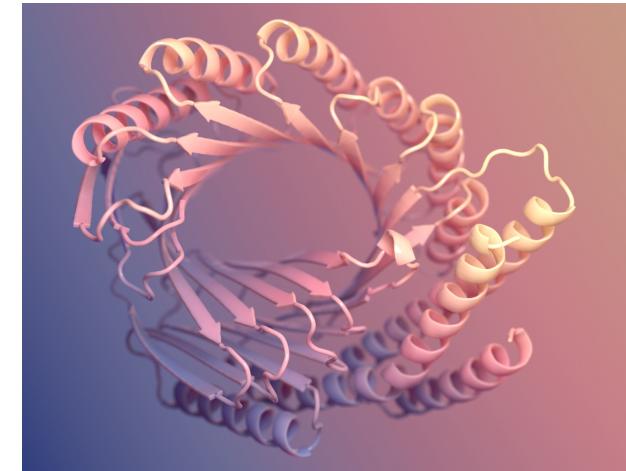
**Designing complex experiments and  
generating hypotheses in science**



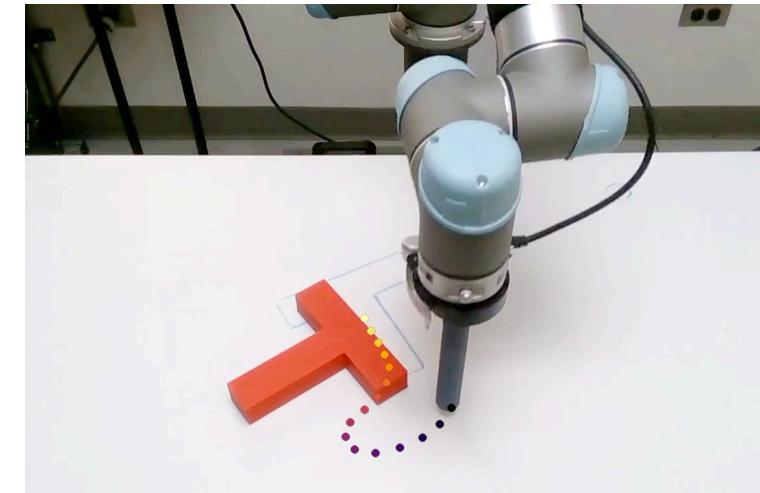
**GPT-4**  
**[OpenAI, 2023]**



**DALL-E 3**  
**[OpenAI, 2023]**



**RFdiffusion**  
**[Watson et al., 2023]**



**Diffusion Policy**  
**[Chi et al., 2023]**

Ground-breaking impact in language modeling, image generation, sciences, robotics

How to make progress in areas where direct supervision signals are limited?

**Reasoning and planning**

**Designing complex experiments and  
generating hypotheses in science**

**Solving data-scarce and  
intricately-structured problems**

# Guidance and Modular Design

**Idea:** train ML models to provide signal for training / inference in other (larger) models

- ▶ Flexible supervision signals
- ▶ Controllable inference (generation under multiple constraints)

# Guidance and Modular Design

**Idea:** train ML models to provide signal for training / inference in other (larger) models

- ▶ Flexible supervision signals
- ▶ Controllable inference (generation under multiple constraints)

**Why?**

- ▶ Need to train and evaluate models without direct supervision signals
- ▶ Need to steer multi-step inference processes
- ▶ Need to re-use and adapt large pre-trained models
- ▶ Need to combine and coordinate multiple models with different areas of expertise

# Guidance and Modular Design

**Idea:** train ML models to provide signal for training / inference in other (larger) models

- ▶ Flexible supervision signals
- ▶ Controllable inference (generation under multiple constraints)

**Why?**

- ▶ Need to train and evaluate models without direct supervision signals
- ▶ Need to steer multi-step inference processes
- ▶ Need to re-use and adapt large pre-trained models
- ▶ Need to combine and coordinate multiple models with different areas of expertise

**Challenges in Guiding Deep Probabilistic Models**

- ▶ Need to manipulate complex probability distributions in high-dimensional spaces
- ▶ Sophisticated and often brittle models, complicated optimization landscapes
- ▶ Require computational efficient training and inference algorithms

# Research Focus & Goals

Thesis: "Guiding Deep Probabilistic Models"

## Research focus areas

- ▶ Addressing complex training dynamics between models trained with different objectives
- ▶ Design of novel training criteria, addressing shortcomings of existing objectives
- ▶ Representing complex probability distributions through generative model combination

## Goals

- ▶ Novel principled algorithms for training and inference in deep probabilistic models
- ▶ Guarantees
  - ▶ Training: optimality of the desired target configurations
  - ▶ Inference: sampling from target distributions

# Chapter II

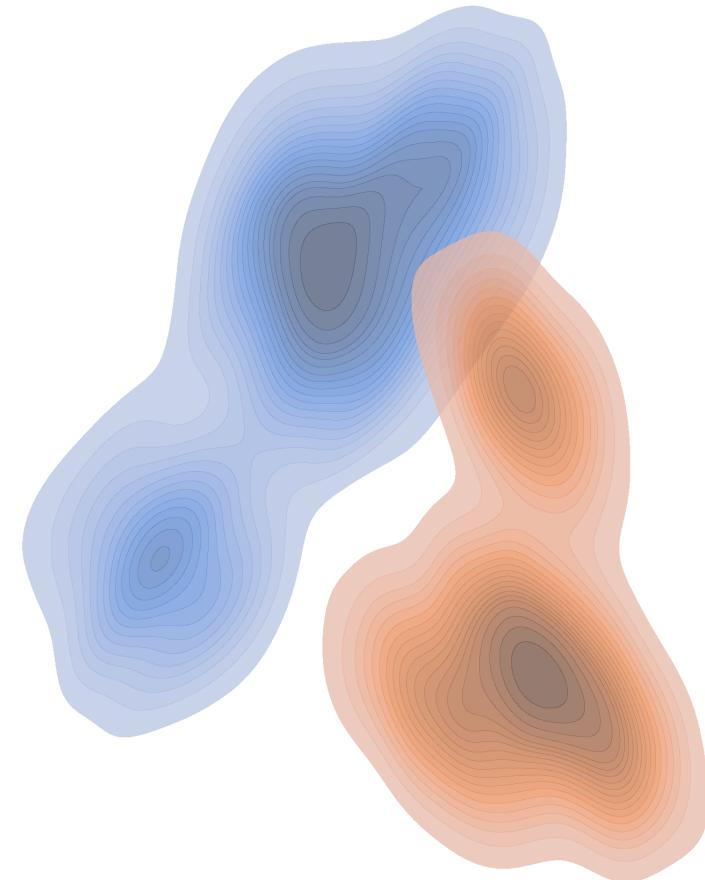
## Pairwise-Discriminator Objectives for Generative Adversarial Networks

---

The Benefits of Pairwise Discriminators for Adversarial Training  
S. Tong\*, T. Garipov\*, T. Jaakkola (arXiv Pre-print, 2020)

# Guidance for Generative Model Training

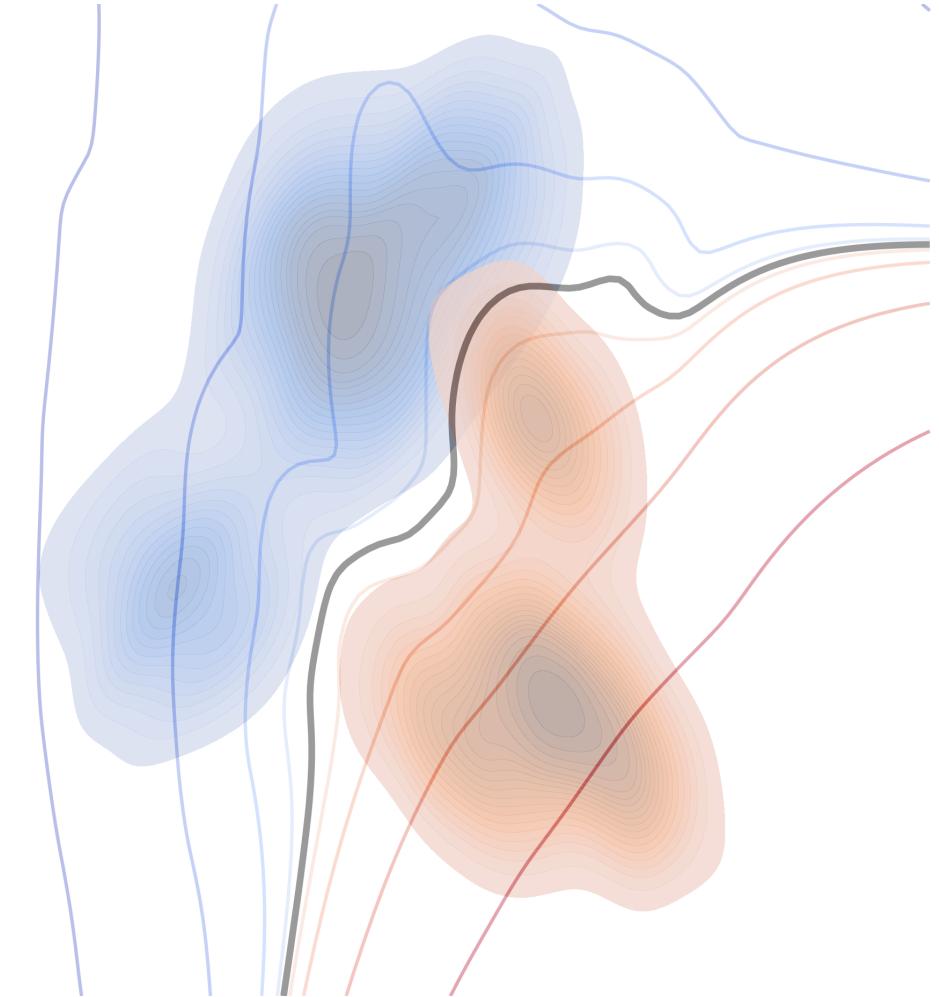
$$\text{JSD}(p \parallel q) = \frac{1}{2} \text{KL} \left( p \parallel \frac{p+q}{2} \right) + \frac{1}{2} \text{KL} \left( q \parallel \frac{p+q}{2} \right)$$



# Guidance for Generative Model Training

$$\text{JSD}(p \parallel q) = \frac{1}{2} \text{KL}\left(p \parallel \frac{p+q}{2}\right) + \frac{1}{2} \text{KL}\left(q \parallel \frac{p+q}{2}\right)$$

**Idea:** train a probabilistic classifier to estimate divergence between distributions



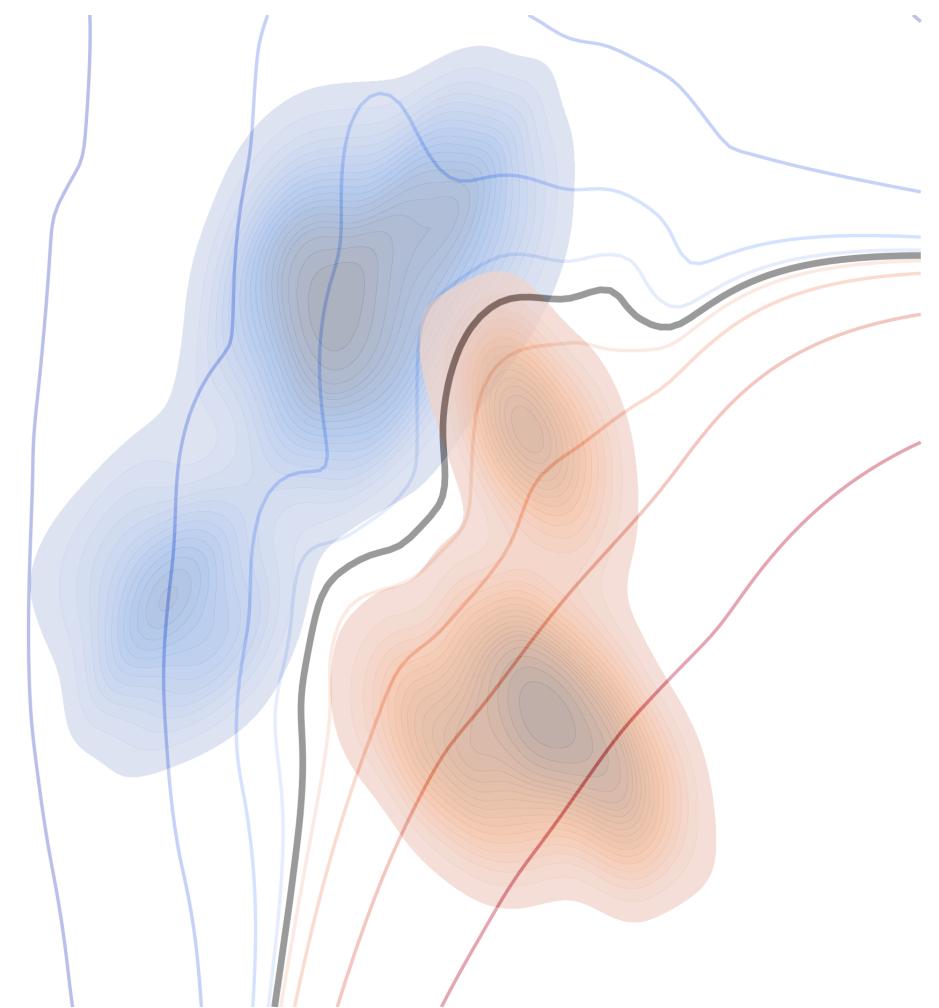
# Guidance for Generative Model Training

$$\text{JSD}(p \parallel q) = \frac{1}{2} \text{KL}\left(p \parallel \frac{p+q}{2}\right) + \frac{1}{2} \text{KL}\left(q \parallel \frac{p+q}{2}\right)$$

$$\text{JSD}(p \parallel q) = \log(2) + \frac{1}{2} \min_{u: \mathcal{X} \rightarrow \mathbb{R}} \mathcal{L}(p, q, u)$$

$$\mathcal{L}(p, q, u) = \mathbb{E}_{p(x)} [\log(1 + \exp(-u(x)))] + \mathbb{E}_{q(x)} [\log(1 + \exp(u(x)))]$$

**Idea:** train a probabilistic classifier to estimate divergence between distributions



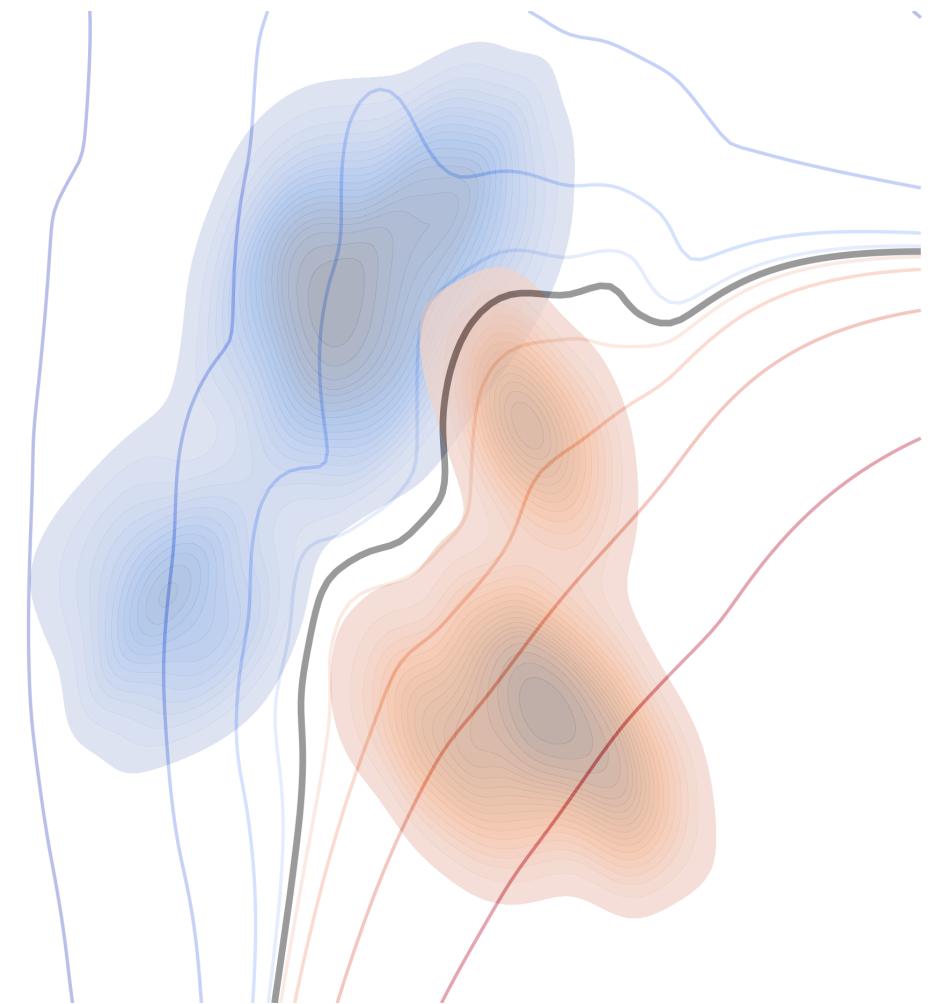
# Guidance for Generative Model Training

$$\text{JSD}(p \| q) = \frac{1}{2} \text{KL} \left( p \left\| \frac{p+q}{2} \right. \right) + \frac{1}{2} \text{KL} \left( q \left\| \frac{p+q}{2} \right. \right)$$

$$\text{JSD}(p \| q) = \log(2) + \frac{1}{2} \min_{u: \mathcal{X} \rightarrow \mathbb{R}} \mathcal{L}(p, q, u)$$

$$\mathcal{L}(p, q, u) = \mathbb{E}_{p(x)} [\log(1 + \exp(-u(x)))] + \mathbb{E}_{q(x)} [\log(1 + \exp(u(x)))]$$

**Idea:** train a probabilistic classifier to estimate divergence between distributions



## Generative Adversarial Networks (GANs)

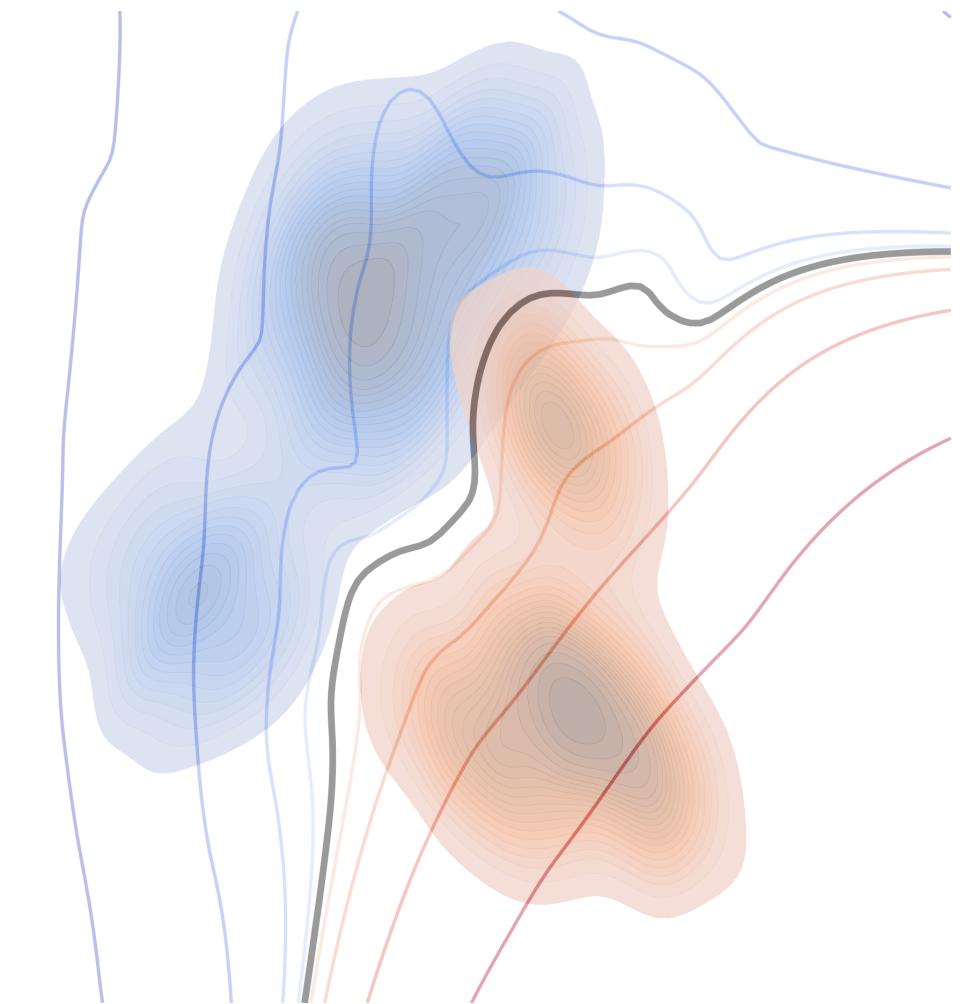
# Guidance for Generative Model Training

$$\text{JSD}(p \| q) = \frac{1}{2} \text{KL} \left( p \left\| \frac{p+q}{2} \right. \right) + \frac{1}{2} \text{KL} \left( q \left\| \frac{p+q}{2} \right. \right)$$

$$\text{JSD}(p \| q) = \log(2) + \frac{1}{2} \min_{u: \mathcal{X} \rightarrow \mathbb{R}} \mathcal{L}(p, q, u)$$

$$\mathcal{L}(p, q, u) = \mathbb{E}_{p(x)} [\log(1 + \exp(-u(x)))] + \mathbb{E}_{q(x)} [\log(1 + \exp(u(x)))]$$

**Idea:** train a probabilistic classifier to estimate divergence between distributions



## Generative Adversarial Networks (GANs)

**Data distribution:**  $p(x)$

**Generator:**  $F_\theta : \mathcal{Z} \rightarrow \mathcal{X}$ ,  $q_\theta(x) : x = F_\theta(z), z \sim q(z)$

**Discriminator:**  $u_\phi : \mathcal{X} \rightarrow \mathbb{R}$ ,  $\hat{P}_\phi(y = \text{data}|x) = \frac{1}{1 + \exp(-u_\phi(x))}$

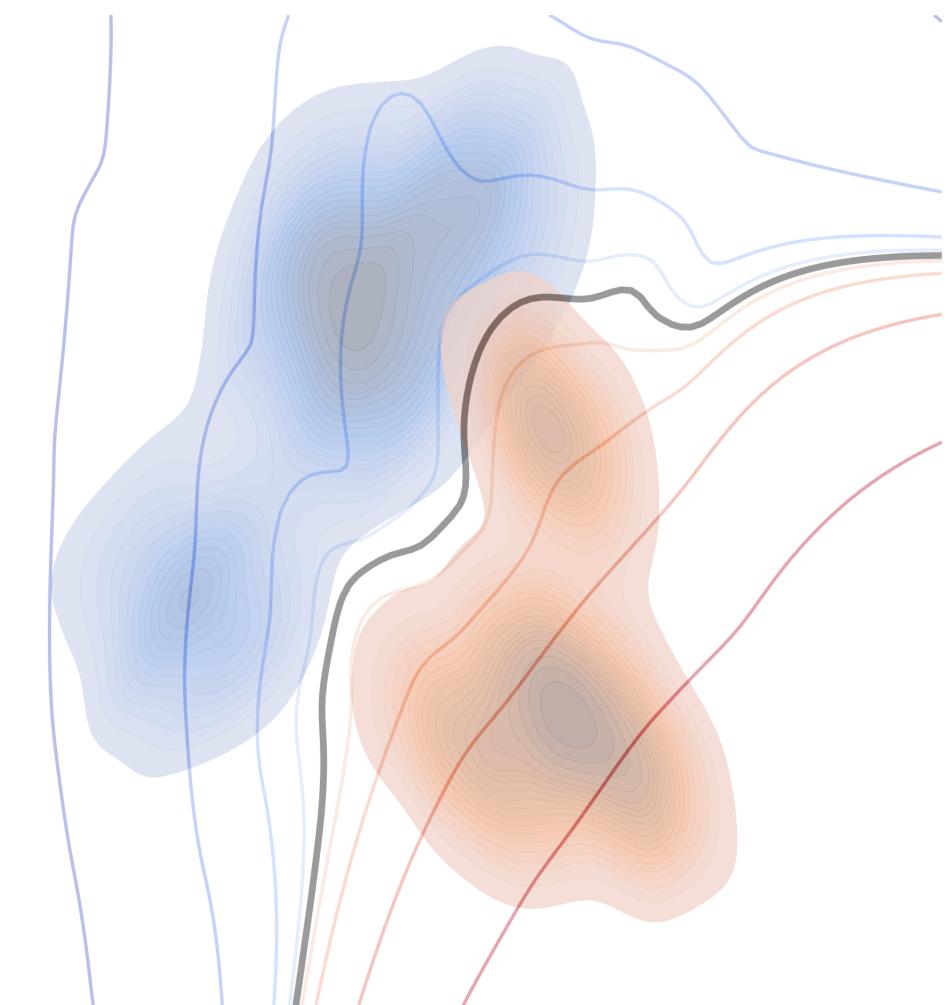
# Guidance for Generative Model Training

$$\text{JSD}(p \| q) = \frac{1}{2} \text{KL} \left( p \left\| \frac{p+q}{2} \right. \right) + \frac{1}{2} \text{KL} \left( q \left\| \frac{p+q}{2} \right. \right)$$

$$\text{JSD}(p \| q) = \log(2) + \frac{1}{2} \min_{u: \mathcal{X} \rightarrow \mathbb{R}} \mathcal{L}(p, q, u)$$

$$\mathcal{L}(p, q, u) = \mathbb{E}_{p(x)} [\log(1 + \exp(-u(x)))] + \mathbb{E}_{q(x)} [\log(1 + \exp(u(x)))]$$

**Idea:** train a probabilistic classifier to estimate divergence between distributions



## Generative Adversarial Networks (GANs)

**Data distribution:**  $p(x)$

**Generator:**  $F_\theta : \mathcal{Z} \rightarrow \mathcal{X}, \quad q_\theta(x) : x = F_\theta(z), \quad z \sim q(z)$

**Discriminator:**  $u_\phi : \mathcal{X} \rightarrow \mathbb{R}, \quad \hat{P}_\phi(y = \text{data}|x) = \frac{1}{1 + \exp(-u_\phi(x))}$

$$u^* = \underset{u: \mathcal{X} \rightarrow \mathbb{R}}{\operatorname{argmin}} \mathcal{L}(p, q, u), \quad u^*(x) = \log \frac{p(x)}{q(x)}, \quad P^*(y = \text{data}|x) = \frac{p(x)}{p(x) + q(x)}$$

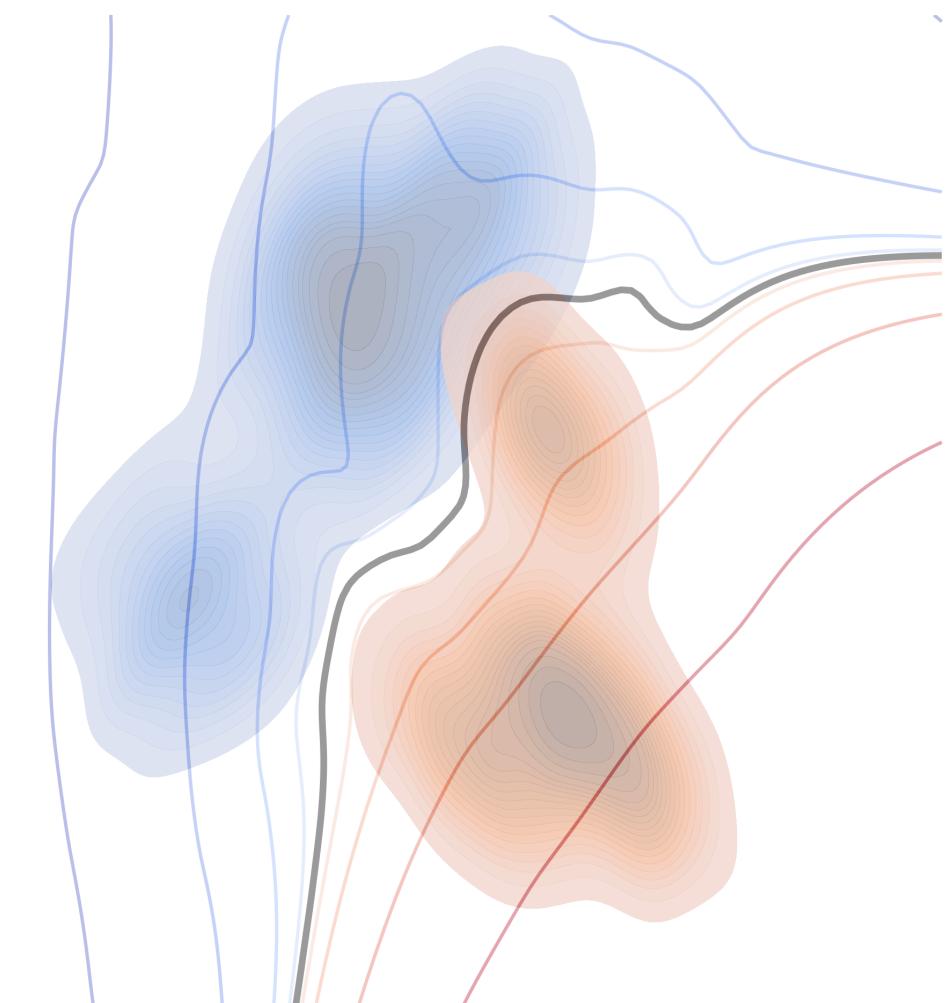
# Guidance for Generative Model Training

$$\text{JSD}(p \| q) = \frac{1}{2} \text{KL}\left(p \middle\| \frac{p+q}{2}\right) + \frac{1}{2} \text{KL}\left(q \middle\| \frac{p+q}{2}\right)$$

$$\text{JSD}(p \| q) = \log(2) + \frac{1}{2} \min_{u: \mathcal{X} \rightarrow \mathbb{R}} \mathcal{L}(p, q, u)$$

$$\mathcal{L}(p, q, u) = \mathbb{E}_{p(x)} [\log(1 + \exp(-u(x)))] + \mathbb{E}_{q(x)} [\log(1 + \exp(u(x)))]$$

**Idea:** train a probabilistic classifier to estimate divergence between distributions



## Generative Adversarial Networks (GANs)

**Data distribution:**  $p(x)$

**Generator:**  $F_\theta : \mathcal{Z} \rightarrow \mathcal{X}, \quad q_\theta(x) : x = F_\theta(z), \quad z \sim q(z)$

**Discriminator:**  $u_\phi : \mathcal{X} \rightarrow \mathbb{R}, \quad \hat{P}_\phi(y = \text{data}|x) = \frac{1}{1 + \exp(-u_\phi(x))}$

$$u^* = \underset{u: \mathcal{X} \rightarrow \mathbb{R}}{\operatorname{argmin}} \mathcal{L}(p, q, u), \quad u^*(x) = \log \frac{p(x)}{q(x)}, \quad P^*(y = \text{data}|x) = \frac{p(x)}{p(x) + q(x)}$$

$$\mathcal{L}(p, q, u^*(p, q)) = 2 \text{JSD}(p \| q) - \log(4)$$

**Max-Min game:**  $\max_{q_\theta} \min_{u_\phi} \mathcal{L}(p, q_\theta, u_\phi)$

**Equilibrium:**  $q^* = p^*, u^* = 0$

# Training Dynamics

**Max-Min game:**  $\max_{q_\theta} \min_{u_\phi} \mathcal{L}(p, q_\theta, u_\phi)$

**Equilibrium:**  $q^* = p^*, u^* = 0$

# Training Dynamics

**Max-Min game:**  $\max_{q_\theta} \min_{u_\phi} \mathcal{L}(p, q_\theta, u_\phi)$

**Equilibrium:**  $q^* = p^*, u^* = 0$

## In Theory:

Train **optimal discriminator**  $u^*$  by minimizing  $\mathcal{L}(p, q_\theta, u)$

Update  $F_\theta$  with gradient of  $\mathcal{L}(p, q_\theta, u^*)$  given **optimal discriminator**  $u^*$

# Training Dynamics

**Max-Min game:**  $\max_{q_\theta} \min_{u_\phi} \mathcal{L}(p, q_\theta, u_\phi)$

**Equilibrium:**  $q^* = p^*, u^* = 0$

## In Theory:

Train **optimal discriminator**  $u^*$  by minimizing  $\mathcal{L}(p, q_\theta, u)$

Update  $F_\theta$  with gradient of  $\mathcal{L}(p, q_\theta, u^*)$  given **optimal discriminator**  $u^*$

## In Practice:

Update  $u_\phi$  with **a few optimization steps** on  $\mathcal{L}(p, q_\theta, u)$

Update  $F_\theta$  with gradient of  $\mathcal{L}(p, q_\theta, u_\phi)$  given **current discriminator**  $u_\phi$

# Training Dynamics

**Max-Min game:**  $\max_{q_\theta} \min_{u_\phi} \mathcal{L}(p, q_\theta, u_\phi)$

**Equilibrium:**  $q^* = p^*, u^* = 0$

## In Theory:

Train **optimal discriminator**  $u^*$  by minimizing  $\mathcal{L}(p, q_\theta, u)$

Update  $F_\theta$  with gradient of  $\mathcal{L}(p, q_\theta, u^*)$  given **optimal discriminator**  $u^*$

- ▶ The generator receives training signal through the discriminator
- ▶ Even if the generator is perfectly aligned with the target, a suboptimal discriminator can destroy the alignment
- ▶ The instability of alignment complicates training dynamics
- ▶ Training can get into a limit cycle and never reach the equilibrium

## In Practice:

Update  $u_\phi$  with **a few optimization steps** on  $\mathcal{L}(p, q_\theta, u)$

Update  $F_\theta$  with gradient of  $\mathcal{L}(p, q_\theta, u_\phi)$  given **current discriminator**  $u_\phi$

# Training Dynamics

**Max-Min game:**  $\max_{q_\theta} \min_{u_\phi} \mathcal{L}(p, q_\theta, u_\phi)$

**Equilibrium:**  $q^* = p^*, u^* = 0$

## In Theory:

Train **optimal discriminator**  $u^*$  by minimizing  $\mathcal{L}(p, q_\theta, u)$

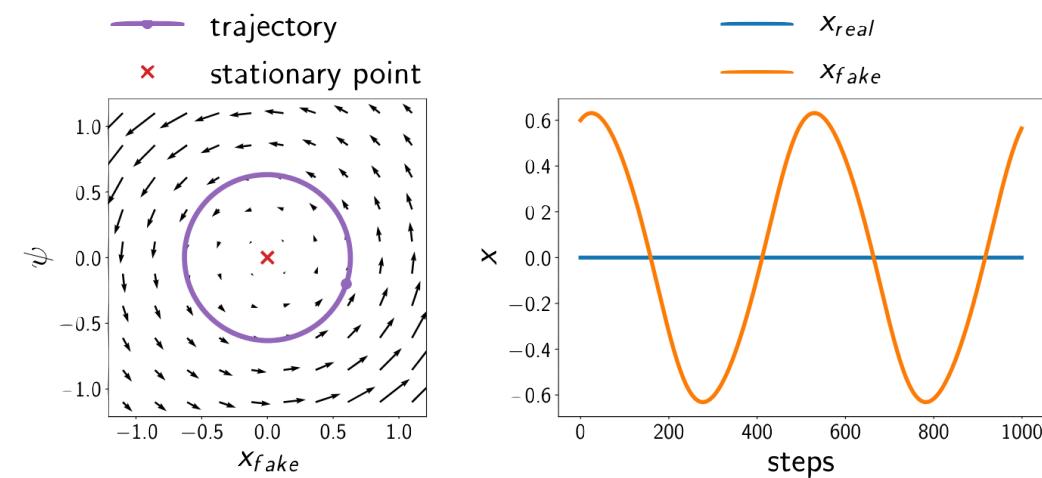
Update  $F_\theta$  with gradient of  $\mathcal{L}(p, q_\theta, u^*)$  given **optimal discriminator**  $u^*$

## In Practice:

Update  $u_\phi$  with **a few optimization steps** on  $\mathcal{L}(p, q_\theta, u)$

Update  $F_\theta$  with gradient of  $\mathcal{L}(p, q_\theta, u_\phi)$  given **current discriminator**  $u_\phi$

- ▶ The generator receives training signal through the discriminator
- ▶ Even if the generator is perfectly aligned with the target, a suboptimal discriminator can destroy the alignment
- ▶ The instability of alignment complicates training dynamics
- ▶ Training can get into a limit cycle and never reach the equilibrium



# Training Dynamics

**Max-Min game:**  $\max_{q_\theta} \min_{u_\phi} \mathcal{L}(p, q_\theta, u_\phi)$

**Equilibrium:**  $q^* = p^*, u^* = 0$

## In Theory:

Train **optimal discriminator**  $u^*$  by minimizing  $\mathcal{L}(p, q_\theta, u)$

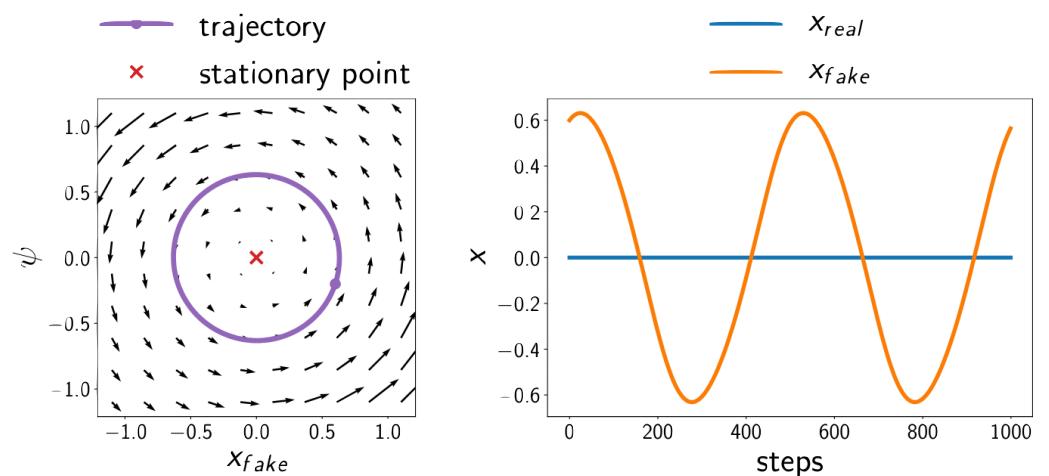
Update  $F_\theta$  with gradient of  $\mathcal{L}(p, q_\theta, u^*)$  given **optimal discriminator**  $u^*$

## In Practice:

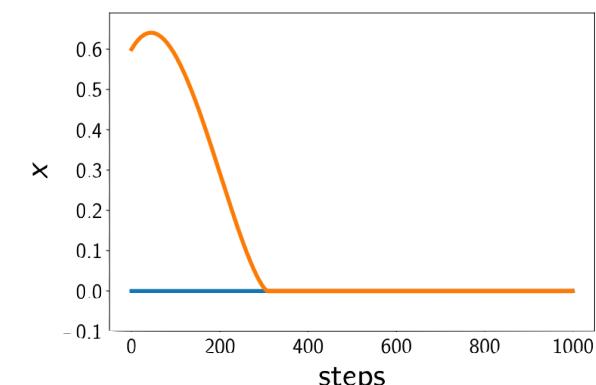
Update  $u_\phi$  with **a few optimization steps** on  $\mathcal{L}(p, q_\theta, u)$

Update  $F_\theta$  with gradient of  $\mathcal{L}(p, q_\theta, u_\phi)$  given **current discriminator**  $u_\phi$

- ▶ The generator receives training signal through the discriminator
- ▶ Even if the generator is perfectly aligned with the target, a suboptimal discriminator can destroy the alignment
- ▶ The instability of alignment complicates training dynamics
- ▶ Training can get into a limit cycle and never reach the equilibrium



**Can we design training objectives  
that preserve distribution alignment?**



# How to Preserve the Alignment?

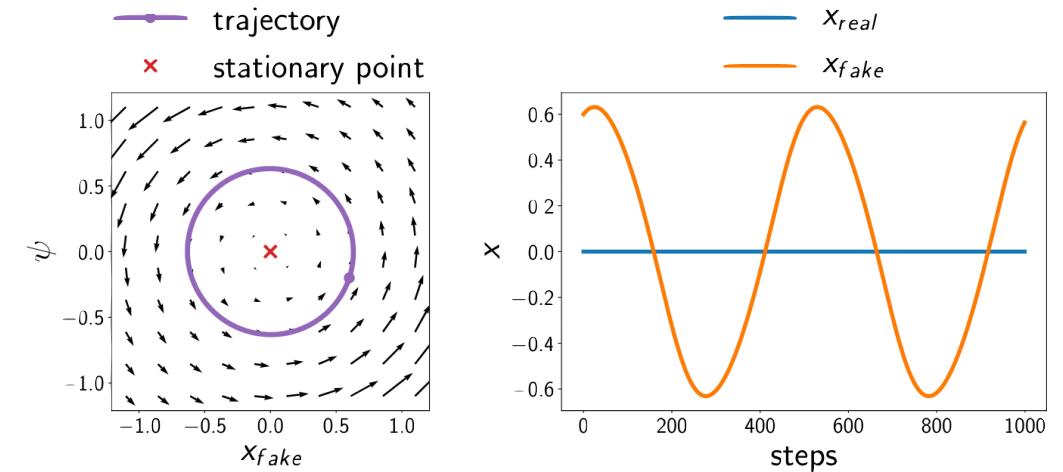
Training with **unary discriminator**  $u(\cdot)$

$u(\cdot)$  distinguishes between **real samples**  $x \sim p(x)$

and **generated samples**  $x \sim q(x)$

$$\mathcal{L}_G(q, u) = \mathbb{E}_{q(x)}[S(u(x))] = \langle a_u^S, q \rangle$$

$$\nabla_q \mathcal{L}_G(q, u) = a_u^S$$



# How to Preserve the Alignment?

Training with **unary discriminator**  $u(\cdot)$

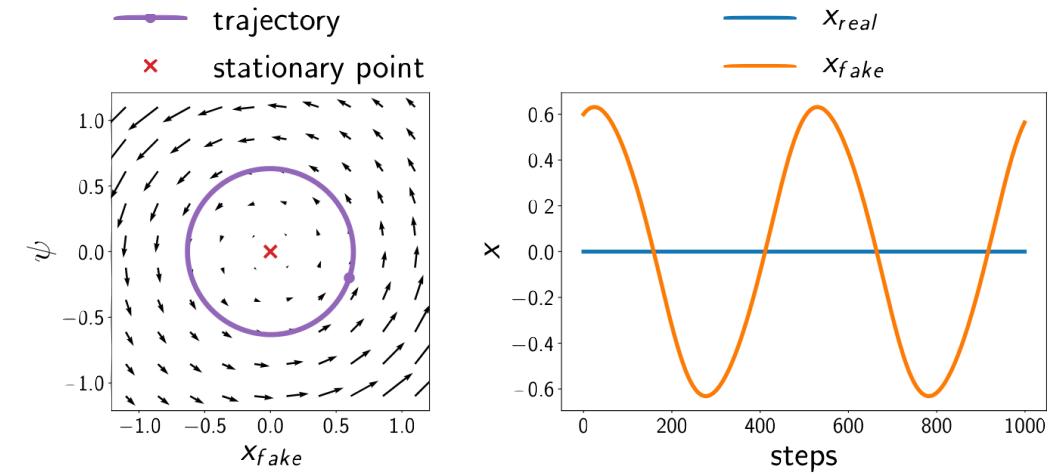
$u(\cdot)$  distinguishes between **real samples**  $x \sim p(x)$

and **generated samples**  $x \sim q(x)$

$$\mathcal{L}_G(q, u) = \mathbb{E}_{q(x)}[S(u(x))] = \langle a_u^S, q \rangle$$

$$\nabla_q \mathcal{L}_G(q, u) = a_u^S$$

- The gradient depends only on the discriminator
- The alignment is preserved only if the discriminator is optimal



# How to Preserve the Alignment?

Training with **unary discriminator**  $u(\cdot)$

$u(\cdot)$  distinguishes between **real samples**  $x \sim p(x)$   
and **generated samples**  $x \sim q(x)$

$$\mathcal{L}_G(q, u) = \mathbb{E}_{q(x)}[S(u(x))] = \langle a_u^S, q \rangle$$

$$\nabla_q \mathcal{L}_G(q, u) = a_u^S$$

- The gradient depends only on the discriminator
- The alignment is preserved only if the discriminator is optimal

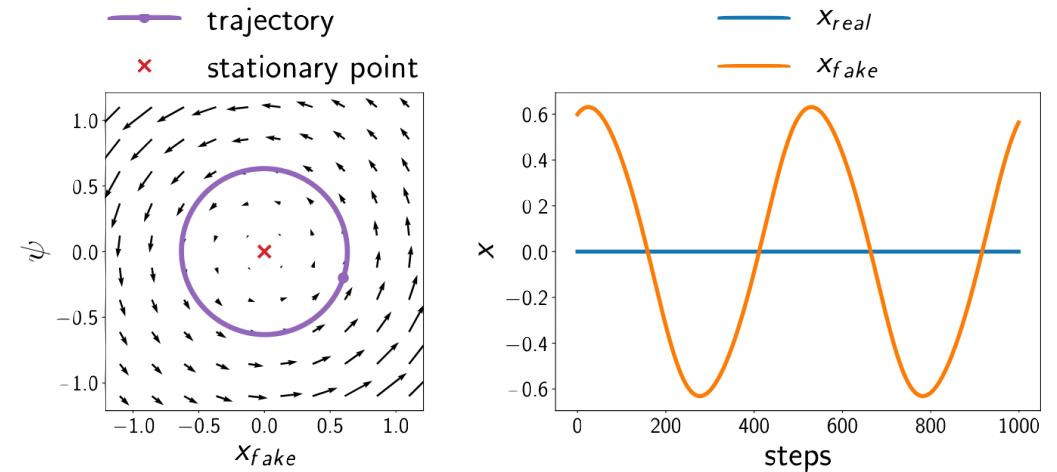
Training with **symmetric binary discriminator**  $U(\cdot, \cdot)$

$U(\cdot, \cdot)$  distinguishes between **same-distribution pairs**  $(x, y) \sim p(x)p(y), q(x)q(y)$

and **different-distribution pairs**  $(x, y) \sim p(x)q(y), p(y)q(x)$

$$\begin{aligned} \mathcal{L}_G(q, U) &= \mathbb{E}_{p \times p}[S(U(x, y))] + \mathbb{E}_{q \times q}[S(U(x, y))] - 2\mathbb{E}_{p \times q}[S(U(x, y))] \\ &= \langle q - p, A_U^S(q - p) \rangle \end{aligned}$$

$$\nabla_q \mathcal{L}_G(q, U) = 2A_U^S(q - p)$$



# How to Preserve the Alignment?

Training with **unary discriminator**  $u(\cdot)$

$u(\cdot)$  distinguishes between **real samples**  $x \sim p(x)$   
and **generated samples**  $x \sim q(x)$

$$\mathcal{L}_G(q, u) = \mathbb{E}_{q(x)}[S(u(x))] = \langle a_u^S, q \rangle$$

$$\nabla_q \mathcal{L}_G(q, u) = a_u^S$$

- The gradient depends only on the discriminator
- The alignment is preserved only if the discriminator is optimal

Training with **symmetric binary discriminator**  $U(\cdot, \cdot)$

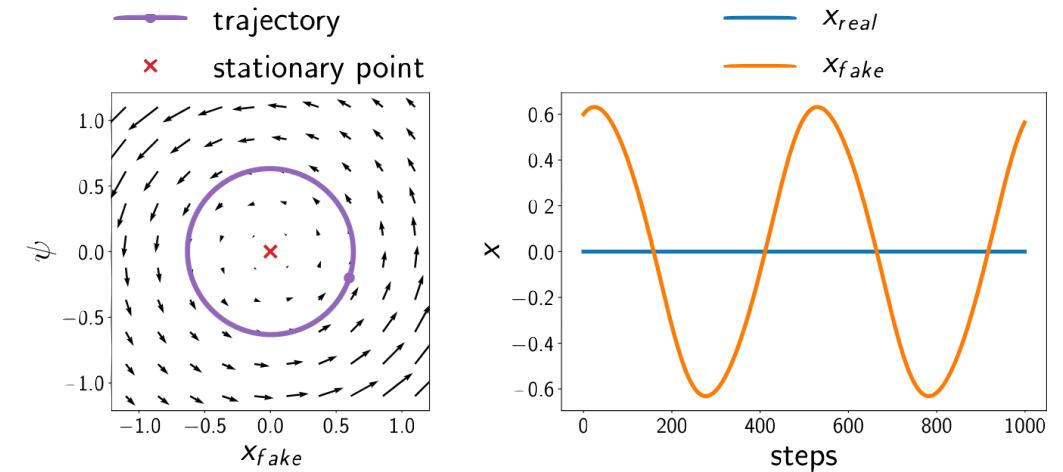
$U(\cdot, \cdot)$  distinguishes between **same-distribution pairs**  $(x, y) \sim p(x)p(y), q(x)q(y)$

and **different-distribution pairs**  $(x, y) \sim p(x)q(y), p(y)q(x)$

$$\begin{aligned} \mathcal{L}_G(q, U) &= \mathbb{E}_{p \times p}[S(U(x, y))] + \mathbb{E}_{q \times q}[S(U(x, y))] - 2\mathbb{E}_{p \times q}[S(U(x, y))] \\ &= \langle q - p, A_U^S(q - p) \rangle \end{aligned}$$

$$\nabla_q \mathcal{L}_G(q, U) = 2A_U^S(q - p)$$

- The gradient is a function of the discriminator and the deviation from the target
- The alignment, once achieved, is preserved with any discriminator



# How to Preserve the Alignment?

Training with **unary discriminator**  $u(\cdot)$

$u(\cdot)$  distinguishes between **real samples**  $x \sim p(x)$   
and **generated samples**  $x \sim q(x)$

$$\mathcal{L}_G(q, u) = \mathbb{E}_{q(x)}[S(u(x))] = \langle a_u^S, q \rangle$$

$$\nabla_q \mathcal{L}_G(q, u) = a_u^S$$

- The gradient depends only on the discriminator
- The alignment is preserved only if the discriminator is optimal

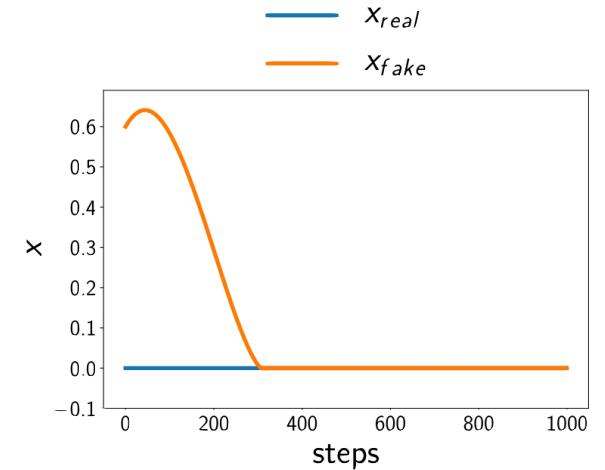
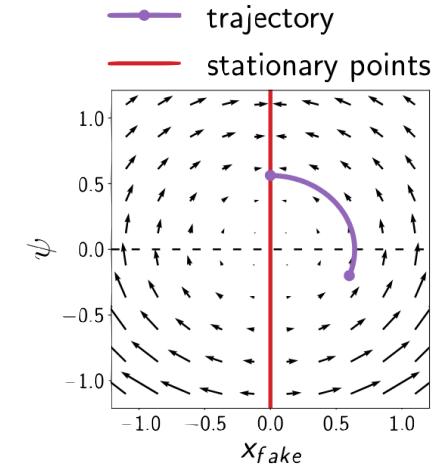
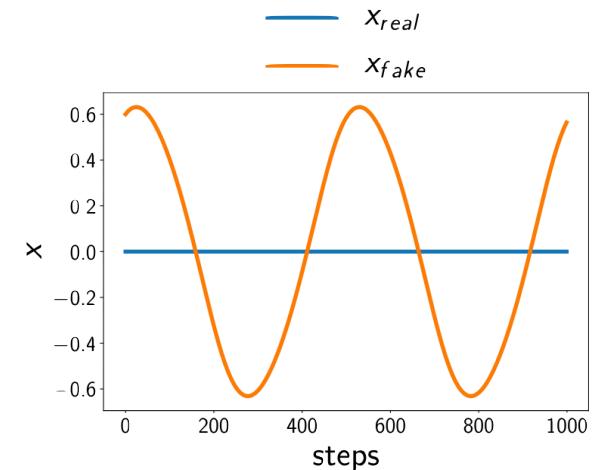
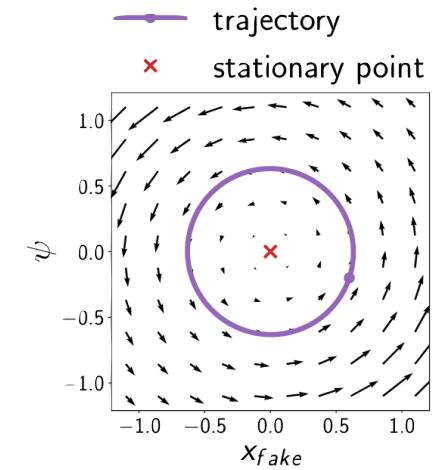
Training with **symmetric binary discriminator**  $U(\cdot, \cdot)$

$U(\cdot, \cdot)$  distinguishes between **same-distribution pairs**  $(x, y) \sim p(x)p(y), q(x)q(y)$   
and **different-distribution pairs**  $(x, y) \sim p(x)q(y), p(y)q(x)$

$$\begin{aligned} \mathcal{L}_G(q, U) &= \mathbb{E}_{p \times p}[S(U(x, y))] + \mathbb{E}_{q \times q}[S(U(x, y))] - 2\mathbb{E}_{p \times q}[S(U(x, y))] \\ &= \langle q - p, A_U^S(q - p) \rangle \end{aligned}$$

$$\nabla_q \mathcal{L}_G(q, U) = 2A_U^S(q - p)$$

- The gradient is a function of the discriminator and the deviation from the target
- The alignment, once achieved, is preserved with any discriminator



# PairGAN

A quadratic form in the space of distributions

$$\mathcal{L}_G(q, U) = \langle q - p, A_U^S(q - p) \rangle$$

corresponds to the Maximum Mean Discrepancy (MMD) distance

$$\text{MMD}^2(p, q) = \langle q - p, A(q - p) \rangle = \|q - p\|_A^2$$

if  $A$  is a positive-definite kernel

# PairGAN

A quadratic form in the space of distributions

$$\mathcal{L}_G(q, U) = \langle q - p, A_U^S(q - p) \rangle$$

corresponds to the Maximum Mean Discrepancy (MMD) distance

$$\text{MMD}^2(p, q) = \langle q - p, A(q - p) \rangle = \|q - p\|_A^2$$

if  $A$  is a positive-definite kernel

In high-dimensional spaces, manually-designed kernels

provide weak signals

# PairGAN

A quadratic form in the space of distributions

$$\mathcal{L}_G(q, U) = \langle q - p, A_U^S(q - p) \rangle$$

corresponds to the Maximum Mean Discrepancy (MMD) distance

$$\text{MMD}^2(p, q) = \langle q - p, A(q - p) \rangle = \|p - q\|_A^2$$

if  $A$  is a positive-definite kernel

Instead, an adaptive pairwise discriminator can be trained  
with the log-likelihood objective

$$\begin{aligned}\mathcal{L}_D(q, U) = & \mathbb{E}_{p \times p}[-\log \sigma(U(x, y))] + \mathbb{E}_{q \times q}[-\log \sigma(U(x, y))] \\ & + \mathbb{E}_{p \times q}[-\log \sigma(-U(x, y))] + \mathbb{E}_{q \times p}[-\log \sigma(-U(x, y))]\end{aligned}$$

$$\begin{aligned}\mathcal{L}_G(q, U) = & \mathbb{E}_{p \times p}[\log \sigma(U(x, y))] + \mathbb{E}_{q \times q}[\log \sigma(U(x, y))] \\ & - 2\mathbb{E}_{p \times q}[\log \sigma(U(x, y))]\end{aligned}$$

In high-dimensional spaces, manually-designed kernels

provide weak signals

# PairGAN

A quadratic form in the space of distributions

$$\mathcal{L}_G(q, U) = \langle q - p, A_U^S(q - p) \rangle$$

corresponds to the Maximum Mean Discrepancy (MMD) distance

$$\text{MMD}^2(p, q) = \langle q - p, A(q - p) \rangle = \|p - q\|_A^2$$

if  $A$  is a positive-definite kernel

In high-dimensional spaces, manually-designed kernels  
provide weak signals

Instead, an adaptive pairwise discriminator can be trained  
with the log-likelihood objective

$$\begin{aligned}\mathcal{L}_D(q, U) = & \mathbb{E}_{p \times p}[-\log \sigma(U(x, y))] + \mathbb{E}_{q \times q}[-\log \sigma(U(x, y))] \\ & + \mathbb{E}_{p \times q}[-\log \sigma(-U(x, y))] + \mathbb{E}_{q \times p}[-\log \sigma(-U(x, y))]\end{aligned}$$

$$\begin{aligned}\mathcal{L}_G(q, U) = & \mathbb{E}_{p \times p}[\log \sigma(U(x, y))] + \mathbb{E}_{q \times q}[\log \sigma(U(x, y))] \\ & - 2\mathbb{E}_{p \times q}[\log \sigma(U(x, y))]\end{aligned}$$

## Non-zero-sum PairGAN

$$\min_{U_\phi} \mathcal{L}_D(q_\theta, U_\phi)$$

$$\min_{q_\theta} \mathcal{L}_G(q_\theta, U_\phi)$$

# PairGAN

A quadratic form in the space of distributions

$$\mathcal{L}_G(q, U) = \langle q - p, A_U^S(q - p) \rangle$$

corresponds to the Maximum Mean Discrepancy (MMD) distance

$$\text{MMD}^2(p, q) = \langle q - p, A(q - p) \rangle = \|p - q\|_A^2$$

if  $A$  is a positive-definite kernel

In high-dimensional spaces, manually-designed kernels  
provide weak signals

Instead, an adaptive pairwise discriminator can be trained  
with the log-likelihood objective

$$\begin{aligned}\mathcal{L}_D(q, U) = & \mathbb{E}_{p \times p}[-\log \sigma(U(x, y))] + \mathbb{E}_{q \times q}[-\log \sigma(U(x, y))] \\ & + \mathbb{E}_{p \times q}[-\log \sigma(-U(x, y))] + \mathbb{E}_{q \times p}[-\log \sigma(-U(x, y))]\end{aligned}$$

$$\begin{aligned}\mathcal{L}_G(q, U) = & \mathbb{E}_{p \times p}[\log \sigma(U(x, y))] + \mathbb{E}_{q \times q}[\log \sigma(U(x, y))] \\ & - 2\mathbb{E}_{p \times q}[\log \sigma(U(x, y))]\end{aligned}$$

## Non-zero-sum PairGAN

$$\min_{U_\phi} \mathcal{L}_D(q_\theta, U_\phi)$$

$$\min_{q_\theta} \mathcal{L}_G(q_\theta, U_\phi)$$

## Zero-sum PairGAN

$$\max_{U_\phi} \mathcal{L}_G(q_\theta, U_\phi)$$

$$\min_{q_\theta} \mathcal{L}_G(q_\theta, U_\phi)$$

# Divergence Minimization

$$\begin{aligned}\mathcal{L}_D(q, U) = & \mathbb{E}_{p \times p} [-\log \sigma(U(x, y))] + \mathbb{E}_{q \times q} [-\log \sigma(U(x, y))] \\ & + \mathbb{E}_{p \times q} [-\log \sigma(-U(x, y))] + \mathbb{E}_{q \times p} [-\log \sigma(-U(x, y))]\end{aligned}$$

$$\begin{aligned}\mathcal{L}_G(q, U) = & \mathbb{E}_{p \times p} [\log \sigma(U(x, y))] + \mathbb{E}_{q \times q} [\log \sigma(U(x, y))] \\ & - 2\mathbb{E}_{p \times q} [\log \sigma(U(x, y))]\end{aligned}$$

## Non-zero-sum PairGAN

$$\min_{U_\phi} \mathcal{L}_D(q_\theta, U_\phi)$$

$$\min_{q_\theta} \mathcal{L}_G(q_\theta, U_\phi)$$

## Zero-sum PairGAN

$$\max_{U_\phi} \mathcal{L}_G(q_\theta, U_\phi)$$

$$\min_{q_\theta} \mathcal{L}_G(q_\theta, U_\phi)$$

# Divergence Minimization

$$\begin{aligned}\mathcal{L}_D(q, U) &= \mathbb{E}_{p \times p}[-\log \sigma(U(x, y))] + \mathbb{E}_{q \times q}[-\log \sigma(U(x, y))] \\ &\quad + \mathbb{E}_{p \times q}[-\log \sigma(-U(x, y))] + \mathbb{E}_{q \times p}[-\log \sigma(-U(x, y))]\end{aligned}$$

$$\begin{aligned}\mathcal{L}_G(q, U) &= \mathbb{E}_{p \times p}[\log \sigma(U(x, y))] + \mathbb{E}_{q \times q}[\log \sigma(U(x, y))] \\ &\quad - 2\mathbb{E}_{p \times q}[\log \sigma(U(x, y))]\end{aligned}$$

## Non-zero-sum PairGAN

$$\min_{U_\phi} \mathcal{L}_D(q_\theta, U_\phi)$$

$$\min_{q_\theta} \mathcal{L}_G(q_\theta, U_\phi)$$

## Zero-sum PairGAN

$$\max_{U_\phi} \mathcal{L}_G(q_\theta, U_\phi)$$

$$\min_{q_\theta} \mathcal{L}_G(q_\theta, U_\phi)$$

**Proposition** (PairGAN Distribution Divergences).

Consider distributions  $p(x)$  and  $q(x)$ . Let  $M_{p,q}^+(x, y)$ ,  $M_{p,q}^-(x, y)$ , and  $M_{p,q}(x, y)$  denote the mixture distributions over pairs

$$M_{p,q}^+(x, y) = \frac{1}{2}p(x)p(y) + \frac{1}{2}q(x)q(y)$$

$$M_{p,q}^-(x, y) = \frac{1}{2}p(x)q(y) + \frac{1}{2}q(x)p(y)$$

$$M_{p,q}(x, y) = \frac{1}{2}M_{p,q}^+(x, y) + \frac{1}{2}M_{p,q}^-(x, y)$$

Given optimal PairGAN discriminators (zero-sum or non-zero-sum), the PairGAN generator objectives are equivalent to the following distribution divergences

$$\left. \begin{array}{l} U_N^*(q) = \underset{U: \mathcal{X} \rightarrow \mathbb{R}}{\operatorname{argmin}} \mathcal{L}_D(q, U) \\ \mathcal{L}_G(q, U_N^*(q)) = 4 \cdot (\text{KL}(M_{p,q}^+ \| M_{p,q}) + \text{KL}(M_{p,q} \| M_{p,q}^+)) \end{array} \right| \begin{array}{l} U_Z^*(q) = \underset{U: \mathcal{X} \rightarrow [\log(\varepsilon), \infty)}{\operatorname{argmax}} \mathcal{L}_G(q, U), \quad (0 < \varepsilon < 1) \\ \mathcal{L}_G(q, U_Z^*(q)) = -\log(\varepsilon) \cdot \text{TV}(M_{p,q}^+ \| M_{p,q}^-). \end{array}$$

Consequently,

$$\mathcal{L}_G(q, U^*(q)) \geq 0$$

$$\mathcal{L}_G(q, U^*(q)) = 0 \text{ if and only if } q = p$$

# **Experimental validation**

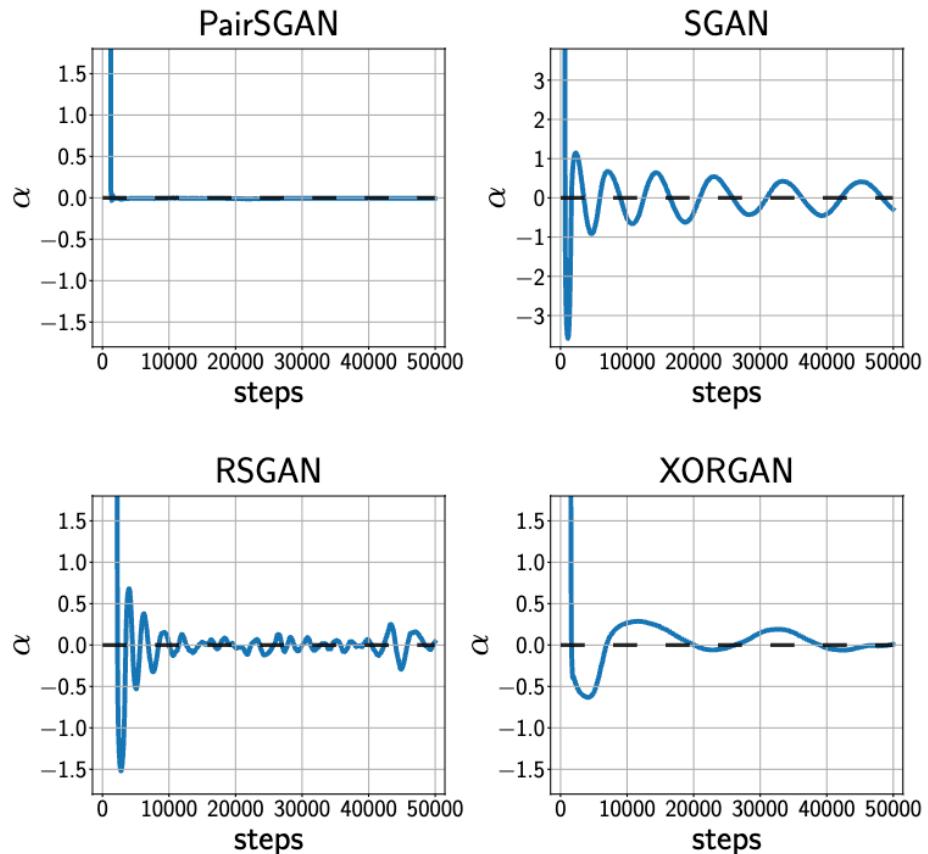
# Experimental validation

- ▶ Verified alignment preservation in DCGAN with restricted generator

# Experimental validation

- Verified alignment preservation in DCGAN with restricted generator

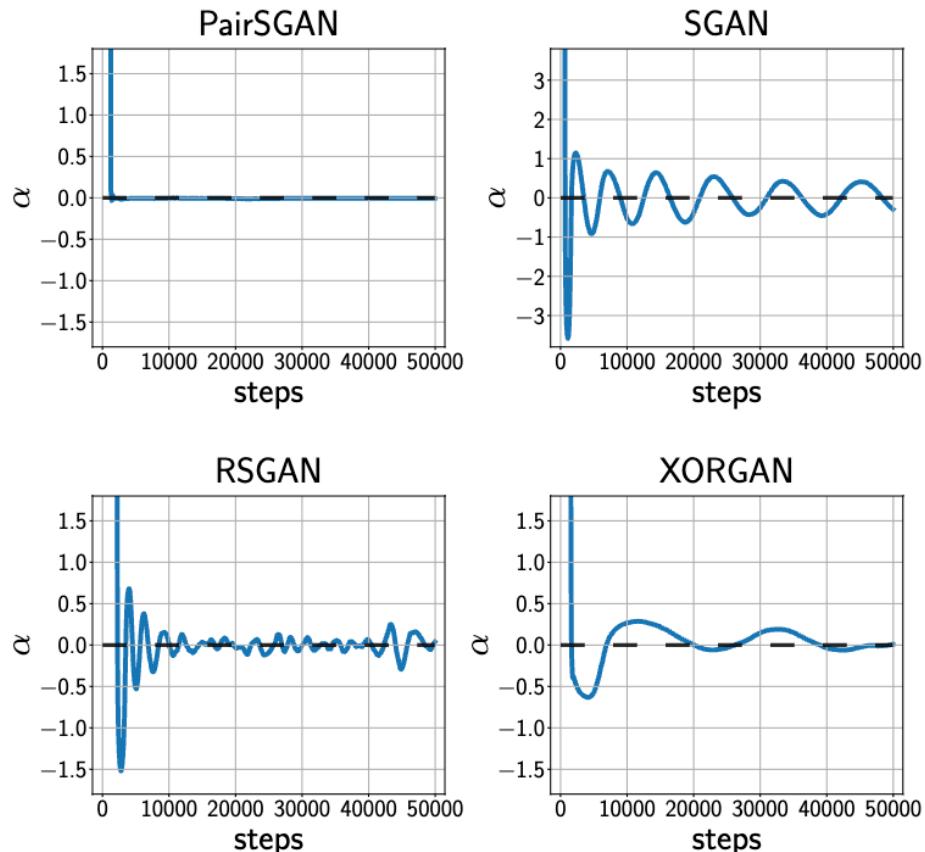
Alignment stability in  
restricted DCGAN generator experiment



# Experimental validation

- ▶ Verified alignment preservation in DCGAN with restricted generator
- ▶ Improved FID curve stability on CIFAR-10 and CAT benchmarks

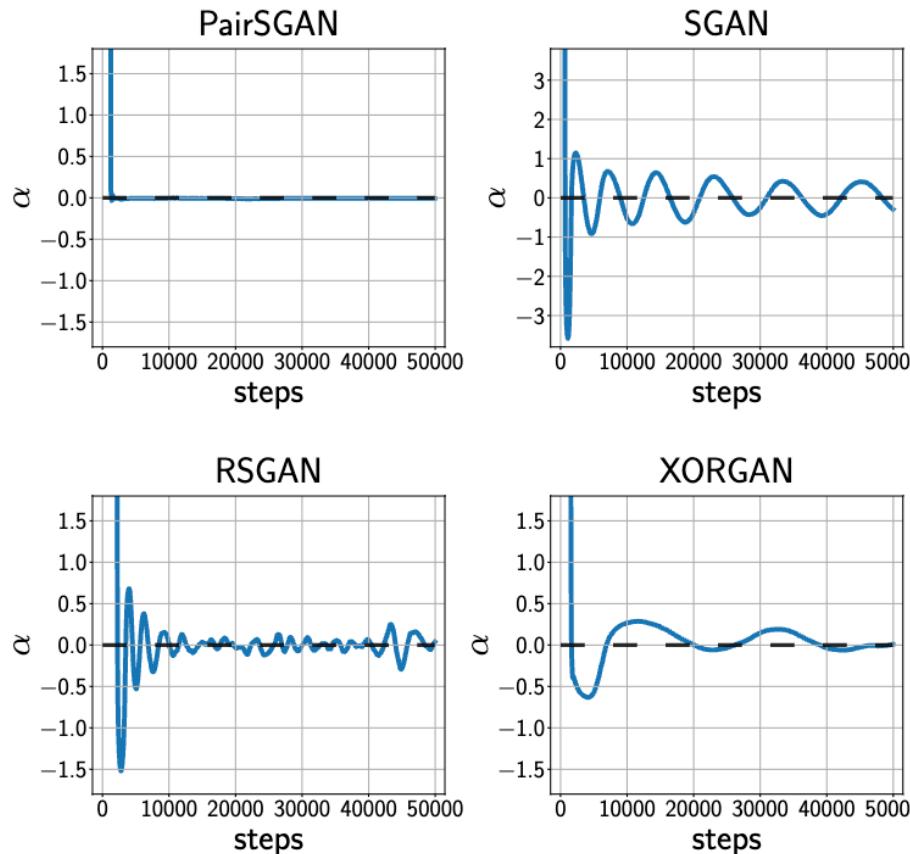
Alignment stability in  
restricted DCGAN generator experiment



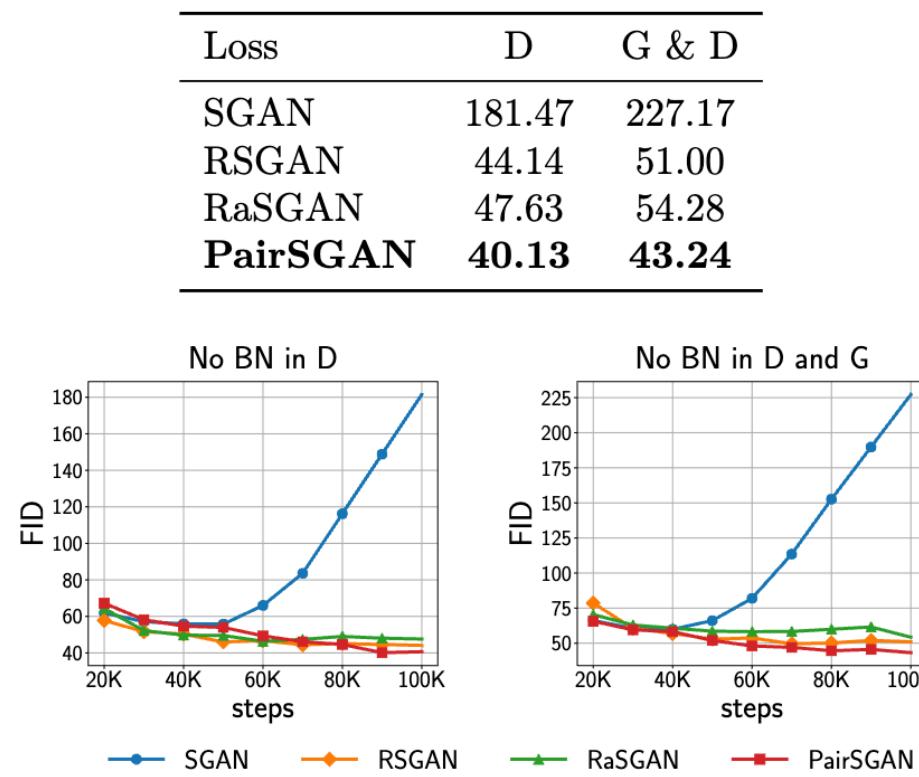
# Experimental validation

- ▶ Verified alignment preservation in DCGAN with restricted generator
- ▶ Improved FID curve stability on CIFAR-10 and CAT benchmarks

Alignment stability in  
restricted DCGAN generator experiment



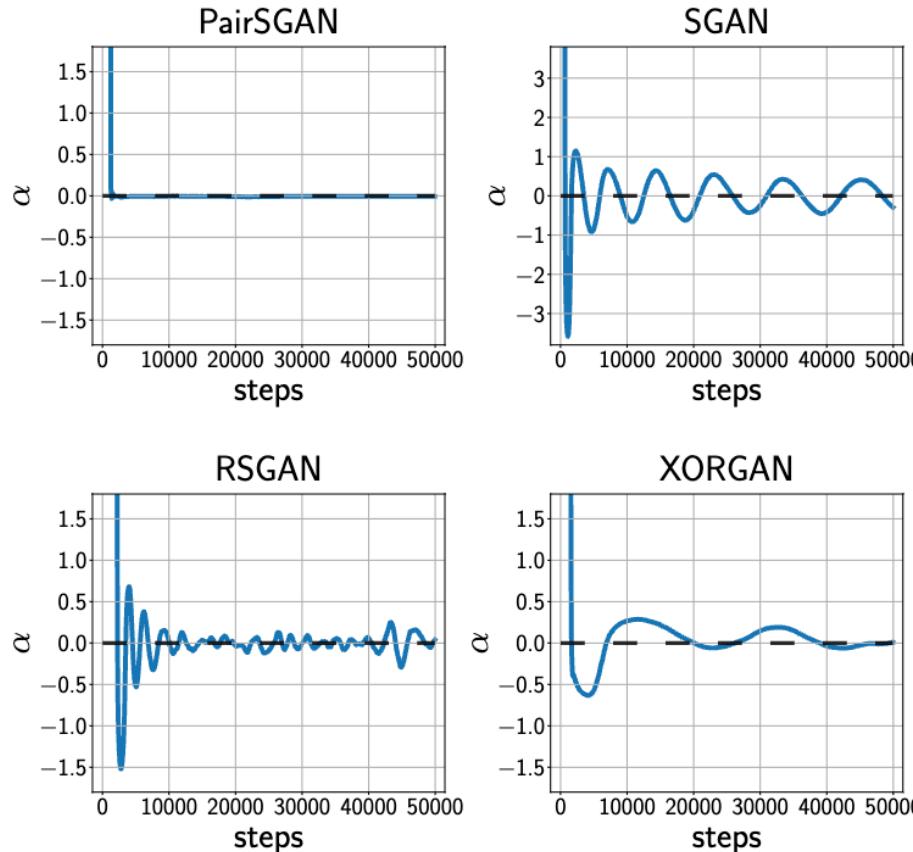
Improved FID on CIFAR 10  
without BN in D / G&D



# Experimental validation

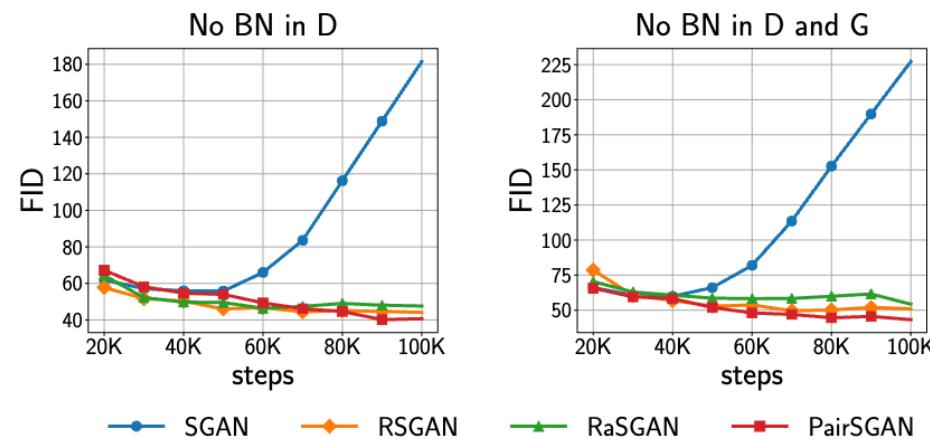
- ▶ Verified alignment preservation in DCGAN with restricted generator
- ▶ Improved FID curve stability on CIFAR-10 and CAT benchmarks

**Alignment stability in restricted DCGAN generator experiment**



**Improved FID on CIFAR 10 without BN in D / G&D**

Loss	D	G & D
SGAN	181.47	227.17
RSGAN	44.14	51.00
RaSGAN	47.63	54.28
<b>PairSGAN</b>	<b>40.13</b>	<b>43.24</b>



**Stable FID curves on CAT data**

64 × 64 images				
Loss	Min	Max	Mean	SD
SGAN	12.27	24.62	16.99	4.07
RSGAN*	19.03	42.05	32.16	7.01
RaSGAN*	15.38	33.11	20.53	5.68
LSGAN*	20.27	224.97	73.62	61.02
RaLSGAN*	11.97	19.29	15.61	2.55
HingeGAN*	17.60	50.94	32.23	14.44
RaHingeGAN*	14.62	27.31	20.29	3.96
RSGAN-GP*	16.41	22.34	18.20	1.82
RaSGAN-GP*	17.32	22	19.58	<b>1.81</b>
<b>PairSGAN</b>	<b>10.28</b>	<b>18.21</b>	<b>13.55</b>	2.24
128 × 128 images				
Loss	Min	Max	Mean	SD
SGAN	19.88	38.68	28.91	6.73
RaSGAN*	21.05	39.65	28.53	6.52
LSGAN*	19.03	51.36	30.28	10.16
RaLSGAN*	<b>15.85</b>	40.26	22.36	7.53
<b>PairSGAN</b>	16.72	<b>25.66</b>	<b>21.43</b>	2.94
256 × 256 images				
Loss	Min	Max	Mean	SD
SGAN	43.30	324.38	171.42	108.47
RaSGAN*	<b>32.11</b>	102.76	56.64	21.03
SpectralSGAN*	54.08	90.43	64.92	12.00
LSGAN*	—	—	—	—
RaLSGAN*	35.21	299.52	70.44	86.01
WGAN-GP*	155.46	437.48	341.91	101.11
<b>PairSGAN</b>	33.94	<b>50.52</b>	<b>41.70</b>	<b>5.23</b>

# Pairwise-Discriminator Objectives for GANs: Summary

Training objectives for GANs which **preserve distribution alignment**

- ▶ Novel zero-sum & non-zero-sum game formulations with pairwise discriminators
  - ▶ New distribution divergences derived from PairGAN game
- ▶ Equilibria stability analysis for parametric generators
  - ▶ Introduced the notion of sufficient discriminators for given parametric generator family
  - ▶ Constructed examples of minimally sufficient discriminators for arbitrary generator families
  - ▶ Established connections to non-parametric MMD objectives and MMD-GAN
- ▶ Extension to multiple distribution alignment scenario

Experimental validation

- ▶ Demonstrated alignment preserving property in DCGAN under restricted generator parameterization
- ▶ Improved FID curve stability on CIFAR-10 and CAT benchmarks

---

The Benefits of Pairwise Discriminators for Adversarial Training

S. Tong\*, T. Garipov\*, T. Jaakkola (arXiv Pre-print, 2020)

# Chapter III

## Adversarial Support Alignment

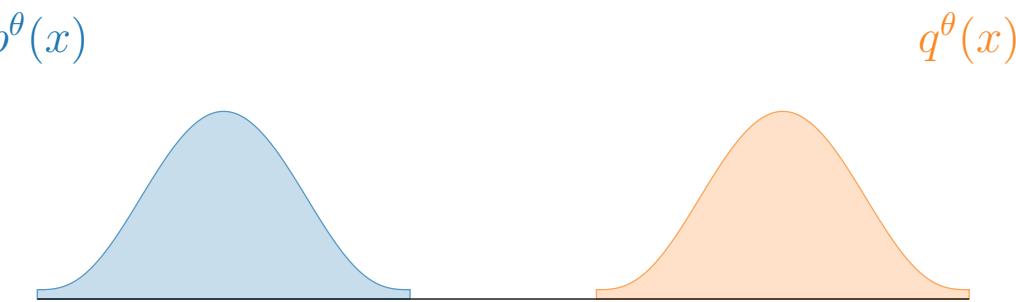
---

Adversarial Support Alignment

S. Tong\*, T. Garipov\*, Y. Zhang, S. Chang, T. Jaakkola (ICLR 2022, Spotlight)

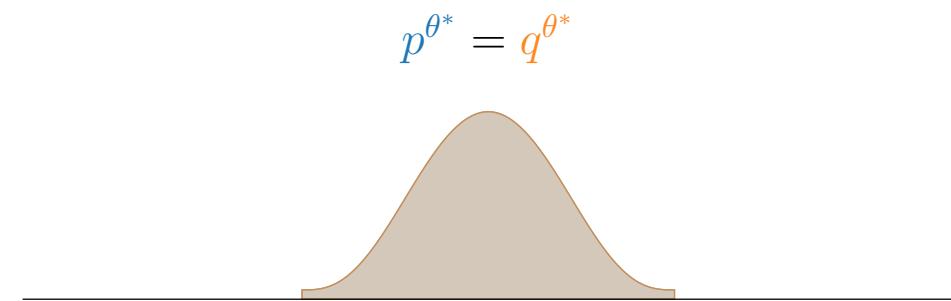
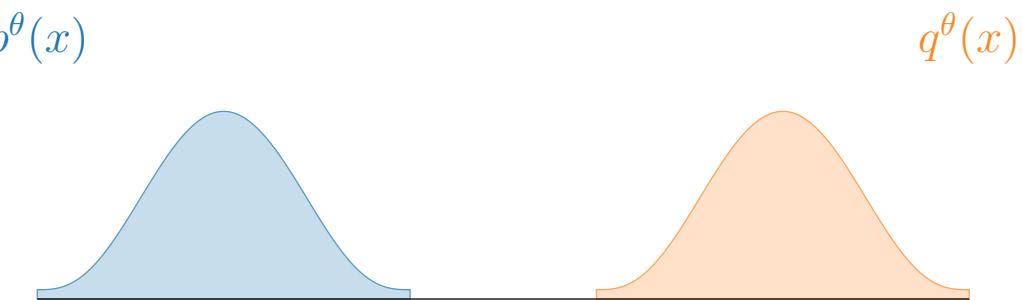
## Distribution Alignment

**Given**  $\mathcal{P} = \{p^\theta \mid \theta \in \Theta\}$   $\mathcal{Q} = \{q^\theta \mid \theta \in \Theta\}$   
**Find**  $\theta^* : p^{\theta^*} = q^{\theta^*}$

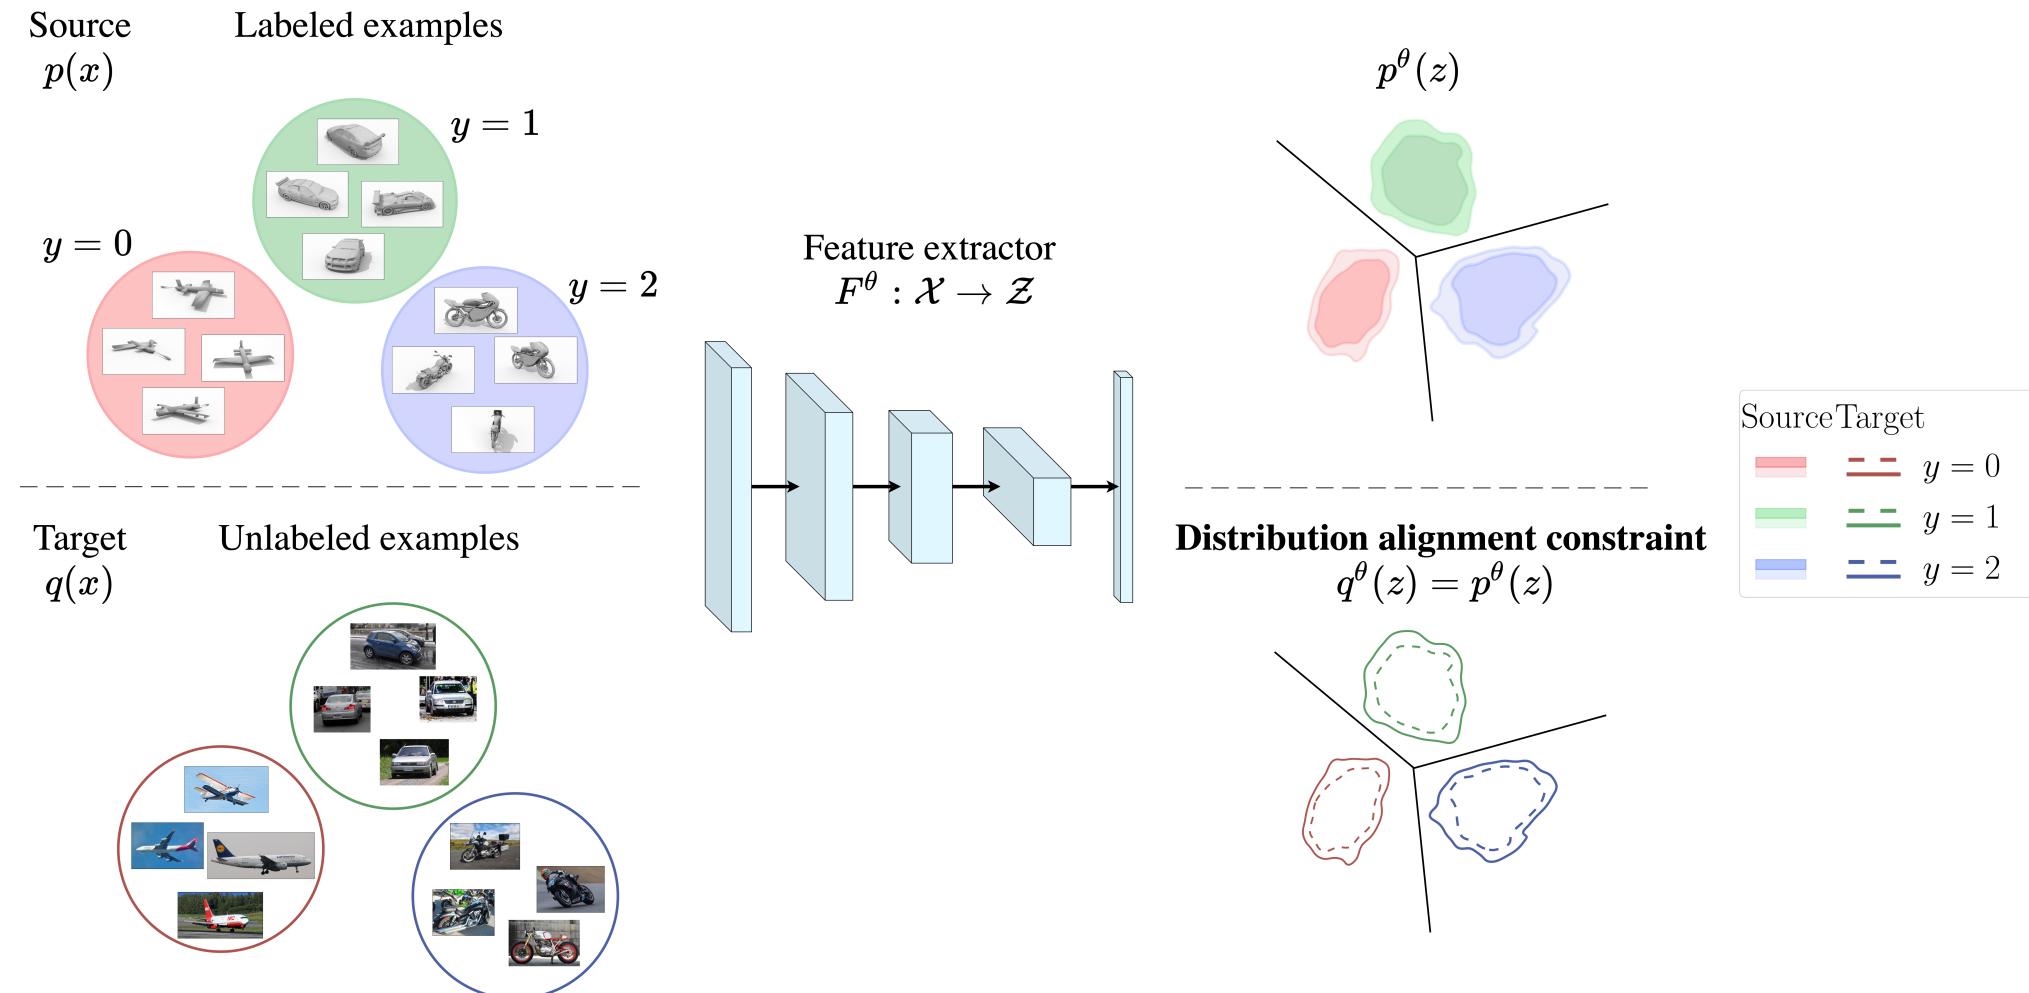


## Distribution Alignment

**Given**  $\mathcal{P} = \{p^\theta \mid \theta \in \Theta\}$   $\mathcal{Q} = \{q^\theta \mid \theta \in \Theta\}$   
**Find**  $\theta^* : p^{\theta^*} = q^{\theta^*}$



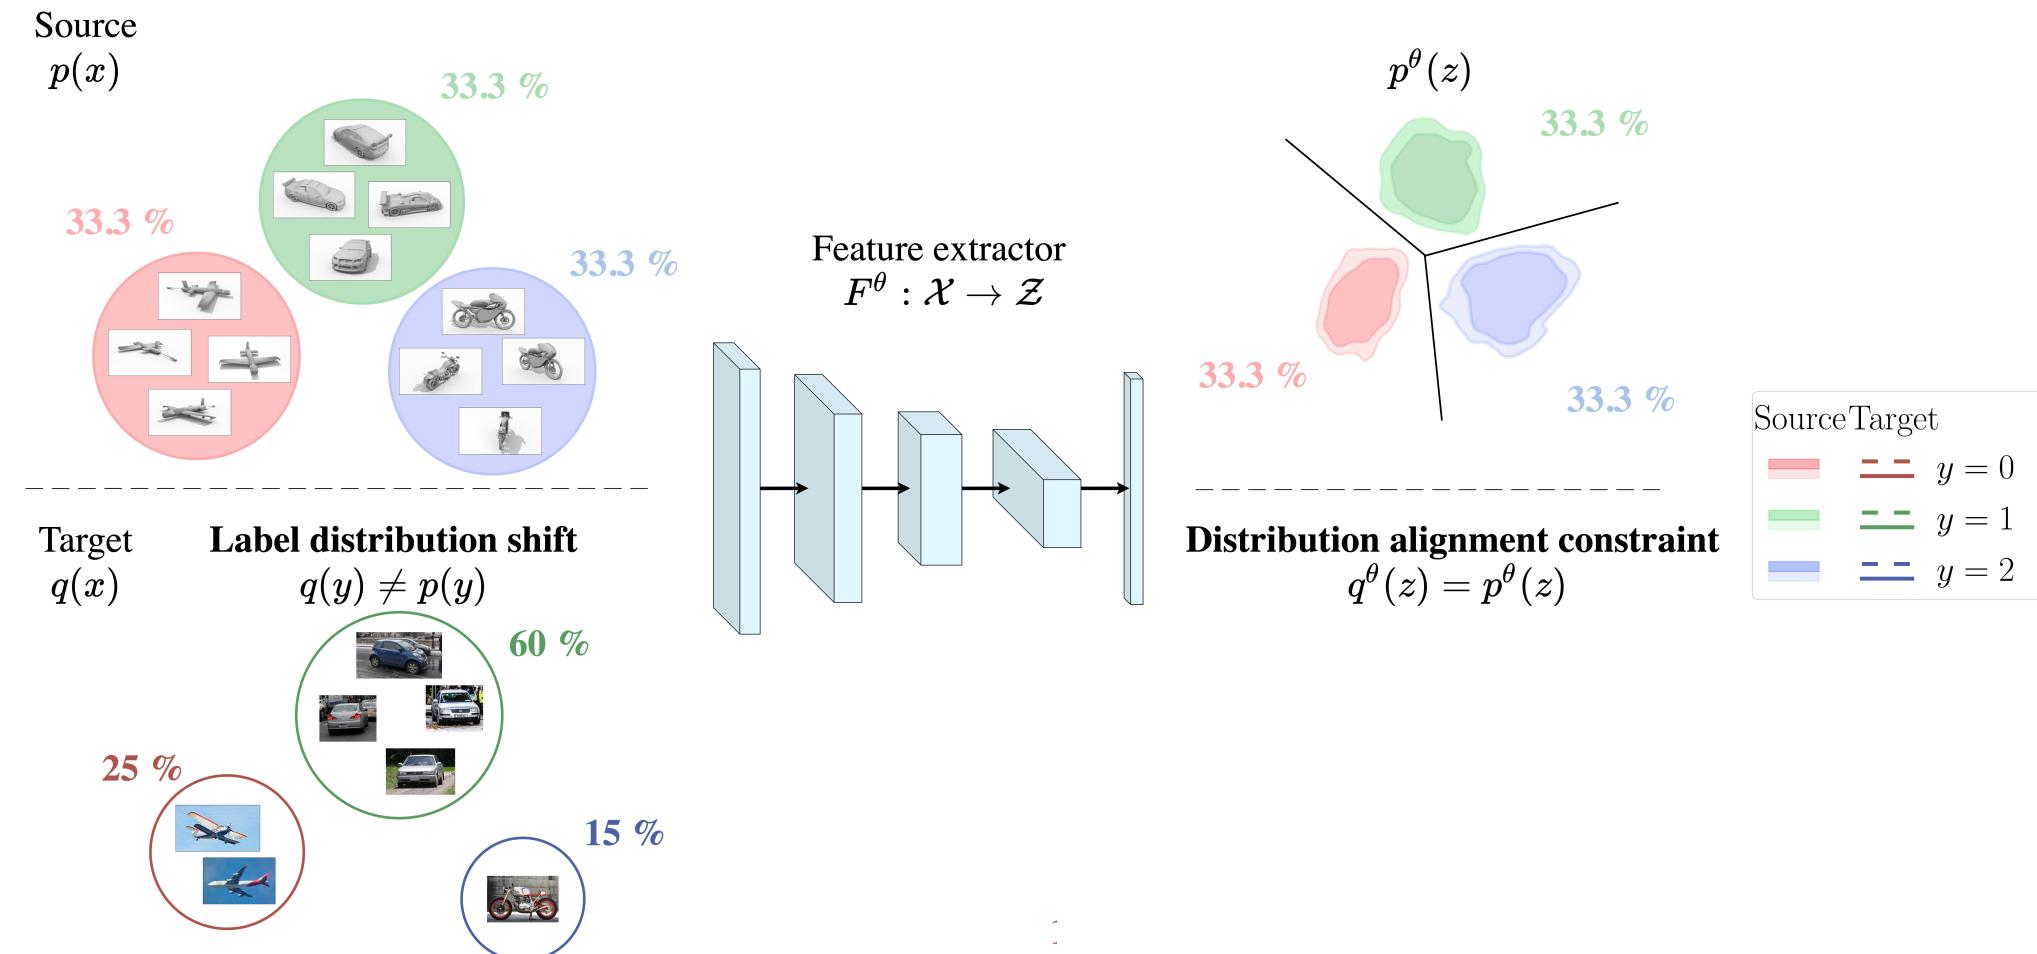
# Distribution Alignment for Domain Adaptation



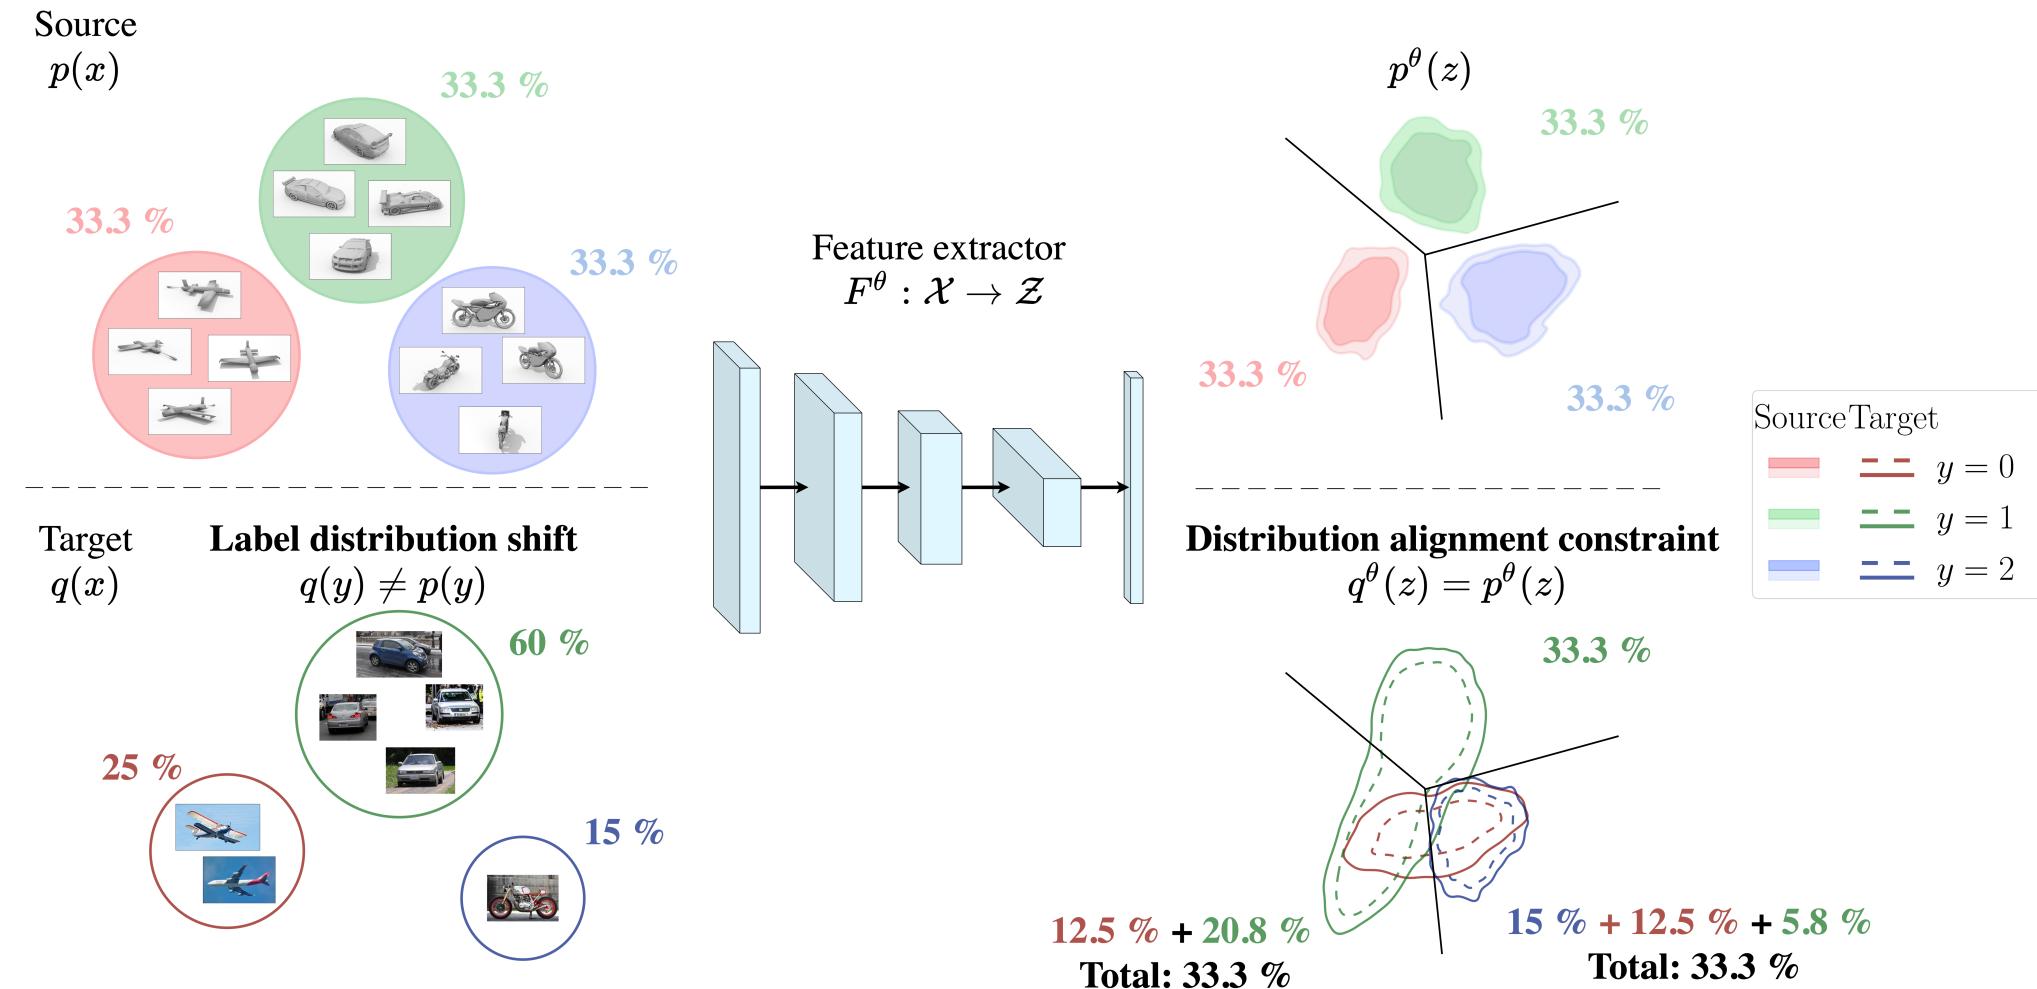
Two training objectives:

- ▶ Source domain classification loss
- ▶ Source vs target discrimination loss (GAN-like game with a discriminator)

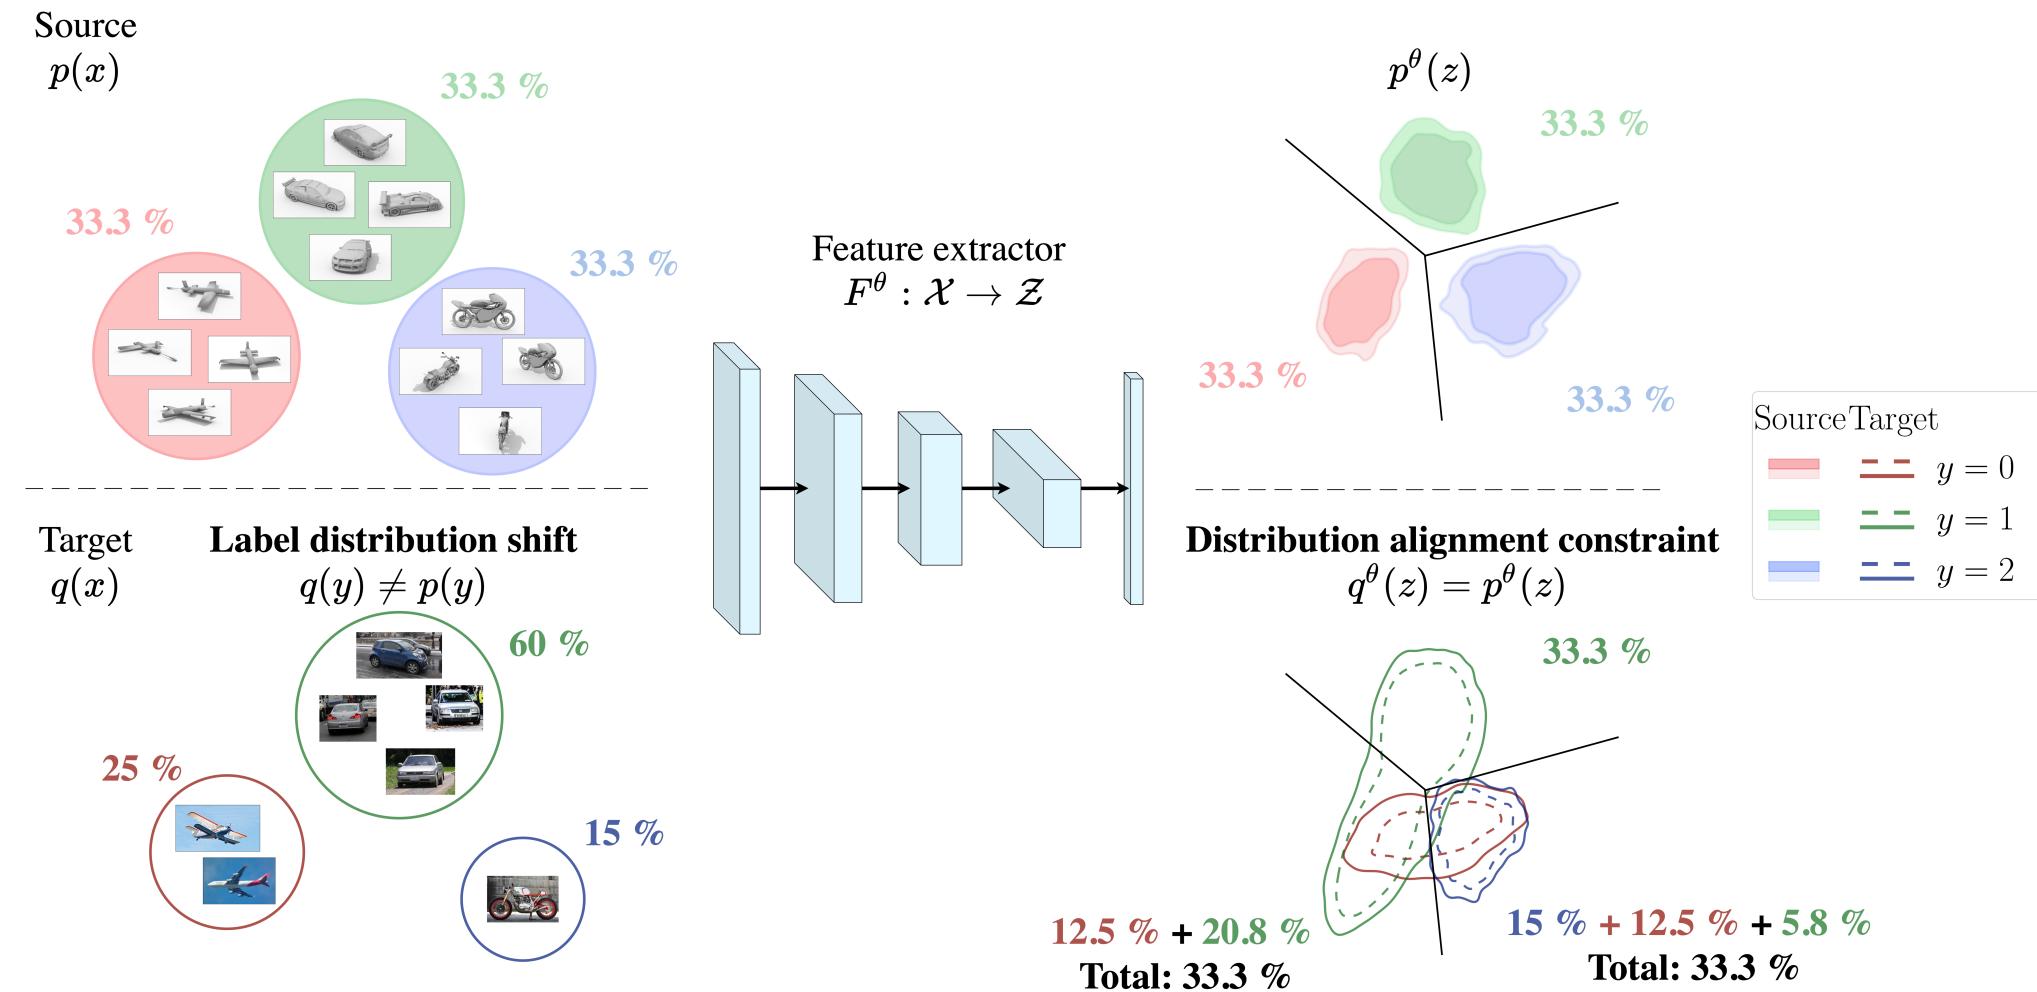
# Issue: Label Distribution Shift



# Issue: Label Distribution Shift



# Issue: Label Distribution Shift

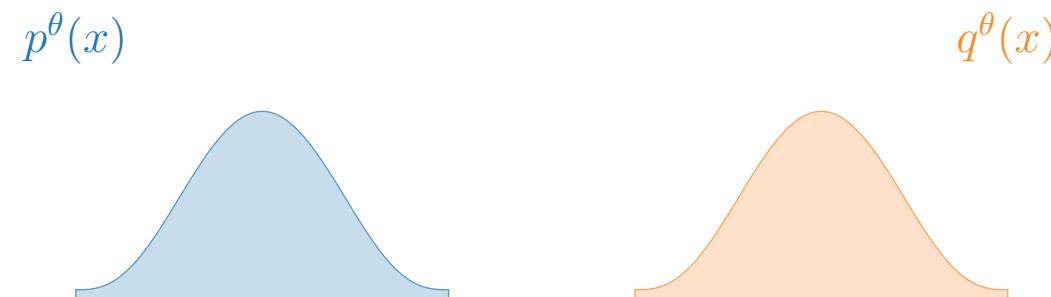


**strict distribution alignment → source-target class mismatch → degraded accuracy**

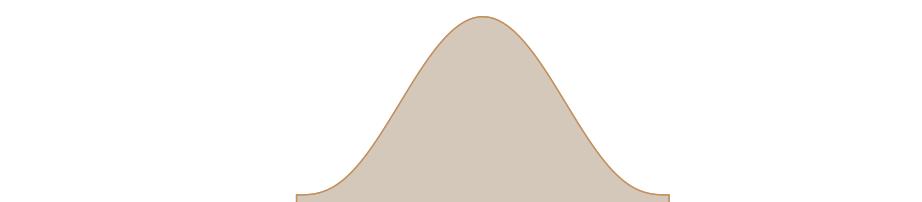
## Distribution Alignment

**Given**  $\mathcal{P} = \{p^\theta \mid \theta \in \Theta\}$   $\mathcal{Q} = \{q^\theta \mid \theta \in \Theta\}$

**Find**  $\theta^* : p^{\theta^*} = q^{\theta^*}$



$$p^{\theta^*} = q^{\theta^*}$$



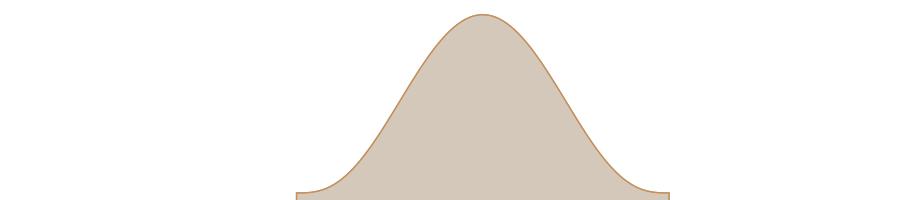
## How can we lift constraints of strict distribution alignment?

### Distribution Alignment

**Given**  $\mathcal{P} = \{p^\theta \mid \theta \in \Theta\}$     $\mathcal{Q} = \{q^\theta \mid \theta \in \Theta\}$   
**Find**  $\theta^* : p^{\theta^*} = q^{\theta^*}$



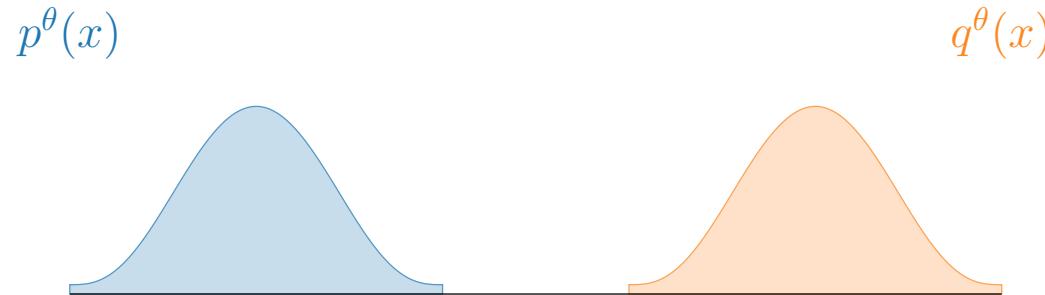
$$p^{\theta^*} = q^{\theta^*}$$



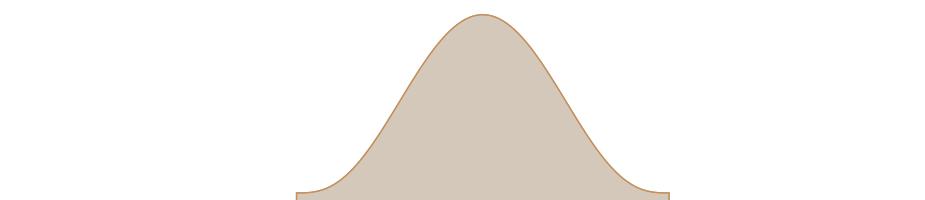
## How can we lift constraints of strict distribution alignment?

### Distribution Alignment

**Given**  $\mathcal{P} = \{p^\theta \mid \theta \in \Theta\}$   $\mathcal{Q} = \{q^\theta \mid \theta \in \Theta\}$   
**Find**  $\theta^* : p^{\theta^*} = q^{\theta^*}$



$$p^{\theta^*} = q^{\theta^*}$$



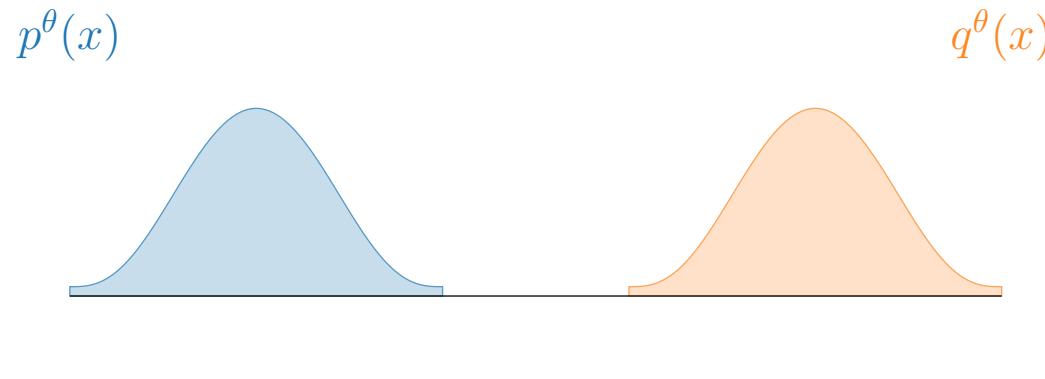
### Support Alignment

**Given**  $\mathcal{P} = \{p^\theta \mid \theta \in \Theta\}$   $\mathcal{Q} = \{q^\theta \mid \theta \in \Theta\}$   
**Find**  $\theta^* : \text{supp}(p^{\theta^*}) = \text{supp}(q^{\theta^*})$

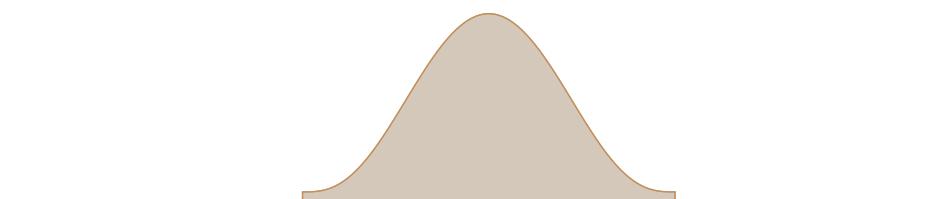
## How can we lift constraints of strict distribution alignment?

### Distribution Alignment

**Given**  $\mathcal{P} = \{p^\theta \mid \theta \in \Theta\}$   $\mathcal{Q} = \{q^\theta \mid \theta \in \Theta\}$   
**Find**  $\theta^* : p^{\theta^*} = q^{\theta^*}$

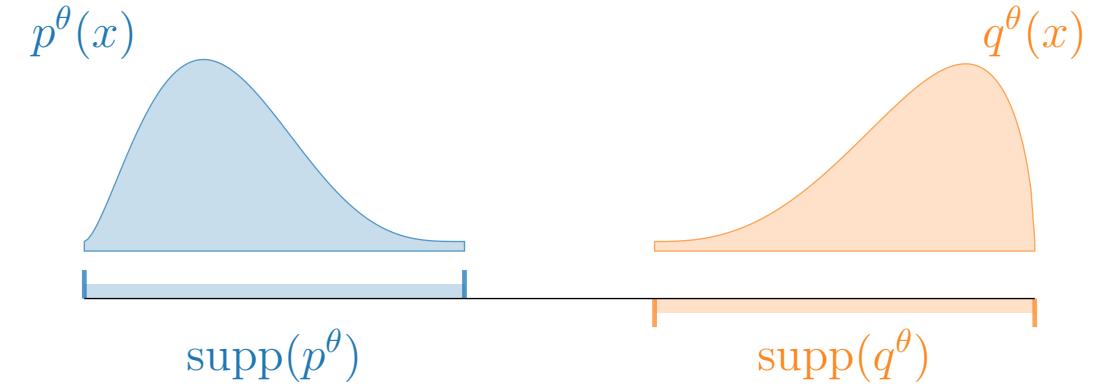


$$p^{\theta^*} = q^{\theta^*}$$



### Support Alignment

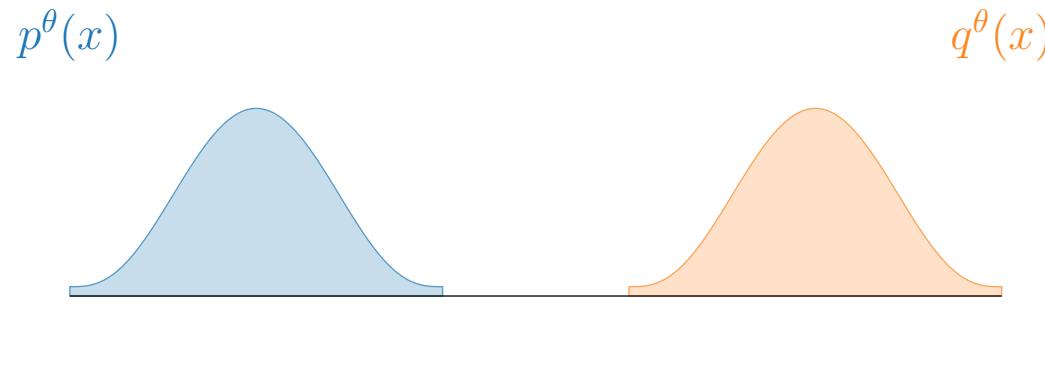
**Given**  $\mathcal{P} = \{p^\theta \mid \theta \in \Theta\}$   $\mathcal{Q} = \{q^\theta \mid \theta \in \Theta\}$   
**Find**  $\theta^* : \text{supp}(p^{\theta^*}) = \text{supp}(q^{\theta^*})$



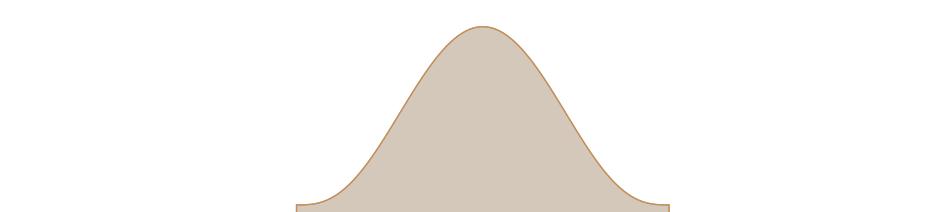
## How can we lift constraints of strict distribution alignment?

### Distribution Alignment

**Given**  $\mathcal{P} = \{p^\theta \mid \theta \in \Theta\}$   $\mathcal{Q} = \{q^\theta \mid \theta \in \Theta\}$   
**Find**  $\theta^* : p^{\theta^*} = q^{\theta^*}$

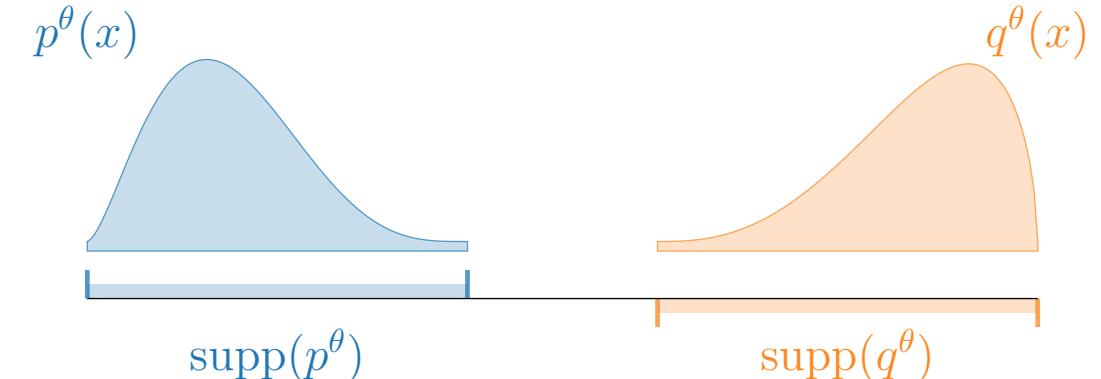


$$p^{\theta^*} = q^{\theta^*}$$

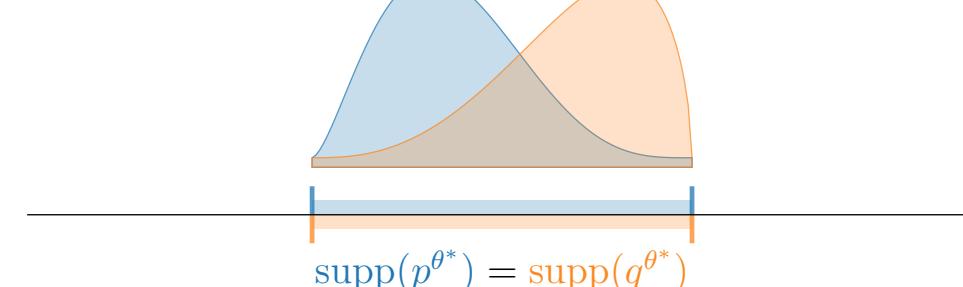


### Support Alignment

**Given**  $\mathcal{P} = \{p^\theta \mid \theta \in \Theta\}$   $\mathcal{Q} = \{q^\theta \mid \theta \in \Theta\}$   
**Find**  $\theta^* : \text{supp}(p^{\theta^*}) = \text{supp}(q^{\theta^*})$



$$p^{\theta^*}(x) = q^{\theta^*}(x)$$



# Method: Support Alignment via Log-Loss Discriminator

$$\sup_{f:\mathcal{X} \rightarrow [0,1]} \mathbb{E}_{x \sim p} [\log f(x)] + \mathbb{E}_{y \sim q} [\log(1 - f(y))]$$

# Method: Support Alignment via Log-Loss Discriminator

$$\sup_{f:\mathcal{X} \rightarrow [0,1]} \mathbb{E}_{x \sim p} [\log f(x)] + \mathbb{E}_{y \sim q} [\log(1 - f(y))]$$

$$f^*(x) = \frac{p(x)}{p(x) + q(x)}$$

# Method: Support Alignment via Log-Loss Discriminator

$$\sup_{f:\mathcal{X} \rightarrow [0,1]} \mathbb{E}_{x \sim p} [\log f(x)] + \mathbb{E}_{y \sim q} [\log(1 - f(y))]$$

$$f^*(x) = \frac{p(x)}{p(x) + q(x)}$$

## Theorem

The mapping  $f^* : \mathcal{X} \rightarrow [0, 1]$   
realized by the optimal discriminator  
preserves support discrepancy

$$\text{supp}(p) = \text{supp}(q) \Leftrightarrow \text{supp}(f^* \sharp p) = \text{supp}(f^* \sharp q)$$

$f^* \sharp p, f^* \sharp q$  — pushforward distributions

# Method: Support Alignment via Log-Loss Discriminator

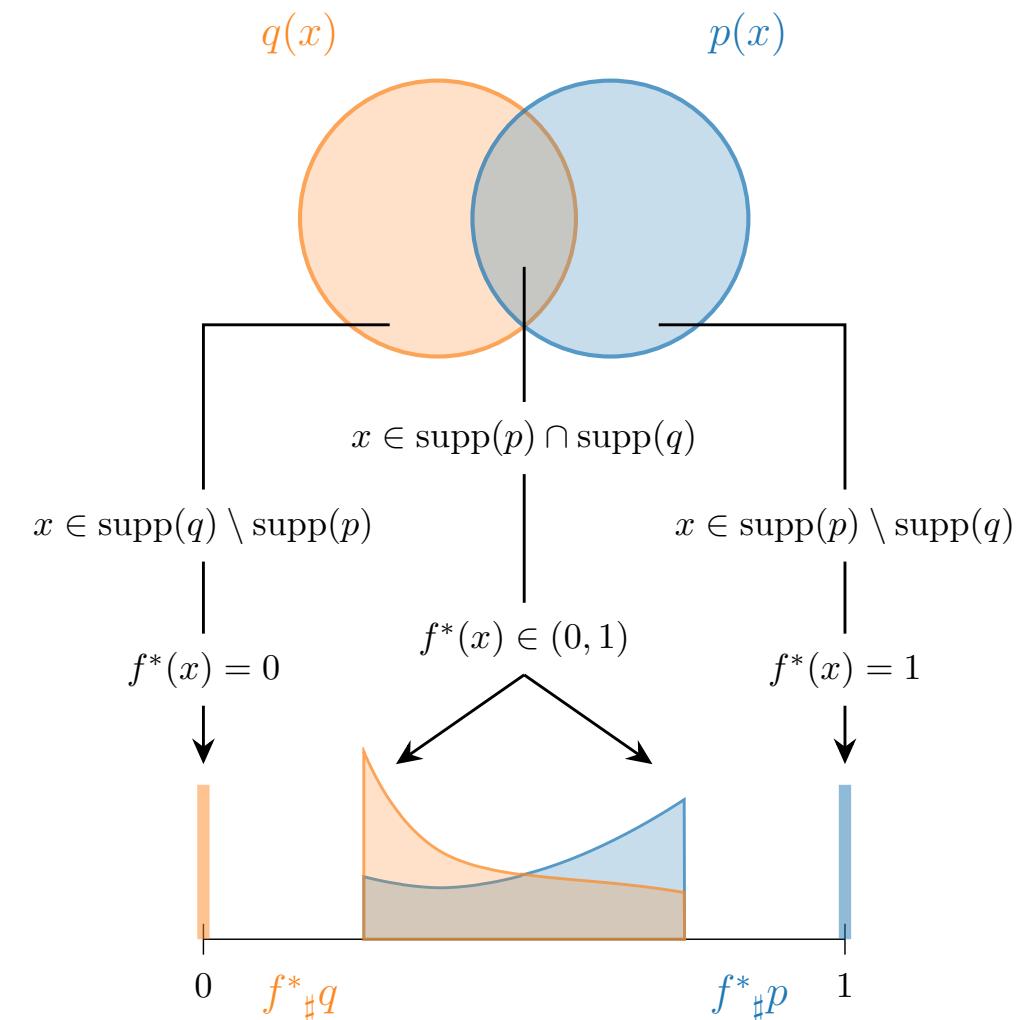
$$\sup_{f:\mathcal{X} \rightarrow [0,1]} \mathbb{E}_{x \sim p} [\log f(x)] + \mathbb{E}_{y \sim q} [\log(1 - f(y))]$$

$$f^*(x) = \frac{p(x)}{p(x) + q(x)}$$

## Theorem

The mapping  $f^* : \mathcal{X} \rightarrow [0, 1]$   
realized by the optimal discriminator  
preserves support discrepancy

$\text{supp}(p) = \text{supp}(q) \Leftrightarrow \text{supp}(f^*\#p) = \text{supp}(f^*\#q)$   
 $f^*\#p, f^*\#q$  — pushforward distributions



# Method: Support Alignment via Log-Loss Discriminator

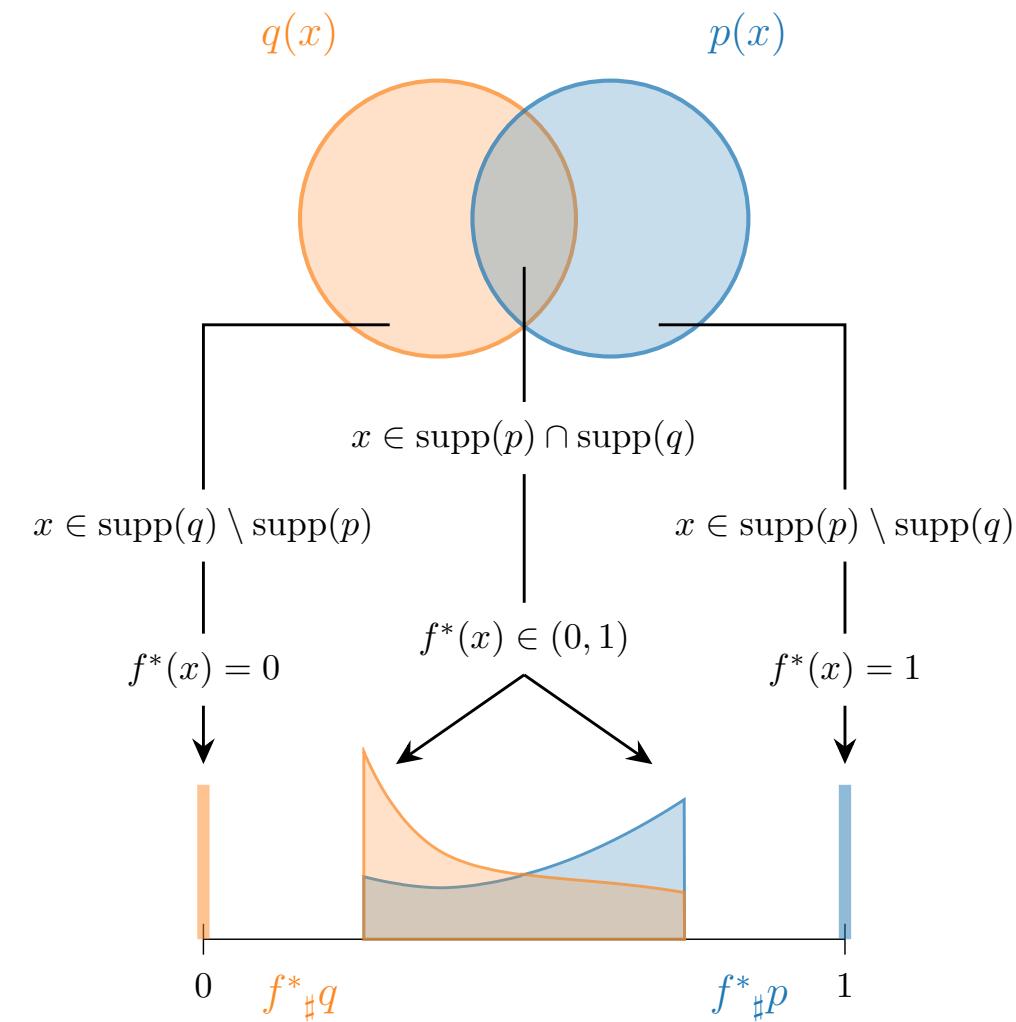
$$\sup_{f:\mathcal{X} \rightarrow [0,1]} \mathbb{E}_{x \sim p} [\log f(x)] + \mathbb{E}_{y \sim q} [\log(1 - f(y))]$$

$$f^*(x) = \frac{p(x)}{p(x) + q(x)}$$

## Theorem

The mapping  $f^* : \mathcal{X} \rightarrow [0, 1]$   
realized by the optimal discriminator  
preserves support discrepancy

$\text{supp}(p) = \text{supp}(q) \Leftrightarrow \text{supp}(f^*\#p) = \text{supp}(f^*\#q)$   
 $f^*\#p, f^*\#q$  — pushforward distributions



**Remark:** the result also holds for  $g : \mathcal{X} \rightarrow \mathbb{R}$

$$g(x) : f(x) = \text{sigmoid}(g(x))$$

# Method: Support Alignment via Log-Loss Discriminator

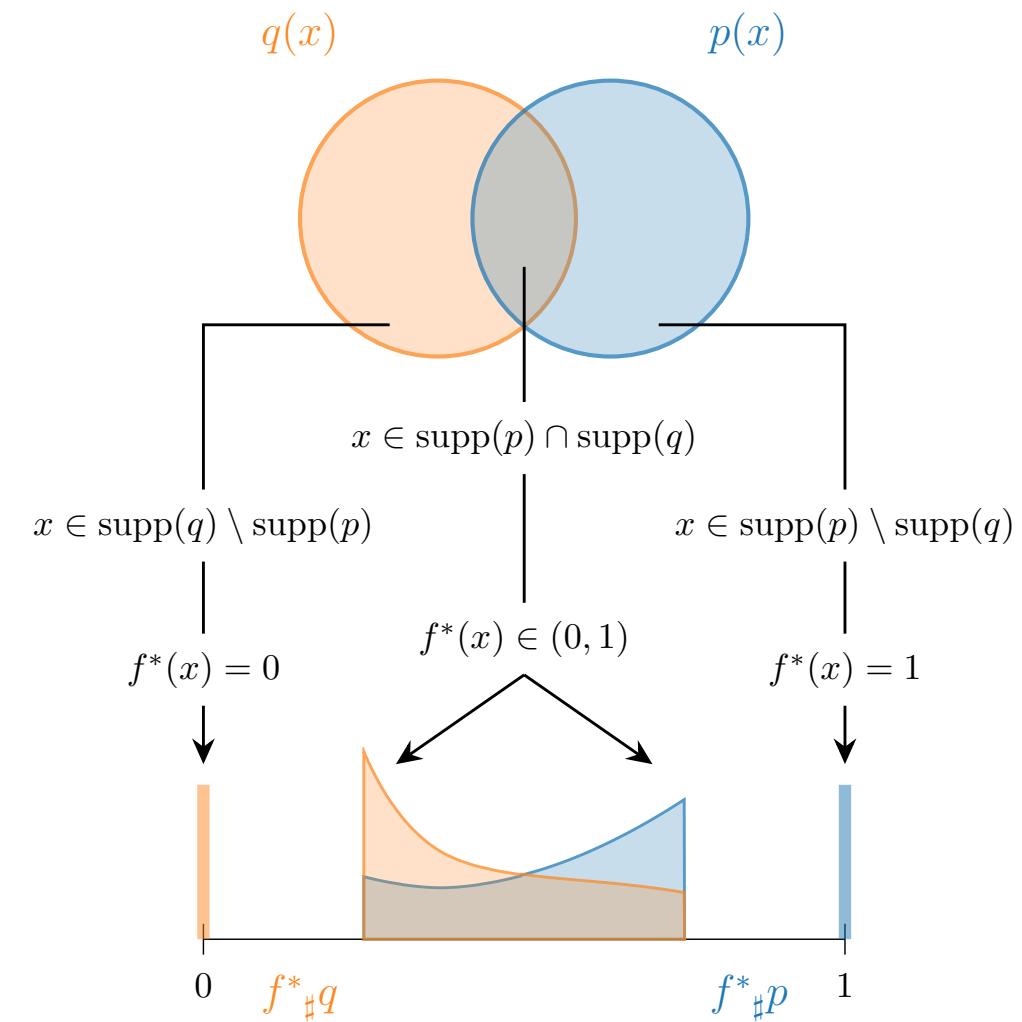
$$\sup_{f:\mathcal{X} \rightarrow [0,1]} \mathbb{E}_{x \sim p} [\log f(x)] + \mathbb{E}_{y \sim q} [\log(1 - f(y))]$$

$$f^*(x) = \frac{p(x)}{p(x) + q(x)}$$

## Theorem

The mapping  $f^* : \mathcal{X} \rightarrow [0, 1]$   
realized by the optimal discriminator  
preserves support discrepancy

$\text{supp}(p) = \text{supp}(q) \Leftrightarrow \text{supp}(f^*\#p) = \text{supp}(f^*\#q)$   
 $f^*\#p, f^*\#q$  — pushforward distributions



**Remark:** the result also holds for  $g : \mathcal{X} \rightarrow \mathbb{R}$   
 $g(x) : f(x) = \text{sigmoid}(g(x))$

**Not all discriminators classes have this property**  
**(details in the paper)**

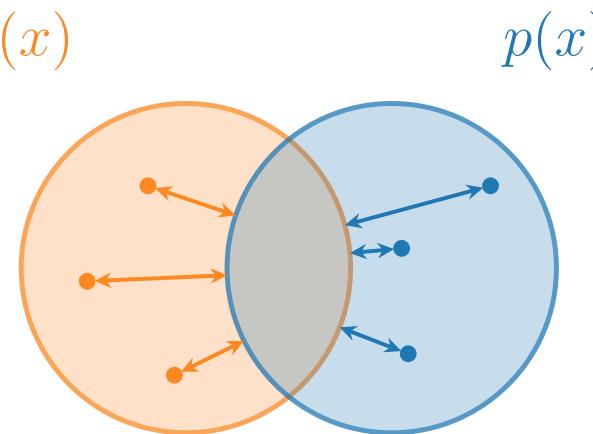
# Support Difference

## Symmetric Support Difference\* / Chamfer Distance

$$\mathcal{D}_\Delta(p, q) = \mathbb{E}_{x^q \sim q} [d(x^q, \text{supp}(p))] + \mathbb{E}_{x^p \sim p} [d(x^p, \text{supp}(q))]$$

$$d(x^q, \text{supp}(p)) = \inf_{x^p \in \text{supp}(p)} d(x^q, x^p)$$

$$d(x^p, \text{supp}(q)) = \inf_{x^q \in \text{supp}(q)} d(x^p, x^q)$$



- 1)  $\mathcal{D}_\Delta(p, q) \geq 0 \quad \forall p, q;$
- 2)  $\mathcal{D}_\Delta(p, q) = 0 \iff \text{supp}(p) = \text{supp}(q)$

\*generalizes Chamfer distance to continuous distributions

# Spectrum of Alignment Criteria

Wasserstein distance

$$\begin{aligned}\mathcal{D}_W(p, q) &= \inf_{\gamma} \mathbb{E}_{(x,y) \sim \gamma} [d(x, y)] \\ \text{s.t } &\int \gamma(x, y) dy = p(x) \\ &\int \gamma(x, y) dx = q(y)\end{aligned}$$

$$\mathcal{D}_W(p, q) = 0 \iff p = q$$

# Spectrum of Alignment Criteria

Wasserstein distance

$$\begin{aligned}\mathcal{D}_W(p, q) &= \inf_{\gamma} \mathbb{E}_{(x,y) \sim \gamma} [d(x, y)] \\ \text{s.t. } &\int \gamma(x, y) dy = p(x) \\ &\int \gamma(x, y) dx = q(y)\end{aligned}$$

$$\mathcal{D}_W(p, q) = 0 \iff p = q$$

$\beta$ -admissible Wasserstein distance ( $\beta > 0$ ) [Wu et al., 2019]

$$\begin{aligned}\mathcal{D}_W^\beta(p, q) &= \inf_{\gamma} \mathbb{E}_{(x,y) \sim \gamma} [d(x, y)] \\ \text{s.t. } &\int \gamma(x, y) dy = p(x) \\ &\int \gamma(x, y) dx \leq (1 + \beta)q(y)\end{aligned}$$

$$\mathcal{D}_W^\beta(p, q) = 0 \iff p(x) \leq (1 + \beta)q(x)$$

# Spectrum of Alignment Criteria

Wasserstein distance

$$\begin{aligned}\mathcal{D}_W(p, q) &= \inf_{\gamma} \mathbb{E}_{(x,y) \sim \gamma} [d(x, y)] \\ \text{s.t. } &\int \gamma(x, y) dy = p(x) \\ &\int \gamma(x, y) dx = q(y)\end{aligned}$$

$$\mathcal{D}_W(p, q) = 0 \iff p = q$$

$\beta$ -admissible Wasserstein distance ( $\beta > 0$ ) [Wu et al., 2019]

$$\begin{aligned}\mathcal{D}_W^\beta(p, q) &= \inf_{\gamma} \mathbb{E}_{(x,y) \sim \gamma} [d(x, y)] \\ \text{s.t. } &\int \gamma(x, y) dy = p(x) \\ &\int \gamma(x, y) dx \leq (1 + \beta)q(y)\end{aligned}$$

$$\mathcal{D}_W^\beta(p, q) = 0 \iff p(x) \leq (1 + \beta)q(x)$$

Support Difference (SD)

$$\begin{aligned}\mathcal{D}_W^\infty(p, q) &= \inf_{\gamma} \mathbb{E}_{(x,y) \sim \gamma} [d(x, y)] = \mathbb{E}_{x \sim p(x)} [\inf_{y \in \text{supp}(q)} d(x, y)] \\ \text{s.t. } &\int \gamma(x, y) dy = p(x) \\ q(y) = 0 &\implies \int \gamma(x, y) dx = 0\end{aligned}$$

$$\mathcal{D}_W^\infty(p, q) = 0 \iff \text{supp}(p) \subset \text{supp}(q)$$

# Spectrum of Alignment Criteria

Wasserstein distance

$$\begin{aligned}\mathcal{D}_W(p, q) &= \inf_{\gamma} \mathbb{E}_{(x,y) \sim \gamma} [d(x, y)] \\ \text{s.t. } &\int \gamma(x, y) dy = p(x) \\ &\int \gamma(x, y) dx = q(y)\end{aligned}$$

$$\mathcal{D}_W(p, q) = 0 \iff p = q$$

$\beta$ -admissible Wasserstein distance ( $\beta > 0$ ) [Wu et al., 2019]

$$\begin{aligned}\mathcal{D}_W^\beta(p, q) &= \inf_{\gamma} \mathbb{E}_{(x,y) \sim \gamma} [d(x, y)] \\ \text{s.t. } &\int \gamma(x, y) dy = p(x) \\ &\int \gamma(x, y) dx \leq (1 + \beta)q(y)\end{aligned}$$

$$\mathcal{D}_W^\beta(p, q) = 0 \iff p(x) \leq (1 + \beta)q(x)$$

Support Difference (SD)

$$\begin{aligned}\mathcal{D}_W^\infty(p, q) &= \inf_{\gamma} \mathbb{E}_{(x,y) \sim \gamma} [d(x, y)] = \mathbb{E}_{x \sim p(x)} [\inf_{y \in \text{supp}(q)} d(x, y)] \\ \text{s.t. } &\int \gamma(x, y) dy = p(x) \\ q(y) = 0 &\implies \int \gamma(x, y) dx = 0\end{aligned}$$

$$\mathcal{D}_W^\infty(p, q) = 0 \iff \text{supp}(p) \subset \text{supp}(q)$$

$$\begin{aligned}\mathcal{D}_\Delta(p, q) &= \mathcal{D}_W^\infty(p, q) + \mathcal{D}_W^\infty(q, p) = \lim_{\beta \rightarrow \infty} \mathcal{D}_W^\beta(p, q) + \mathcal{D}_W^\beta(q, p) \\ \mathcal{D}_\Delta(p, q) = 0 &\iff \text{supp}(p) = \text{supp}(q)\end{aligned}$$

# Spectrum of Alignment Criteria

**Proposition** (Alignment Conditions are Preserved by Log-loss Discriminator).

Let  $f^*$  be the optimal discriminator  $f^*(x) = \frac{p(x)}{p(x)+q(x)}$  for distributions  $p, q$ .

Let  $\mathcal{D}_W^{\beta_1, \beta_2}(p, q) = \mathcal{D}_W^{\beta_1}(p, q) + \mathcal{D}_W^{\beta_2}(q, p)$ . Then,

1.  $\mathcal{D}_W(p, q) = 0$  if and only if  $\mathcal{D}_W(f^*\#p, f^*\#q) = 0$ ; (distribution alignment)
2.  $\mathcal{D}_W^{\beta_1, \beta_2}(p, q) = 0$  if and only if  $\mathcal{D}_W^{\beta_1, \beta_2}(f^*\#p, f^*\#q) = 0$ ; (relaxed distribution alignment)
3.  $\mathcal{D}_\Delta(p, q) = 0$  if and only if  $\mathcal{D}_\Delta(f^*\#p, f^*\#q) = 0$ . (support alignment)

# Adversarial Support Alignment

$$\mathcal{L}_D(\theta, g) = \mathbb{E}_{x \sim p^\theta} \left[ \log \left( 1 + e^{-g(x)} \right) \right] + \mathbb{E}_{x \sim q^\theta} \left[ \log \left( 1 + e^{g(x)} \right) \right]$$

# Adversarial Support Alignment

$$\mathcal{L}_D(\theta, g) = \mathbb{E}_{x \sim p^\theta} \left[ \log \left( 1 + e^{-g(x)} \right) \right] + \mathbb{E}_{x \sim q^\theta} \left[ \log \left( 1 + e^{g(x)} \right) \right]$$

Alignment objective:  $\min_{\theta} \mathcal{L}_A(\theta, g)$

# Adversarial Support Alignment

$$\mathcal{L}_D(\theta, g) = \mathbb{E}_{x \sim p^\theta} \left[ \log \left( 1 + e^{-g(x)} \right) \right] + \mathbb{E}_{x \sim q^\theta} \left[ \log \left( 1 + e^{g(x)} \right) \right]$$

Alignment objective:  $\min_{\theta} \mathcal{L}_A(\theta, g)$

## Adversarial Distribution Alignment (GAN/DANN)

$$\mathcal{L}_A(\theta, g) = \mathbb{E}_{x \sim p^\theta} \left[ \log \left( 1 + e^{g(x)} \right) \right] + \mathbb{E}_{x \sim q^\theta} \left[ \log \left( 1 + e^{-g(x)} \right) \right]$$

# Adversarial Support Alignment

$$\mathcal{L}_D(\theta, g) = \mathbb{E}_{x \sim p^\theta} \left[ \log \left( 1 + e^{-g(x)} \right) \right] + \mathbb{E}_{x \sim q^\theta} \left[ \log \left( 1 + e^{g(x)} \right) \right]$$

Alignment objective:  $\min_{\theta} \mathcal{L}_A(\theta, g)$

**Adversarial Distribution Alignment  
(GAN/DANN)**

$$\mathcal{L}_A(\theta, g) = \mathbb{E}_{x \sim p^\theta} \left[ \log \left( 1 + e^{g(x)} \right) \right] + \mathbb{E}_{x \sim q^\theta} \left[ \log \left( 1 + e^{-g(x)} \right) \right]$$

**Adversarial Support Alignment  
(ASA)**

$$\mathcal{L}_A(\theta, g) = \mathbb{E}_{x \sim p^\theta} \left[ d(g(x), \text{supp}(g \sharp q)) \right] + \mathbb{E}_{x \sim q^\theta} \left[ d(g(x), \text{supp}(g \sharp p)) \right]$$

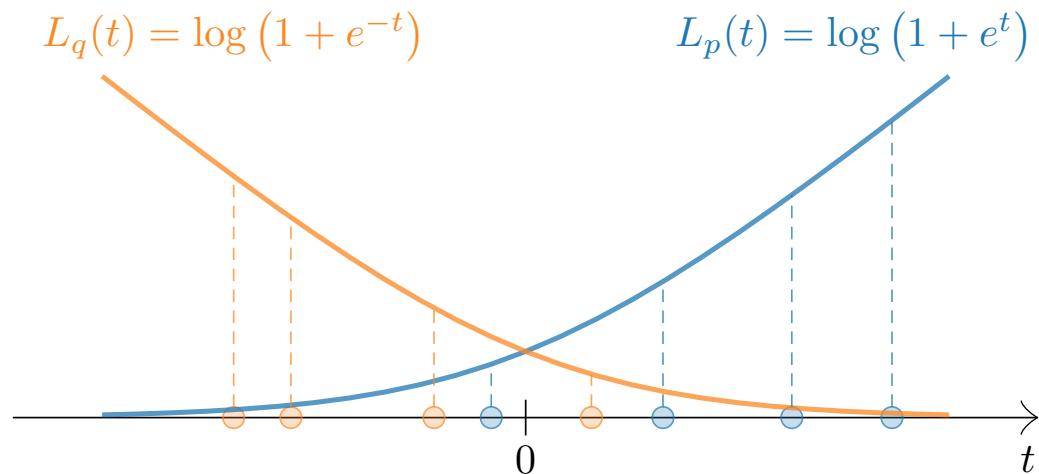
# Adversarial Support Alignment

$$\mathcal{L}_D(\theta, g) = \mathbb{E}_{x \sim p^\theta} \left[ \log \left( 1 + e^{-g(x)} \right) \right] + \mathbb{E}_{x \sim q^\theta} \left[ \log \left( 1 + e^{g(x)} \right) \right]$$

Alignment objective:  $\min_{\theta} \mathcal{L}_A(\theta, g)$

**Adversarial Distribution Alignment  
(GAN/DANN)**

$$\mathcal{L}_A(\theta, g) = \mathbb{E}_{x \sim p^\theta} \left[ \log \left( 1 + e^{g(x)} \right) \right] + \mathbb{E}_{x \sim q^\theta} \left[ \log \left( 1 + e^{-g(x)} \right) \right]$$



**Adversarial Support Alignment  
(ASA)**

$$\mathcal{L}_A(\theta, g) = \mathbb{E}_{x \sim p^\theta} \left[ d(g(x), \text{supp}(g \sharp q)) \right] + \mathbb{E}_{x \sim q^\theta} \left[ d(g(x), \text{supp}(g \sharp p)) \right]$$

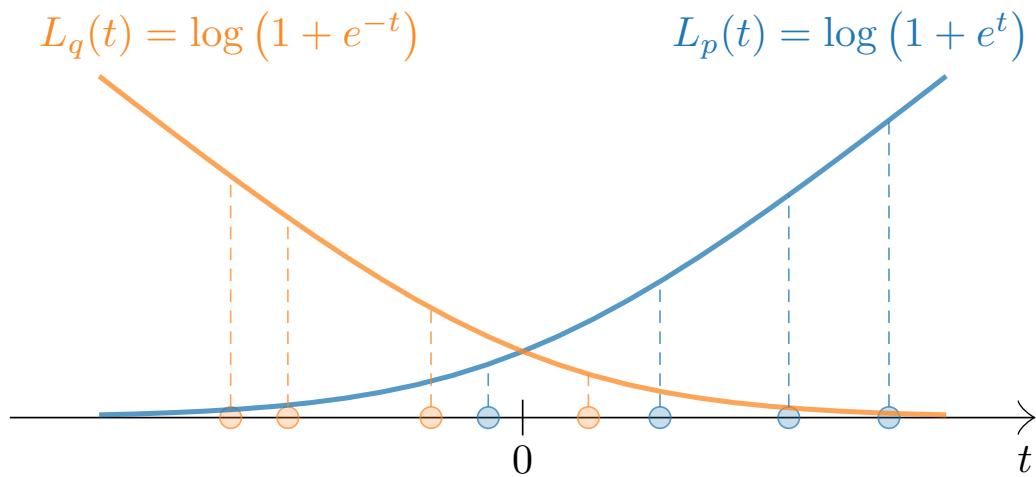
# Adversarial Support Alignment

$$\mathcal{L}_D(\theta, g) = \mathbb{E}_{x \sim p^\theta} \left[ \log \left( 1 + e^{-g(x)} \right) \right] + \mathbb{E}_{x \sim q^\theta} \left[ \log \left( 1 + e^{g(x)} \right) \right]$$

Alignment objective:  $\min_{\theta} \mathcal{L}_A(\theta, g)$

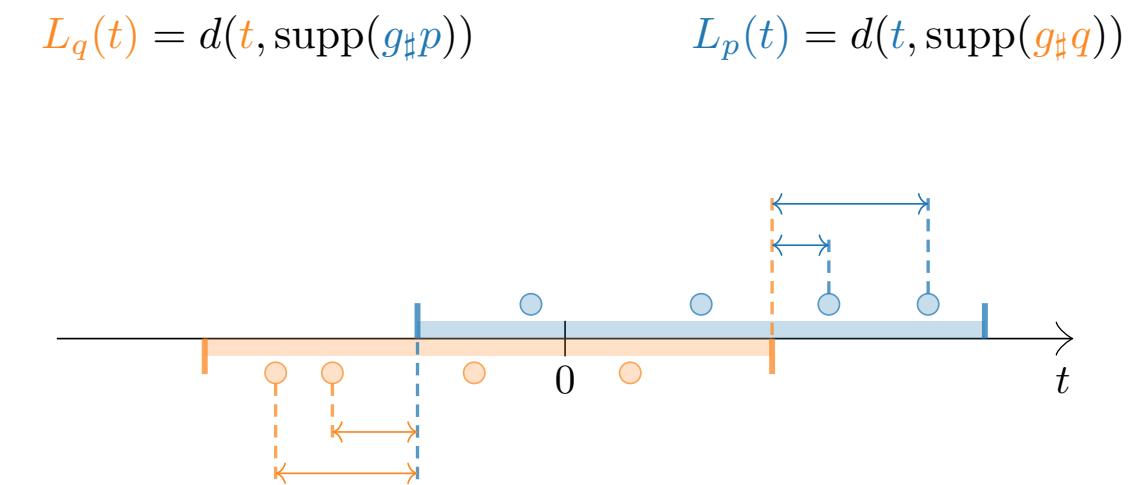
## Adversarial Distribution Alignment (GAN/DANN)

$$\mathcal{L}_A(\theta, g) = \mathbb{E}_{x \sim p^\theta} \left[ \log \left( 1 + e^{g(x)} \right) \right] + \mathbb{E}_{x \sim q^\theta} \left[ \log \left( 1 + e^{-g(x)} \right) \right]$$



## Adversarial Support Alignment (ASA)

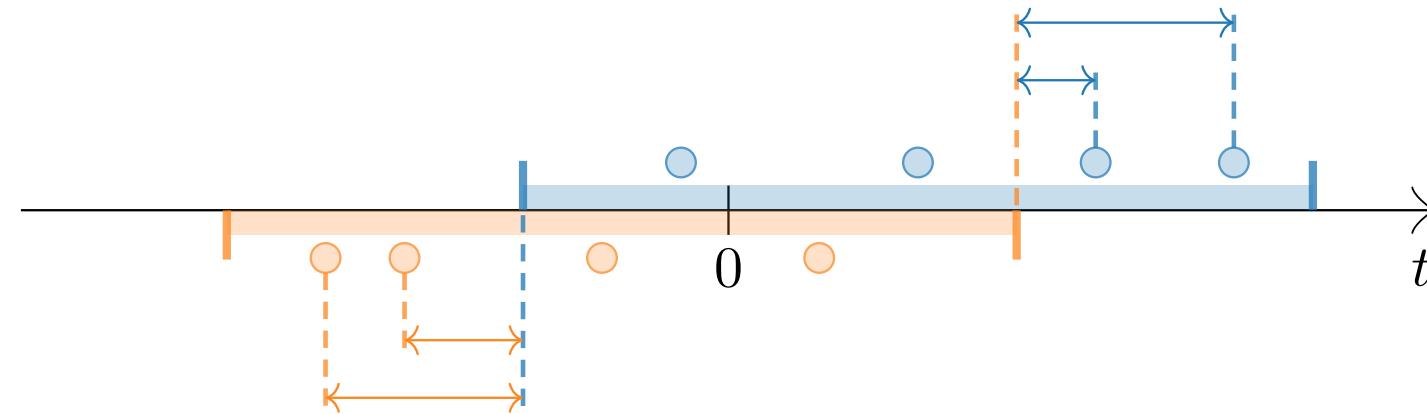
$$\mathcal{L}_A(\theta, g) = \mathbb{E}_{x \sim p^\theta} \left[ d(g(x), \text{supp}(g \sharp q)) \right] + \mathbb{E}_{x \sim q^\theta} \left[ d(g(x), \text{supp}(g \sharp p)) \right]$$



$$\mathcal{L}_A(\theta, g) = \mathbb{E}_{x \sim p^\theta} \left[ d(g(x), \text{supp}(g_\sharp q)) \right] + \mathbb{E}_{x \sim q^\theta} \left[ d(g(x), \text{supp}(g_\sharp p)) \right]$$

$$L_q(t) = d(\textcolor{brown}{t}, \text{supp}(\textcolor{blue}{g}_\sharp p))$$

$$L_p(t) = d(\textcolor{blue}{t}, \text{supp}(\textcolor{brown}{g}_\sharp q))$$

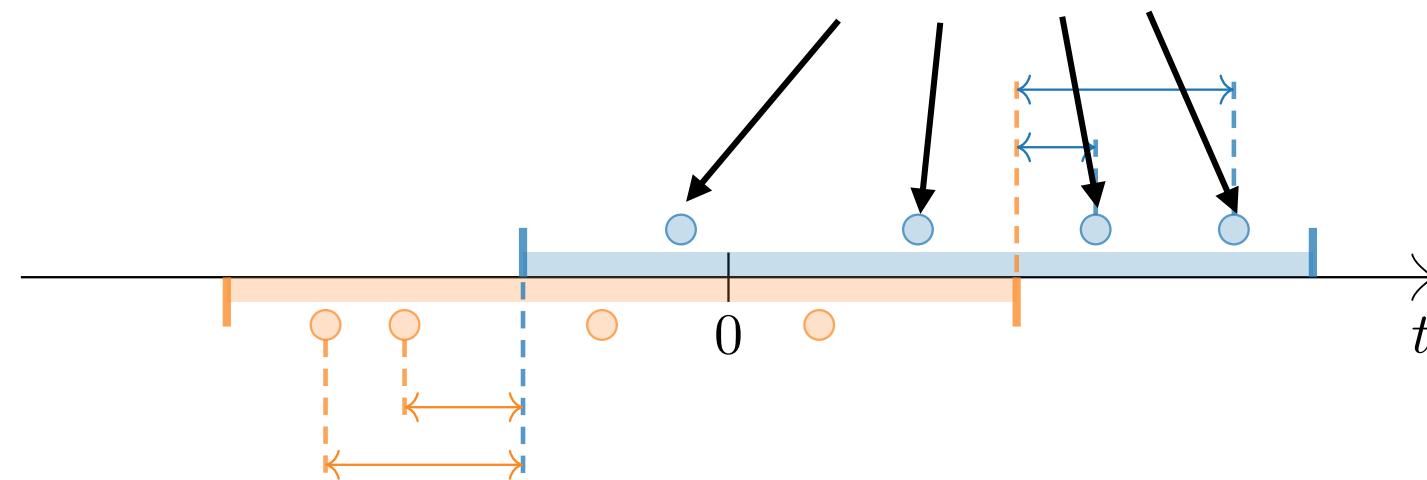


$$\mathcal{L}_A(\theta, g) = \mathbb{E}_{x \sim p^\theta} \left[ d(g(x), \text{supp}(g_\sharp q)) \right] + \mathbb{E}_{x \sim q^\theta} \left[ d(g(x), \text{supp}(g_\sharp p)) \right]$$

$$L_q(t) = d(\textcolor{brown}{t}, \text{supp}(\textcolor{blue}{g}_\sharp p))$$

$$L_p(t) = d(\textcolor{blue}{t}, \text{supp}(\textcolor{brown}{g}_\sharp q))$$

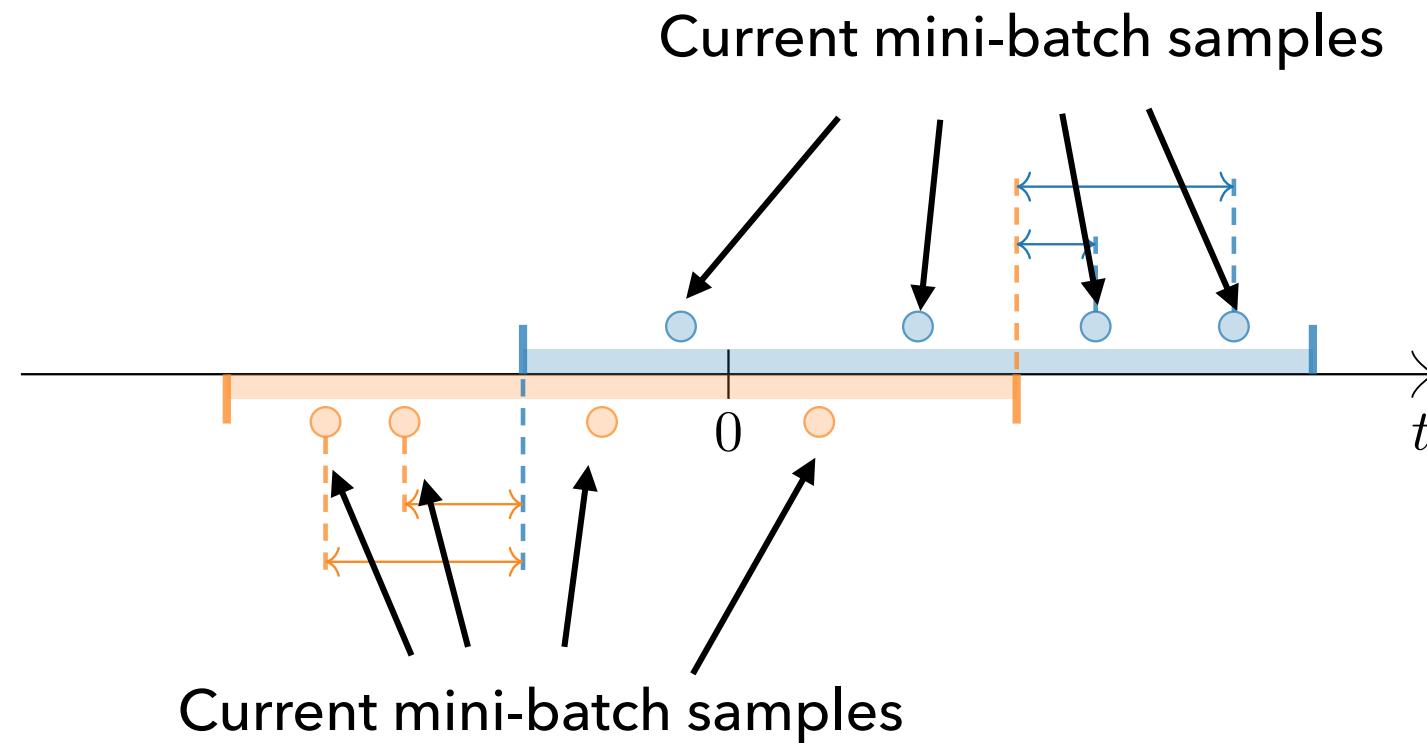
Current mini-batch samples



$$\mathcal{L}_A(\theta, g) = \mathbb{E}_{x \sim p^\theta} \left[ d(g(x), \text{supp}(g_\sharp q)) \right] + \mathbb{E}_{x \sim q^\theta} \left[ d(g(x), \text{supp}(g_\sharp p)) \right]$$

$$L_q(t) = d(\textcolor{brown}{t}, \text{supp}(\textcolor{blue}{g}_\sharp p))$$

$$L_p(t) = d(\textcolor{teal}{t}, \text{supp}(\textcolor{brown}{g}_\sharp q))$$



$$\mathcal{L}_A(\theta, g) = \mathbb{E}_{x \sim p^\theta} \left[ d(g(x), \text{supp}(g_\sharp q)) \right] + \mathbb{E}_{x \sim q^\theta} \left[ d(g(x), \text{supp}(g_\sharp p)) \right]$$

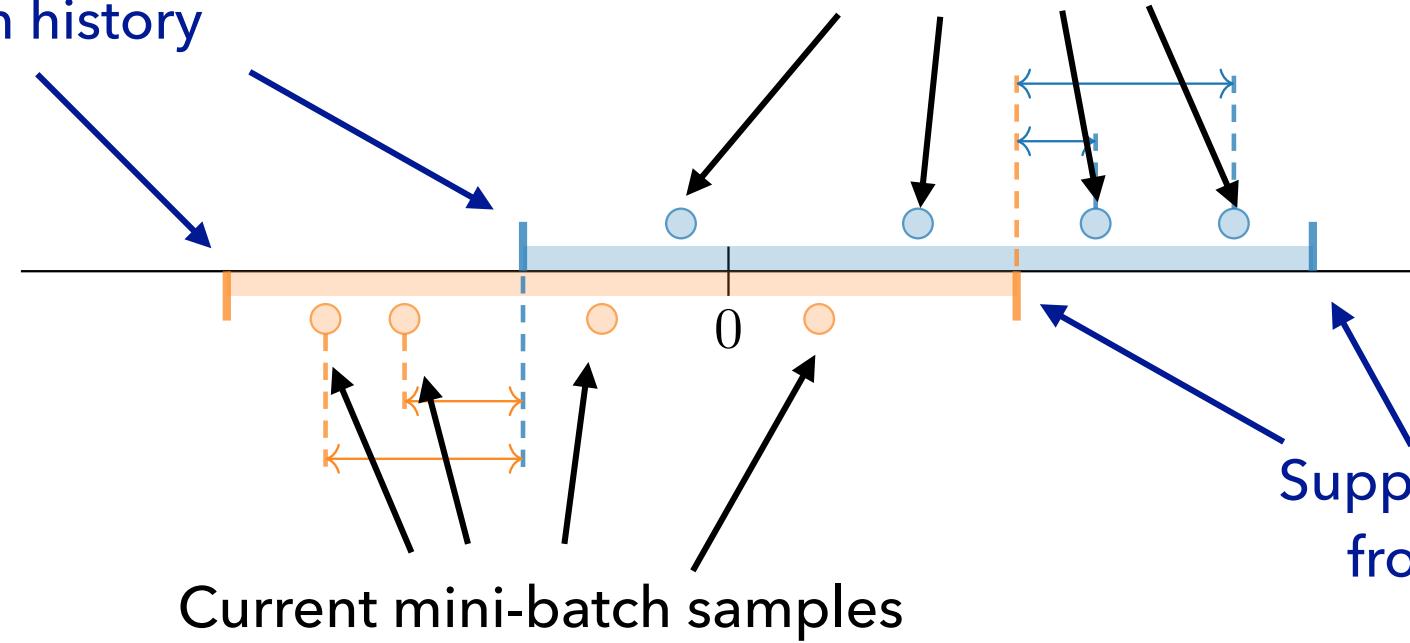
$$L_q(t) = d(\textcolor{brown}{t}, \text{supp}(\textcolor{blue}{g}_\sharp p))$$

$$L_p(t) = d(\textcolor{blue}{t}, \text{supp}(\textcolor{brown}{g}_\sharp q))$$

Support bounds are estimated  
from rolling batch history

Current mini-batch samples

Support bounds are estimated  
from rolling batch history



# Results: 2D Embeddings Visualization

Toy problem: USPS→MNIST, 3 classes, 2D Embeddings

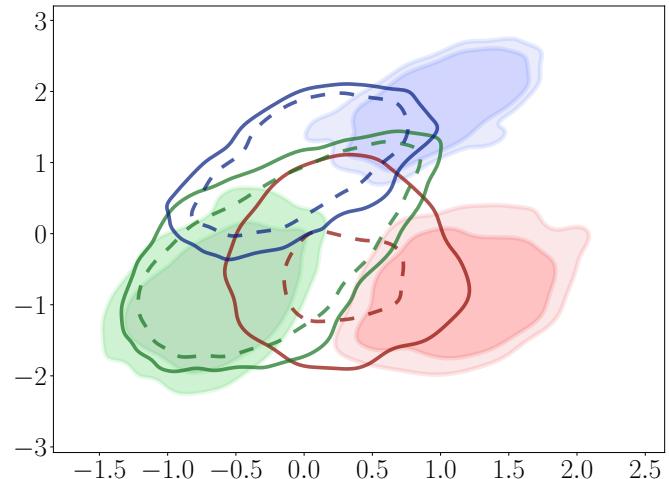
**Source class distribution**  
[33%, 33%, 33%]

**Target class distribution**  
[23%, 65%, 12%]

# Results: 2D Embeddings Visualization

Toy problem: USPS $\rightarrow$ MNIST, 3 classes, 2D Embeddings

**Source class distribution**  
[33%, 33%, 33%]



(a) No DA (avg acc: 63%)

$$\mathcal{D}_W(p_Z^\theta, q_Z^\theta) = 0.78$$

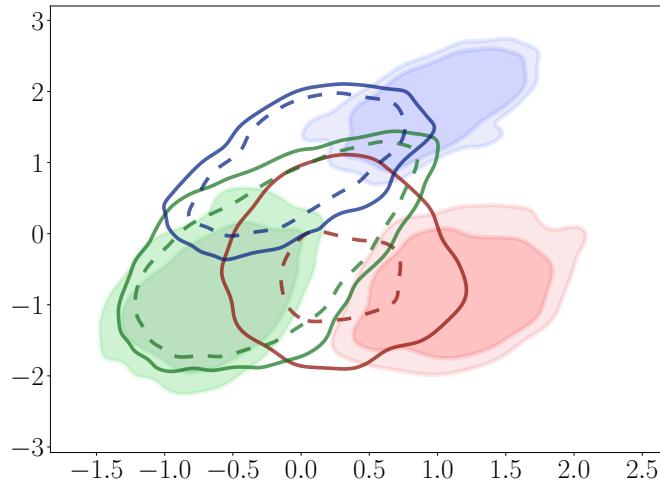
$$\mathcal{D}_\Delta(p_Z^\theta, q_Z^\theta) = 0.10$$

**Target class distribution**  
[23%, 65%, 12%]

# Results: 2D Embeddings Visualization

Toy problem: USPS→MNIST, 3 classes, 2D Embeddings

**Source class distribution**  
[33%, 33%, 33%]

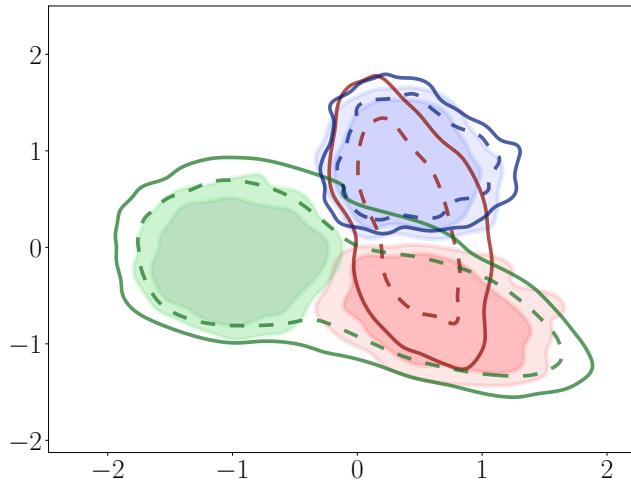


(a) No DA (avg acc: 63%)

$$\mathcal{D}_W(p_Z^\theta, q_Z^\theta) = 0.78$$

$$\mathcal{D}_\Delta(p_Z^\theta, q_Z^\theta) = 0.10$$

**Target class distribution**  
[23%, 65%, 12%]



(b) DANN (avg acc: 75%)

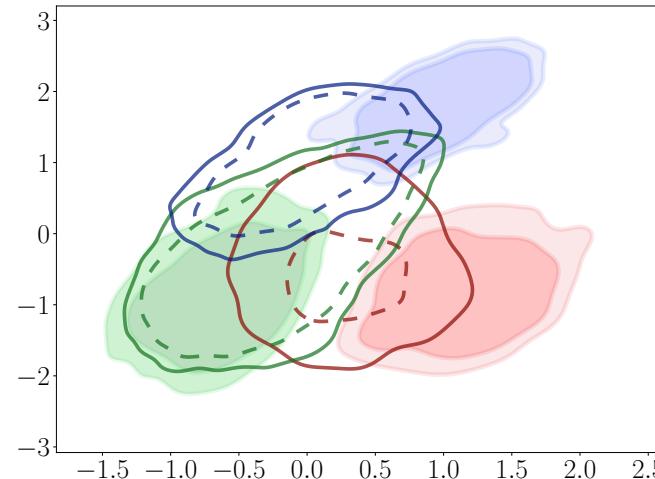
$$\mathcal{D}_W(p_Z^\theta, q_Z^\theta) = 0.07$$

$$\mathcal{D}_\Delta(p_Z^\theta, q_Z^\theta) = 0.02$$

# Results: 2D Embeddings Visualization

Toy problem: USPS→MNIST, 3 classes, 2D Embeddings

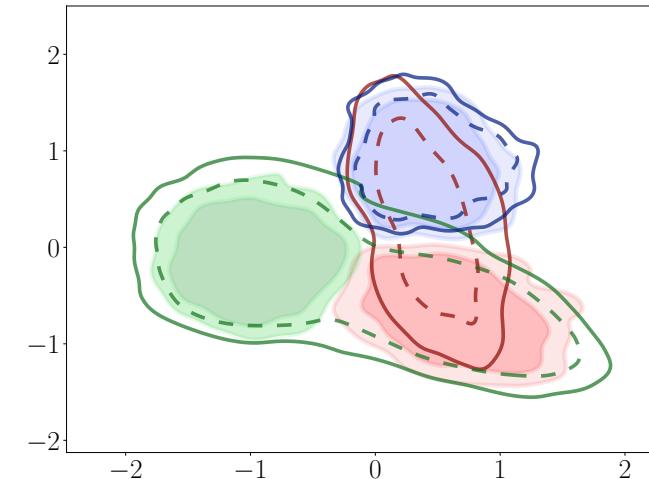
**Source class distribution**  
[33%, 33%, 33%]



(a) No DA (avg acc: 63%)

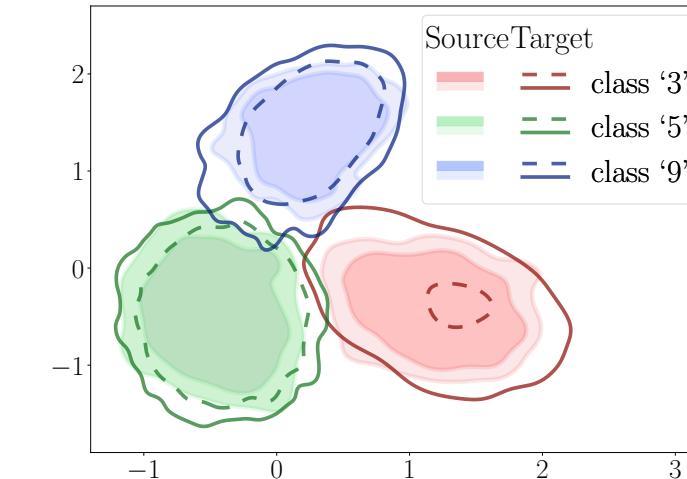
$$\mathcal{D}_W(p_Z^\theta, q_Z^\theta) = 0.78$$
$$\mathcal{D}_\Delta(p_Z^\theta, q_Z^\theta) = 0.10$$

**Target class distribution**  
[23%, 65%, 12%]



(b) DANN (avg acc: 75%)

$$\mathcal{D}_W(p_Z^\theta, q_Z^\theta) = 0.07$$
$$\mathcal{D}_\Delta(p_Z^\theta, q_Z^\theta) = 0.02$$



(c) ASA-abs (avg acc: 94%)

$$\mathcal{D}_W(p_Z^\theta, q_Z^\theta) = 0.59$$
$$\mathcal{D}_\Delta(p_Z^\theta, q_Z^\theta) = 0.03$$

# Results: USPS → MNIST, SmallCNN

Average and minimum class accuracy (%) on USPS→MNIST  
across different levels of shifts in label distributions ( $\alpha$ ).

Algorithm	$\alpha = 0.0$ no shift		$\alpha = 1.0$		$\alpha = 1.5$		$\alpha = 2.0$ severe shift	
	average	min	average	min	average	min	average	min
No DA	71.9	20.3	72.9	25.8	71.3	27.5	71.3	16.6
DANN	97.8	96.0	83.5	25.1	70.0	01.1	57.8	00.9
VADA	<b>98.0</b>	96.2	88.2	48.9	78.2	06.6	61.9	01.4
IWDAN	97.5	95.7	95.7	81.3	86.5	15.2	74.4	07.3
IWCDAN	<b>98.0</b>	<b>96.6</b>	<b>96.7</b>	85.1	91.3	66.5	77.5	22.2
sDANN-4	87.4	05.6	94.9	<b>85.7</b>	86.8	21.6	81.5	39.3
ASA-sq	93.7	89.2	92.3	83.5	90.9	69.9	87.2	62.5
ASA-abs	94.1	88.9	92.8	78.9	<b>92.5</b>	<b>82.4</b>	<b>90.4</b>	<b>68.4</b>

Distribution Alignment  
Relaxed Distribution Alignment  
Support Alignment (ours)

~0% worst-class accuracy under severe shift

Our method is robust to distribution shifts

# Results: Larger Datasets

STL10 → CIFAR10  
DeepCNN

Algorithm	$\alpha = 0.0$ no shift		$\alpha = 2.0$ severe shift	
	average	min	average	min
No DA	69.9	49.8	65.8	43.7
DANN	75.3	54.6	63.3	27.0
VADA	<b>76.7</b>	<b>56.9</b>	63.2	25.5
IWDAN	69.9	50.5	64.4	36.8
IWCDAN	70.1	47.8	64.5	37.0
sDANN-4	71.8	52.1	66.4	39.0
ASA-sq	71.7	52.9	<b>68.1</b>	<b>44.7</b>
ASA-abs	71.6	49.0	67.8	40.9

VisDA-17  
ResNet50

Algorithm	$\alpha = 0.0$ no shift		$\alpha = 2.0$ severe shift	
	average	min	average	min
No DA	49.5	22.2	45.3	19.5
DANN	<b>75.4</b>	36.7	43.1	03.6
VADA	75.3	40.5	43.9	08.5
IWDAN	73.2	31.7	45.1	04.6
IWCDAN	71.6	27.6	38.3	00.6
sDANN-4	72.4	37.8	50.7	18.6
ASA-sq	64.9	35.7	51.9	18.3
ASA-abs	64.8	<b>40.6</b>	<b>52.5</b>	<b>19.7</b>

# Adversarial Support Alignment: Summary

**Support alignment: novel training criterion**, an alternative to distribution alignment

- ▶ **Support divergence** defined for continuous distributions
  - ▶ **Spectrum of alignment criteria** within **relaxed OT** framework
  - ▶ Distribution alignment / relaxed distribution alignment / support alignment
- ▶ Analysis of support discrepancy in the discriminator output space
  - ▶ Log-loss discriminator **preserves support discrepancy**
  - ▶ Not all discriminators have this property (e.g. linear discriminators, Wasserstein discriminators)
- ▶ Practical method: **log-loss discriminator + support difference + history buffers**

Experimental validation: unsupervised domain adaptation under label distribution shift

- ▶ Image classification UDA benchmarks, USPS-MNIST, CIFAR-STL, VisDA17
- ▶ Robust performance in the face of label distribution shift, **improved worst class accuracy**

# Chapter IV

## Compositional Sculpting of Iterative Generative Processes

---

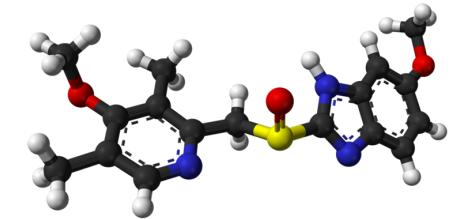
Compositional Sculpting of Iterative Generative Processes

**T. Garipov\***, S. De Peuter, G. Yang, V Targ, S. Kaski, T. Jaakkola (NeurIPS 2023)

# Composition of Generative Models

Large-scale general-purpose pre-training is becoming ubiquitous

- ▶ Need to re-use and adapt **pre-trained models** for **new tasks**

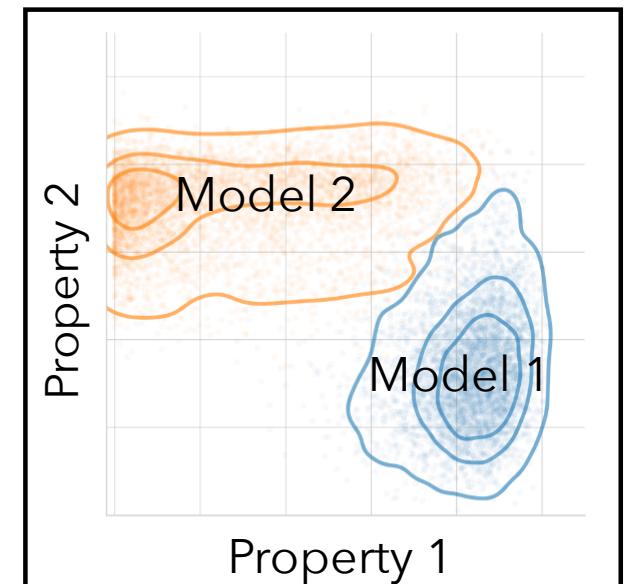


Applications: multi-objective generation (e.g. drug-like molecules)

- ▶ Need to combine knowledge from multiple sources (models, datasets)
- ▶ Need to explore trade-offs between multiple criteria

Composition is a powerful modeling tool

- ▶ Increases model capacity
- ▶ Enables control of sampling distributions



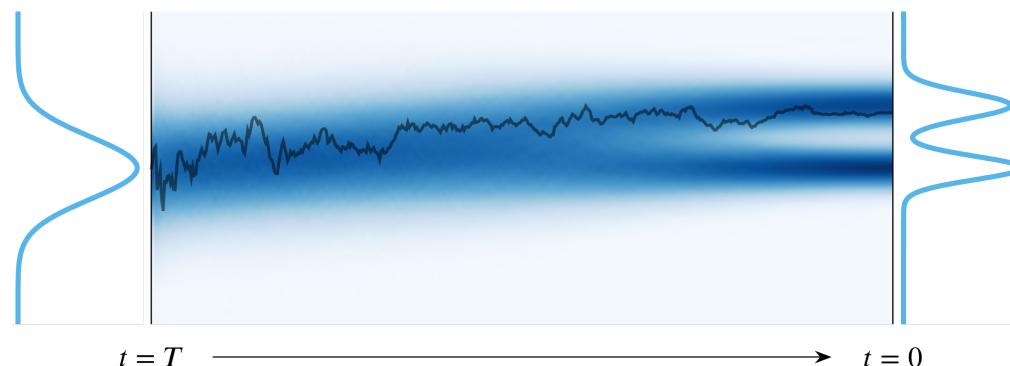
(Property 1) AND (Property 2)

# Iterative Generative Processes

**Diffusion model** generates trajectories  $\tau = (x_T \rightarrow \{x_t\}_{t=0}^T \rightarrow x_0)$  with terminal state distribution  $p(x_0)$  by following a backward SDE

$$dx_t = [f_t(x_t) - g_t^2 \underbrace{s_t(x_t; \theta)}_{\nabla_{x_t} \log p_t(x_t)}] dt + g_t d\bar{w}_t,$$

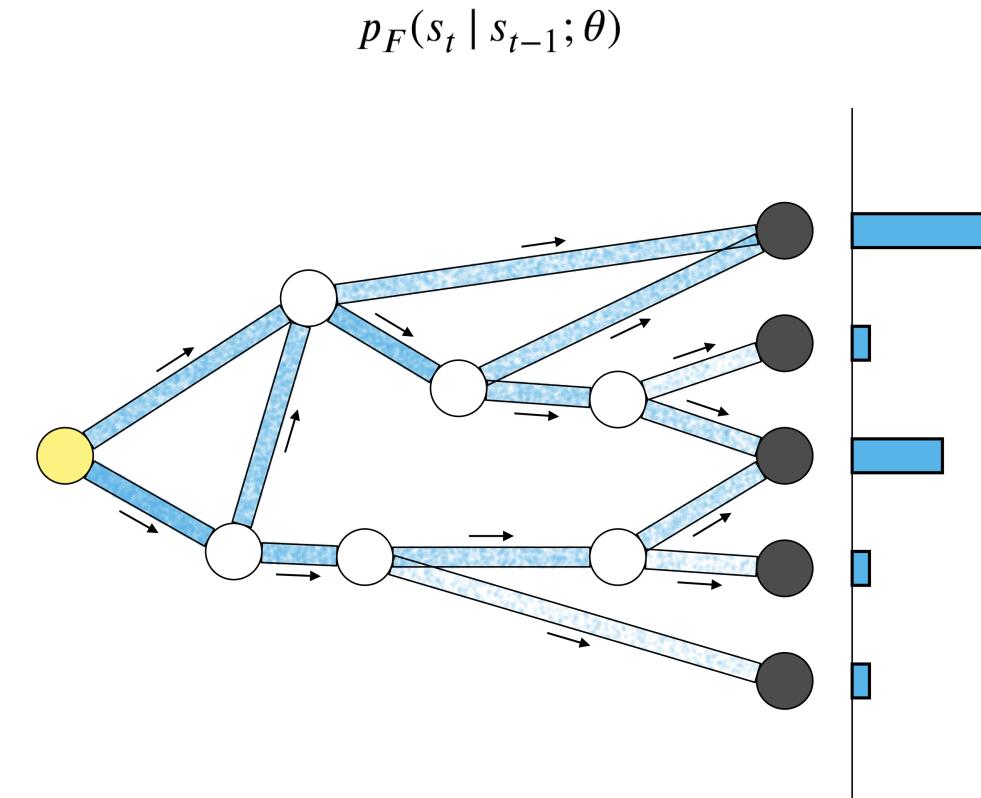
corresponding to a forward noising process  $dx_t = f_t(x_t) dt + g_t dw_t$



"Score-Based Generative Modeling through Stochastic Differential Equations"

[Song et al, ICLR 2021]

**GFlowNet** generates trajectories  $\tau = (s_0 \rightarrow \dots \rightarrow s_{n-1} \rightarrow x)$  with terminal distribution  $p(x) = \frac{R(x)}{Z}$  by following a forward policy



"Flow Network based Generative Models for Non-Iterative Diverse Candidate Generation"

[Bengio et al, NeurIPS 2021]

	<b>Models</b>	<b>Composition Operations</b>	<b>Sampling Algorithm</b>
[Hinton, Neural Computation 2002] [Du et al, NeurIPS 2020]	Energy-based models (EBMs) $p_i(x) \propto \exp(-E_i(x; \theta))$	Principle: energy-function arithmetic Product: $\frac{1}{Z} p_1(x) p_2(x)$ Negation: $\frac{1}{Z} \frac{p_1(x)}{(p_2(x))^\gamma}$	MCMC Langevin dynamics

Base models:  $p_1(x) = \frac{1}{Z_1} \exp \{ -E_1(x) \}, \quad p_2(x) = \frac{1}{Z_2} \exp \{ -E_2(x) \}$

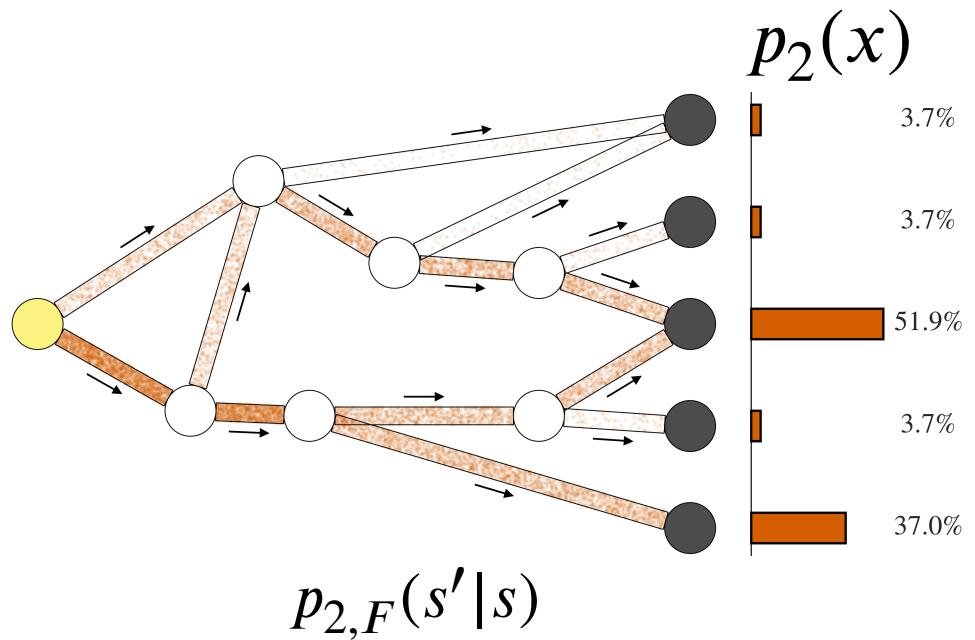
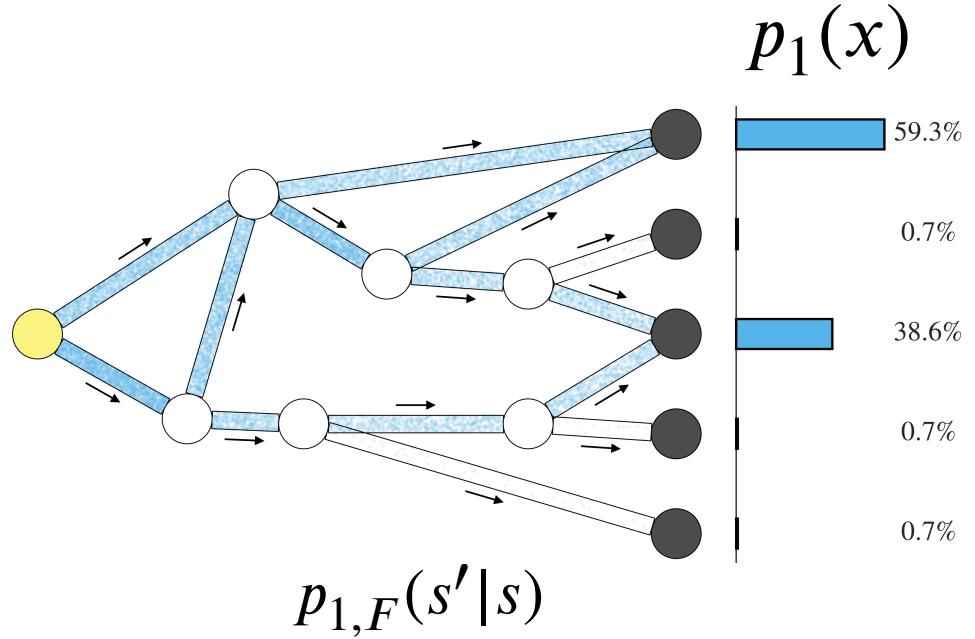
**Product:**  $p_{\text{prod}}(x) \propto p_1(x)p_2(x) \propto \exp \{ - (E_1(x) + E_2(x)) \}$

**Negation:**  $p_{\text{neg}}(x) \propto \frac{p_1(x)}{(p_2(x))^\gamma} \propto \exp \{ - (E_1(x) - \gamma E_2(x)) \}$

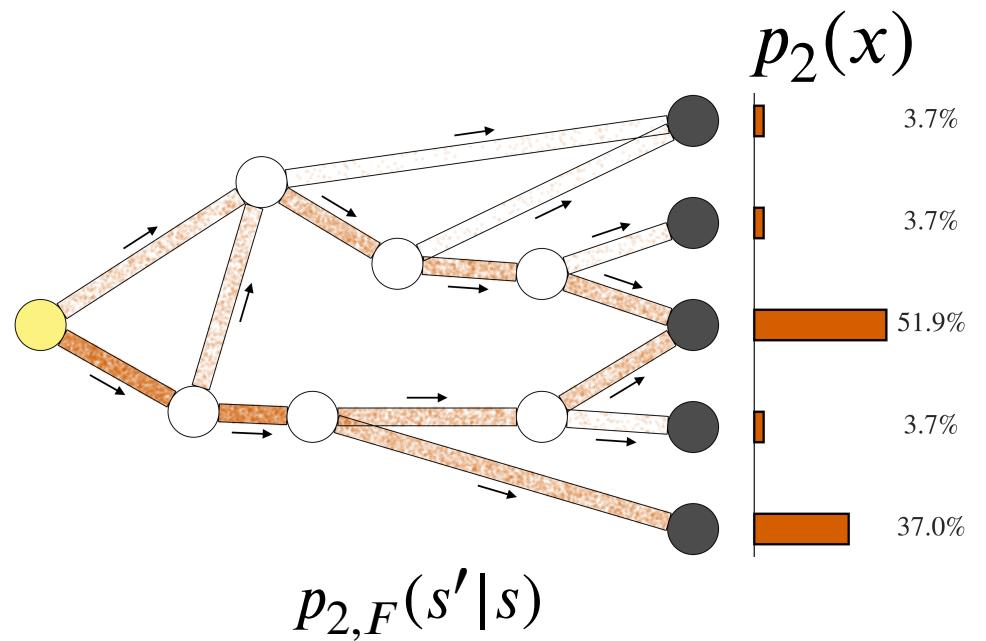
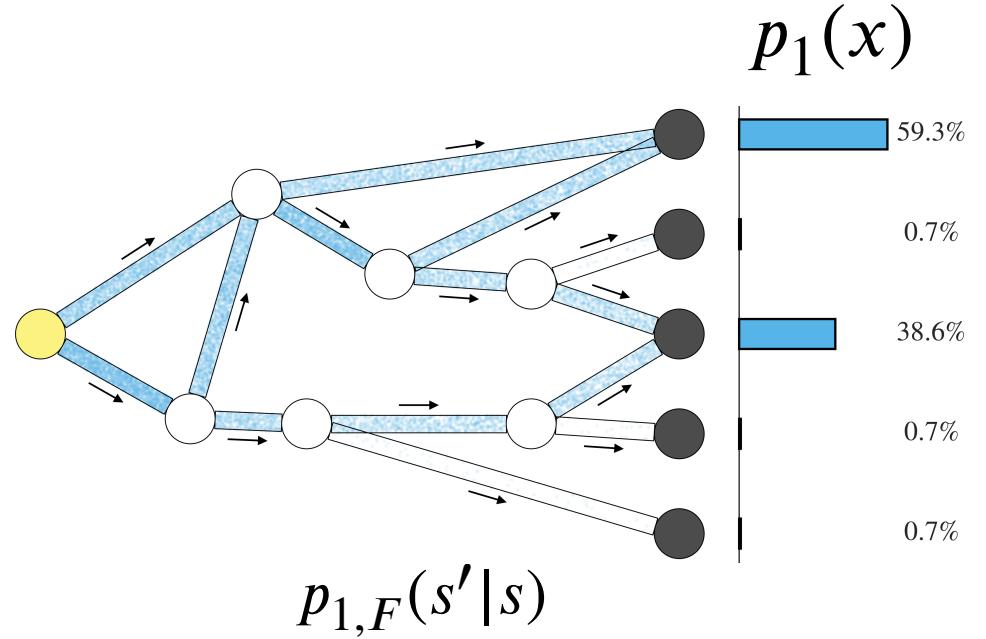
	<b>Models</b>	<b>Composition Operations</b>	<b>Sampling Algorithm</b>
<a href="#">[Hinton, Neural Computation 2002]</a> <a href="#">[Du et al, NeurIPS 2020]</a>	Energy-based models (EBMs) $p_i(x) \propto \exp(-E_i(x; \theta))$	Principle: energy-function arithmetic Product: $\frac{1}{Z} p_1(x) p_2(x)$ Negation: $\frac{1}{Z} \frac{p_1(x)}{(p_2(x))^\gamma}$	MCMC Langevin dynamics
<a href="#">[Liu et al, ECCV 2022]</a> <a href="#">[Du et al, ICML 2023]</a>	Diffusion models $p_i(x): s_{i,t}(x_t; \theta) \approx \nabla_{x_t} \log p_{i,t}(x_t)$	Principle: score-function arithmetic Product: $\frac{1}{Z} p_1(x) p_2(x)$ Negation: $\frac{1}{Z} \frac{p_1(x)}{(p_2(x))^\gamma}$	Diffusion sampling + annealed MCMC

**Challenge:** iterative generative processes (Diffusion models & GFlowNets) impose delicate balance conditions

# Composition Tools: Mixtures



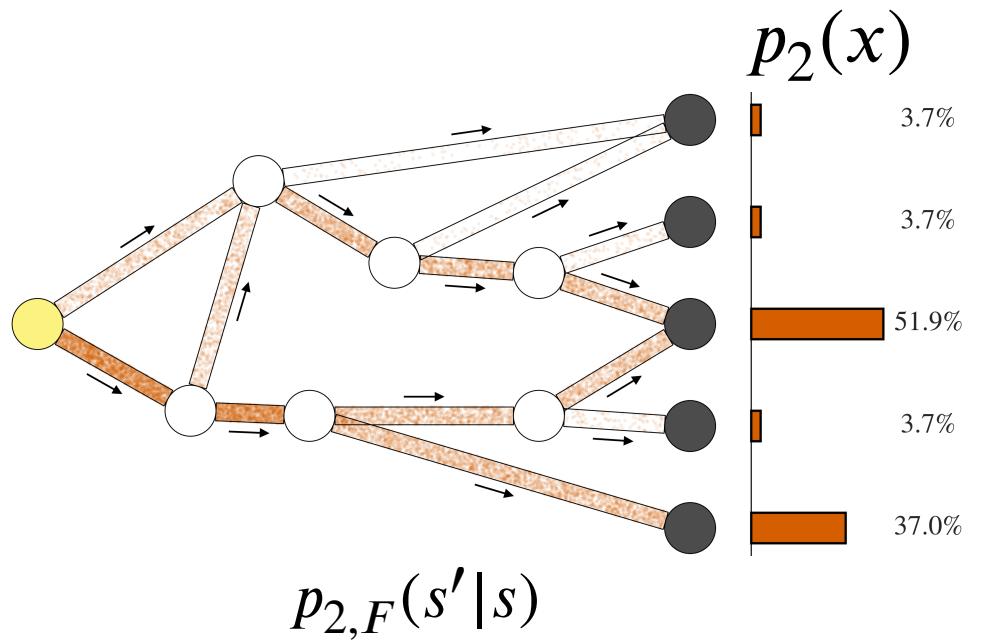
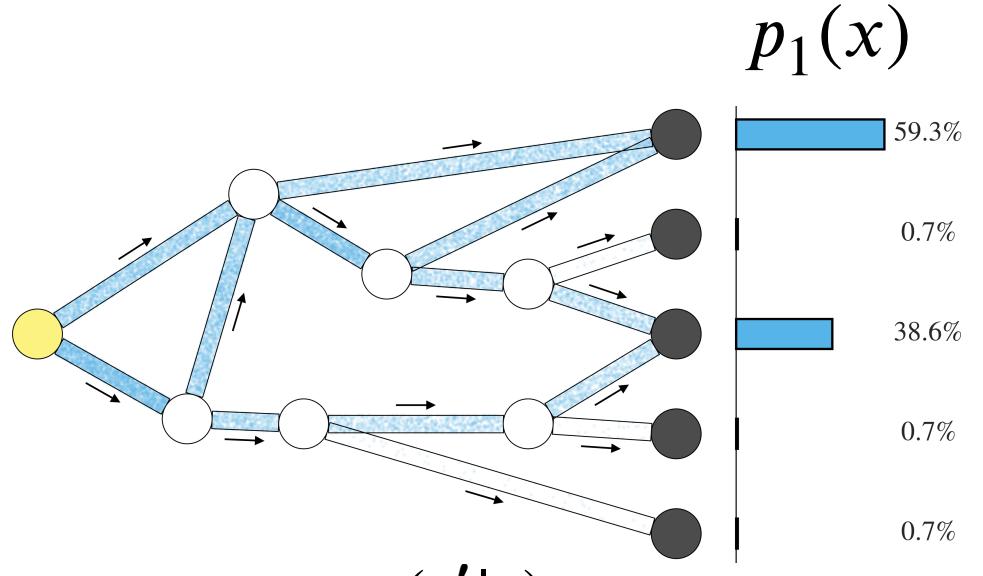
# Composition Tools: Mixtures



# Mixture

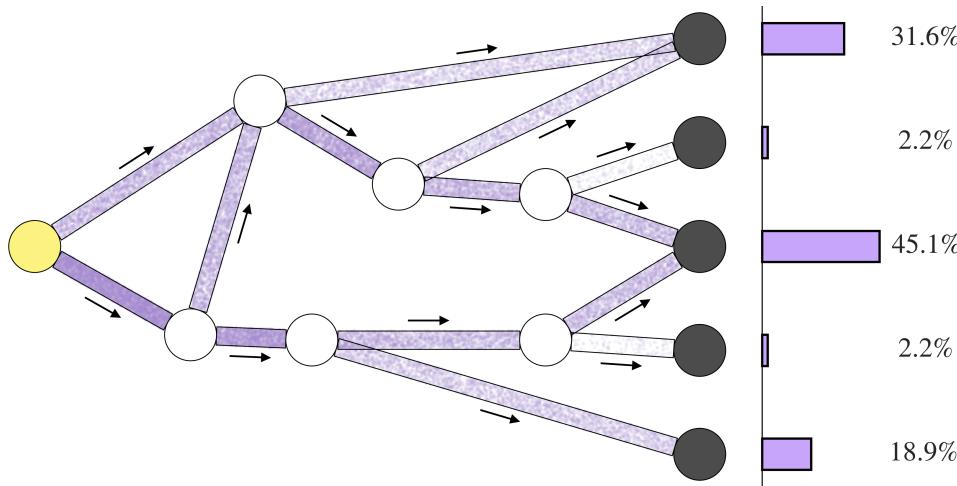
$$p_M(x) = \frac{1}{2}p_1(x) + \frac{1}{2}p_2(x)$$

# Composition Tools: Mixtures



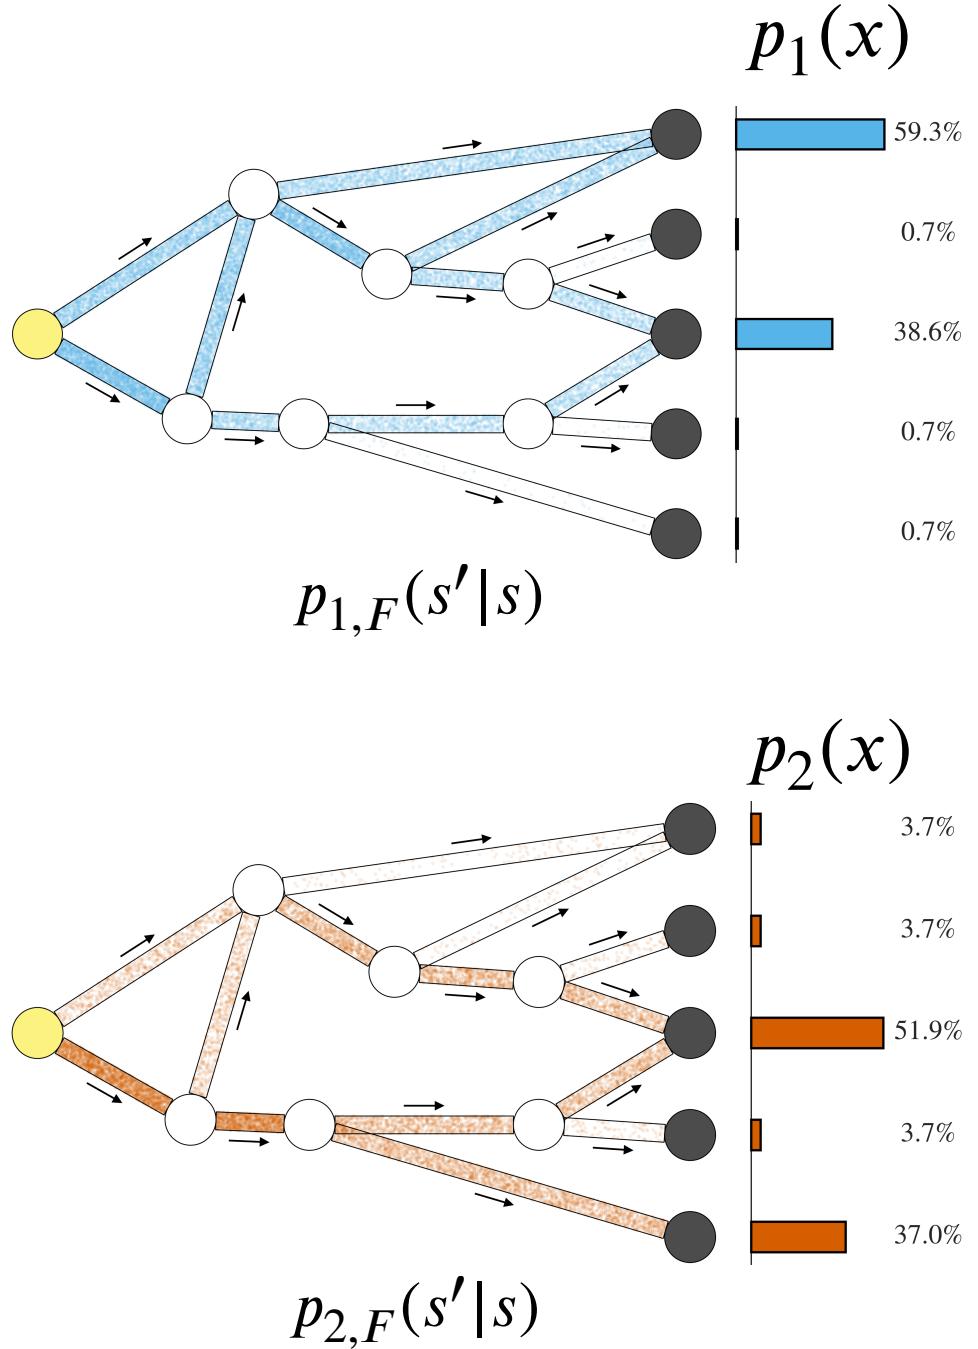
Mixture

$$p_M(x) = \frac{1}{2}p_1(x) + \frac{1}{2}p_2(x)$$



$$p_{M,F}(s'|s) = \sum_{i=1}^2 p(y=i|s)p_{i,F}(s'|s)$$

# Composition Tools: Mixtures



**Mixture**

$$p_M(x) = \frac{1}{2}p_1(x) + \frac{1}{2}p_2(x)$$

Bar charts showing feature distributions for the mixture model:

- $p_M(x)$ : 31.6% (purple), 2.2% (grey), 45.1% (purple), 2.2% (grey), 18.9% (purple)

$$p_{M,F}(s'|s) = \sum_{i=1}^2 p(y=i|s)p_{i,F}(s'|s)$$

**Classifier**

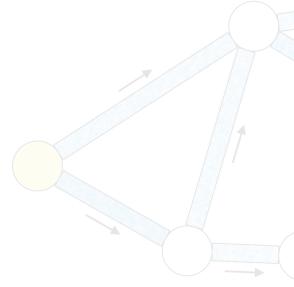
Final output proportions (top to bottom):

- 94/6
- 15/85
- 43/57
- 15/85
- 2/98
- 40/60
- 24/76
- 42/58
- 50/50
- 61/39
- 70/30

$$p(y=i|s) = \frac{p_i(s)}{p_1(s) + p_2(s)}$$

# Composition Tools: Mixtures

$p_1(x)$



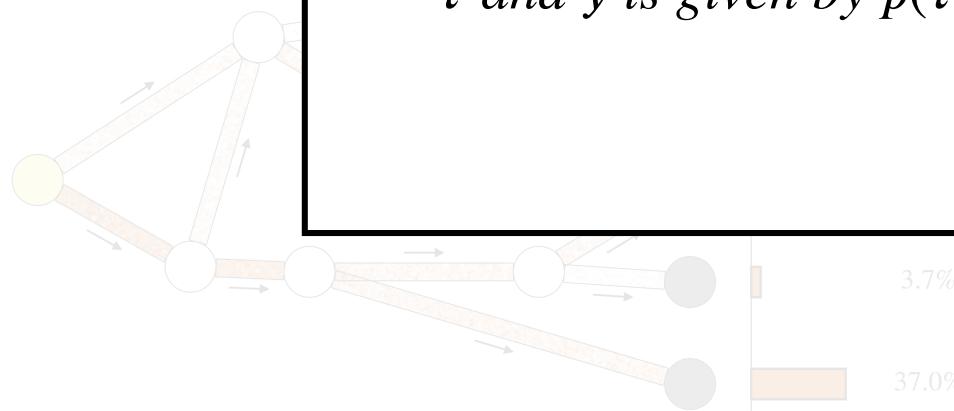
**Proposition (GFlowNet mixture policy).**

Suppose distributions  $p_1(x), \dots, p_m(x)$  are realized by GFlowNets with forward policies  $p_{1,F}(\cdot|\cdot), \dots, p_{m,F}(\cdot|\cdot)$ . Then, the mixture distribution  $p_M(x) = \sum_{i=1}^m \omega_i p_i(x)$  with  $\omega_1, \dots, \omega_m \geq 0$  and  $\sum_{i=1}^m \omega_i = 1$  is realized by the GFlowNet forward policy

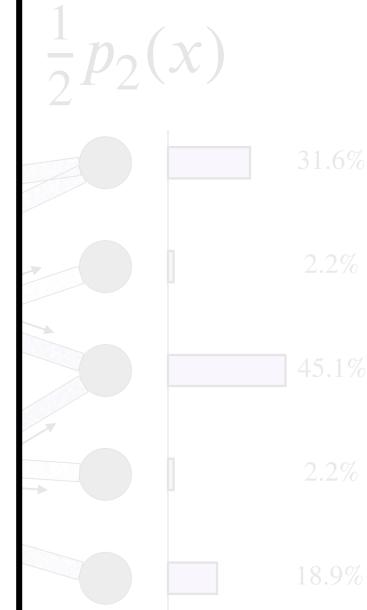
$$p_{M,F}(s'|s) = \sum_{i=1}^m p(y=i|s)p_{i,F}(s'|s),$$

where  $y$  is a random variable such that the joint distribution of a GFlowNet trajectory  $\tau$  and  $y$  is given by  $p(\tau, y=i) = \omega_i p_i(\tau)$  for  $i \in \{1, \dots, m\}$ .

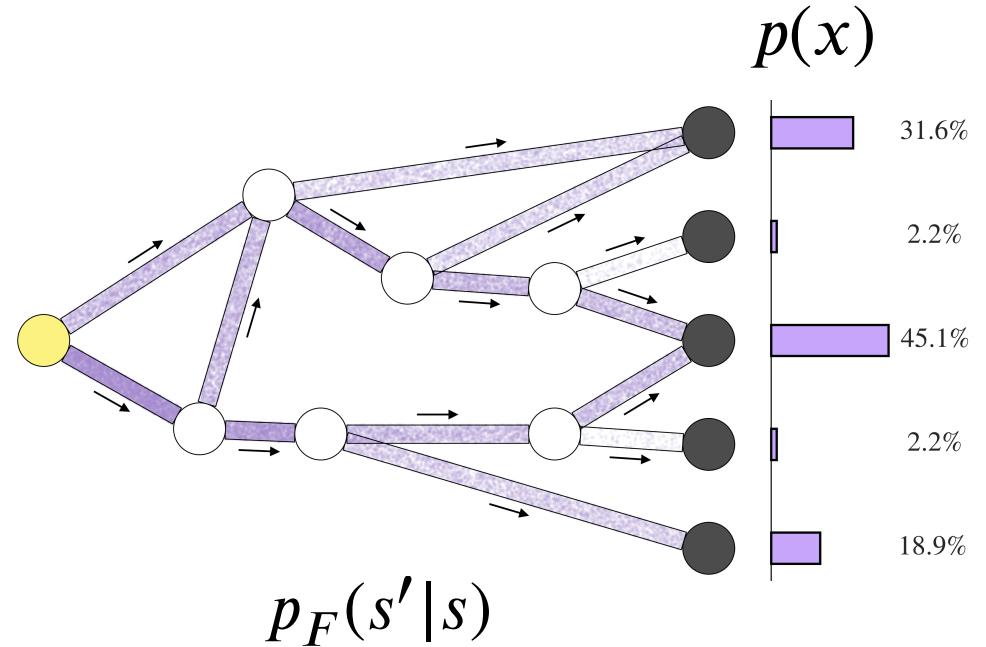
New result for GFlowNets!



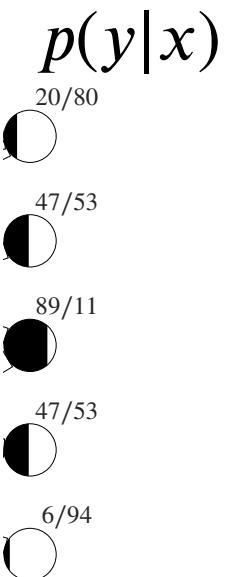
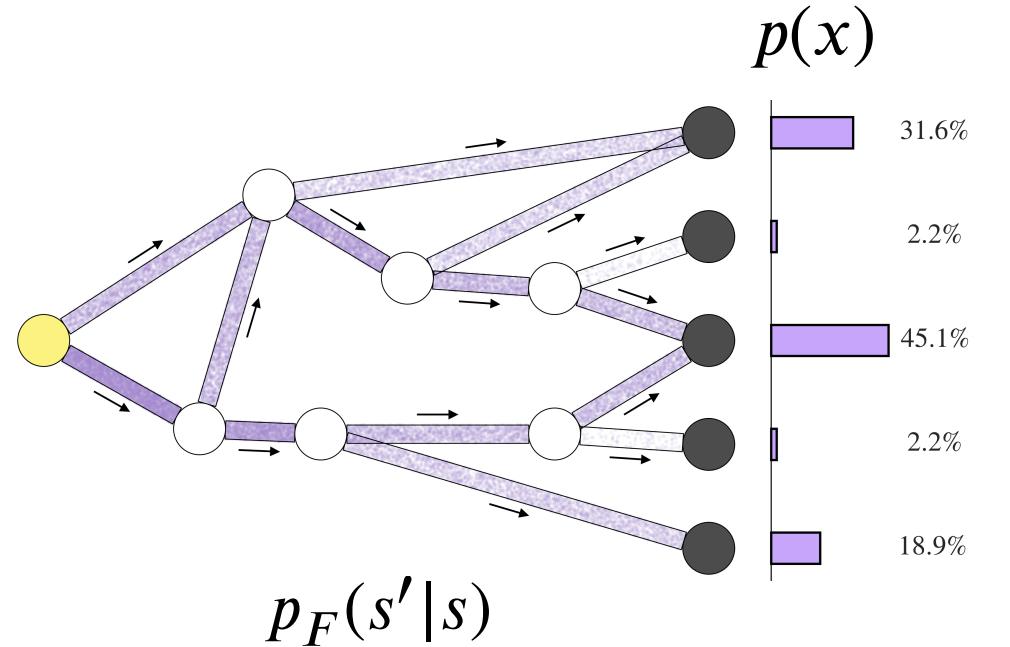
$p_{2,F}(s'|s)$



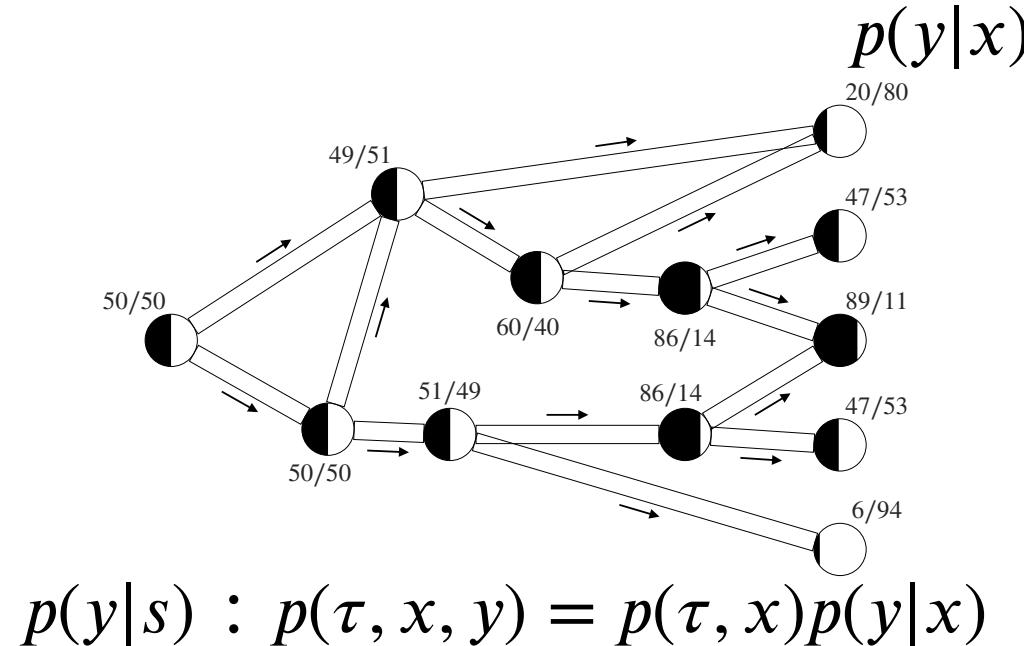
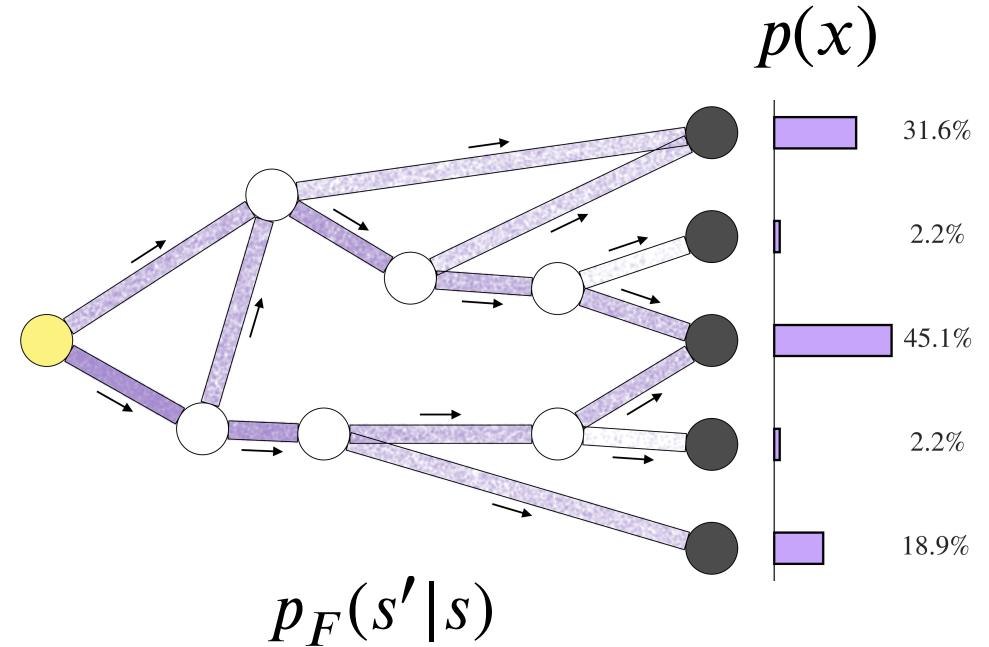
# Composition Tools: Classifier Guidance



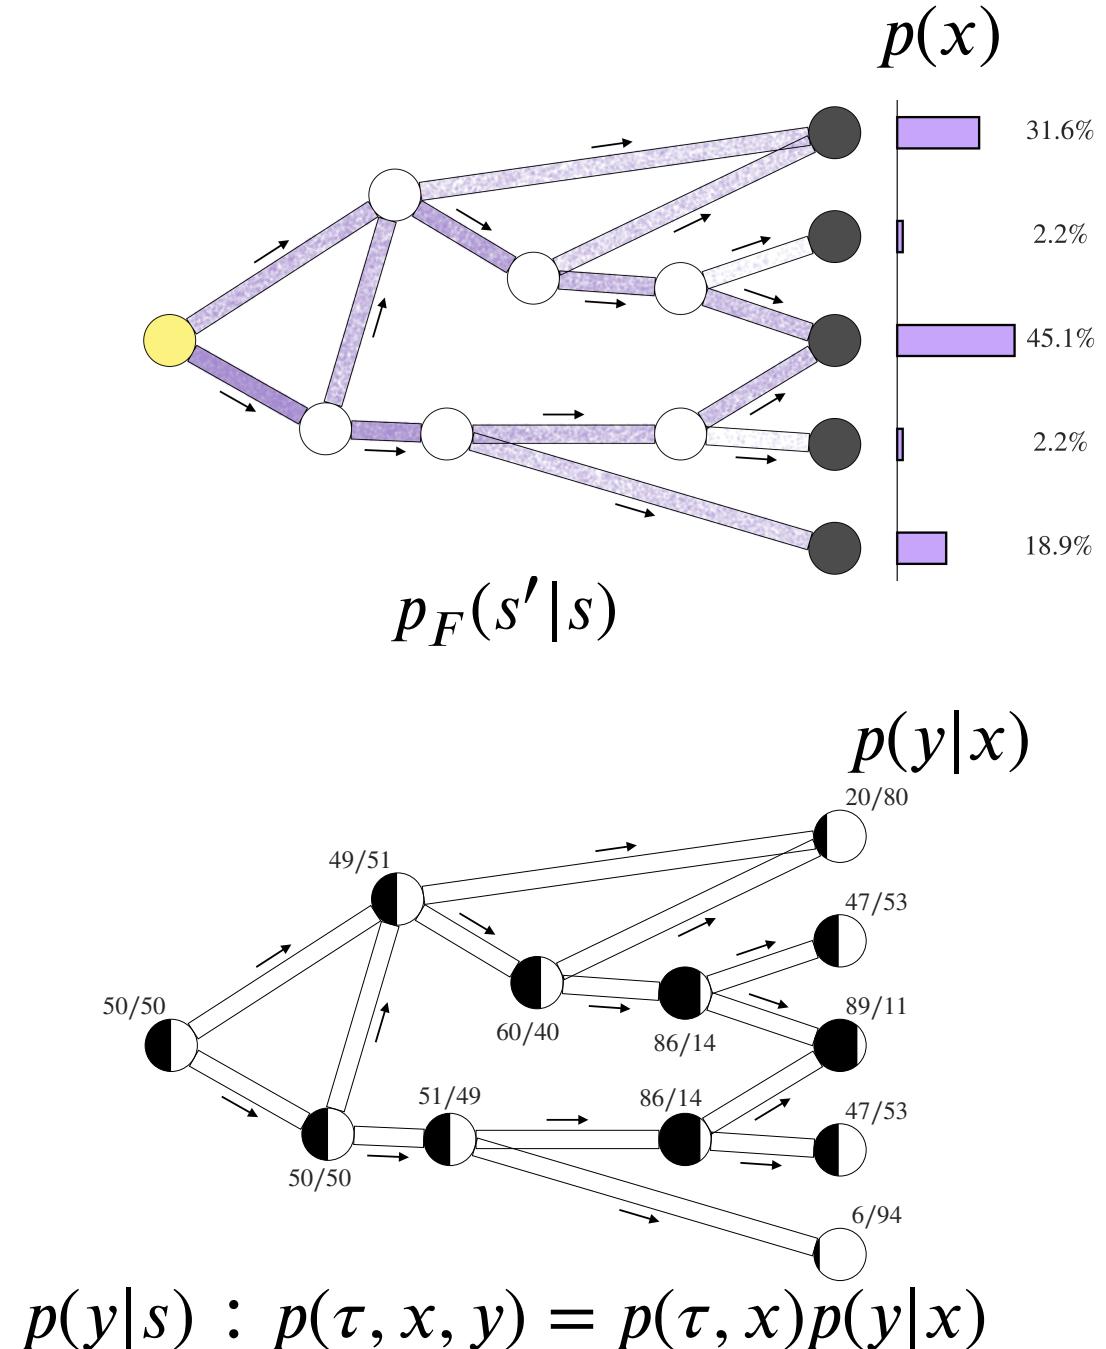
# Composition Tools: Classifier Guidance



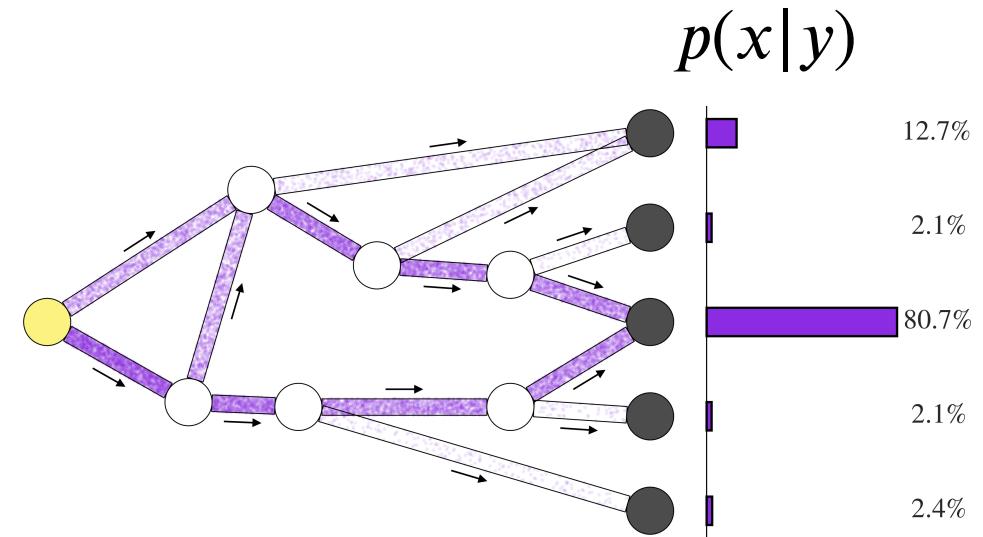
# Composition Tools: Classifier Guidance



# Composition Tools: Classifier Guidance



Classifier-guided GFlowNet



$$p_F(s' | s, y) = p_F(s' | s) \frac{p(y|s')}{p(y|s)}$$

# Composition Tools: Classifier Guidance

$p(x)$

**Proposition** (GFlowNet classifier guidance).

Consider a joint distribution  $p(x, y)$  over a discrete space  $\mathcal{X} \times \mathcal{Y}$  such that the marginal distribution  $p(x)$  is realized by a GFlowNet with forward policy  $p_F(\cdot | \cdot)$ . Further, assume that the joint distribution of  $x$ ,  $y$ , and GFlowNet trajectories  $\tau = (s_0 \rightarrow \dots \rightarrow s_n = x)$  decomposes as  $p(\tau, x, y) = p(\tau, x)p(y|x)$ , i.e.  $y$  is independent of the intermediate states  $s_0, \dots, s_{n-1}$  in  $\tau$  given  $x$ . Then,

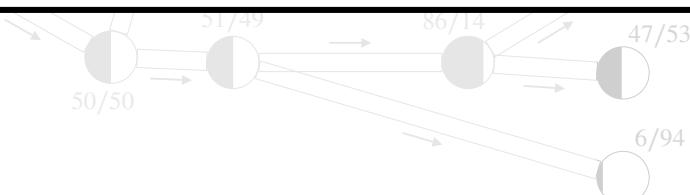
1. For all non-terminal nodes  $s \in \mathcal{S} \setminus \mathcal{X}$  in the GFlowNet DAG  $(\mathcal{S}, \mathcal{A})$ , the probabilities  $p(y|s)$  satisfy

$$p(y|s) = \sum_{s' : (s \rightarrow s') \in \mathcal{A}} p_F(s'|s)p(y|s').$$

2. The conditional distribution  $p(x|y)$  is realized by the classifier-guided policy

$$p_F(s'|s, y) = p_F(s'|s) \frac{p(y|s')}{p(y|s)}.$$

New result for GFlowNets!



$$p(y|s) : p(\tau, x, y) = p(\tau, x)p(y|x)$$

# Composition Operations

Given: 2 distributions  $p_1(x)$ ,  $p_2(x)$

Prior

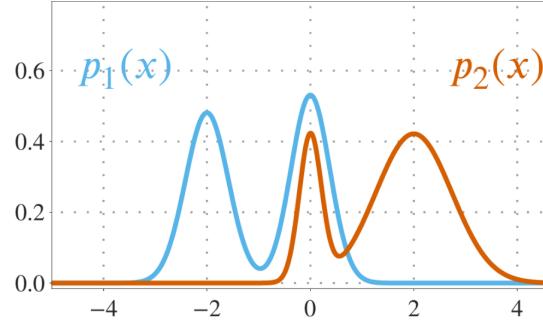
$$\tilde{p}(x) = \frac{1}{2}p_1(x) + \frac{1}{2}p_2(x)$$

Observations

$$\tilde{p}(y_k=i | x) = \frac{p_i(x)}{p_1(x) + p_2(x)}, \quad i \in \{1, 2\}$$

Posterior (composition)

$$\tilde{p}(x | y_1=i, y_2=j) = \frac{p_i(x)p_j(x)}{p_1(x) + p_2(x)}$$



# Composition Operations

Given: 2 distributions  $p_1(x)$ ,  $p_2(x)$

Prior

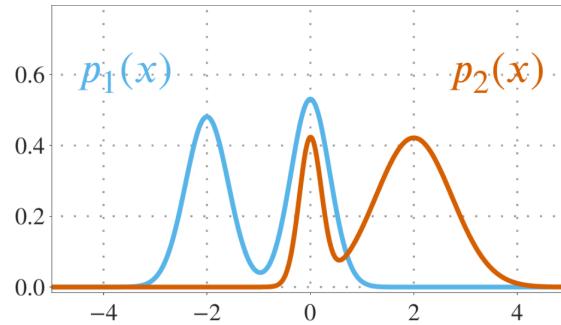
$$\tilde{p}(x) = \frac{1}{2}p_1(x) + \frac{1}{2}p_2(x)$$

Observations

$$\tilde{p}(y_k=i | x) = \frac{p_i(x)}{p_1(x) + p_2(x)}, \quad i \in \{1, 2\}$$

Posterior (composition)

$$\tilde{p}(x | y_1=i, y_2=j) = \frac{p_i(x)p_j(x)}{p_1(x) + p_2(x)}$$



“Harmonic Mean”

$$(p_1 \otimes p_2)(x) = \tilde{p}(x | y_1=1, y_2=2) \propto \frac{p_1(x)p_2(x)}{p_1(x) + p_2(x)}$$

# Composition Operations

Given: 2 distributions  $p_1(x), p_2(x)$

**Prior**

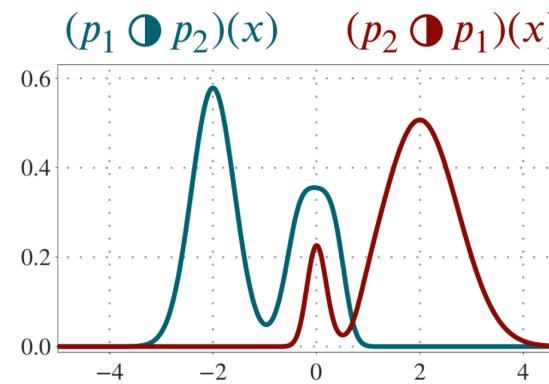
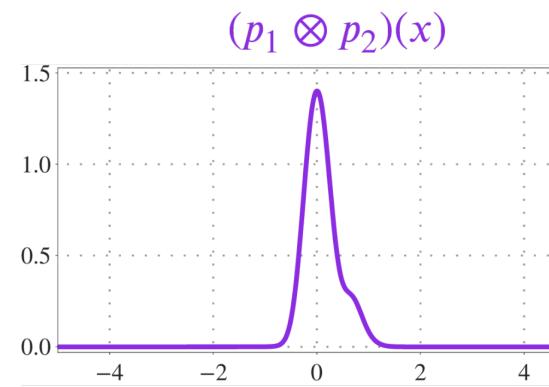
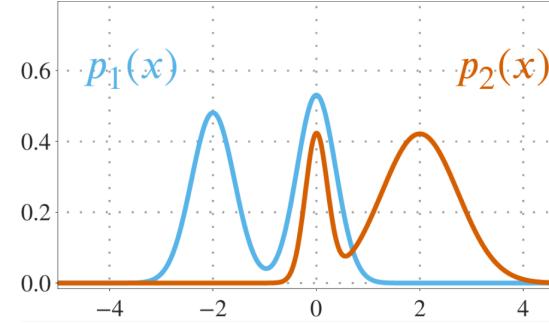
$$\tilde{p}(x) = \frac{1}{2}p_1(x) + \frac{1}{2}p_2(x)$$

**Observations**

$$\tilde{p}(y_k=i | x) = \frac{p_i(x)}{p_1(x) + p_2(x)}, \quad i \in \{1, 2\}$$

**Posterior (composition)**

$$\tilde{p}(x | y_1=i, y_2=j) = \frac{p_i(x)p_j(x)}{p_1(x) + p_2(x)}$$



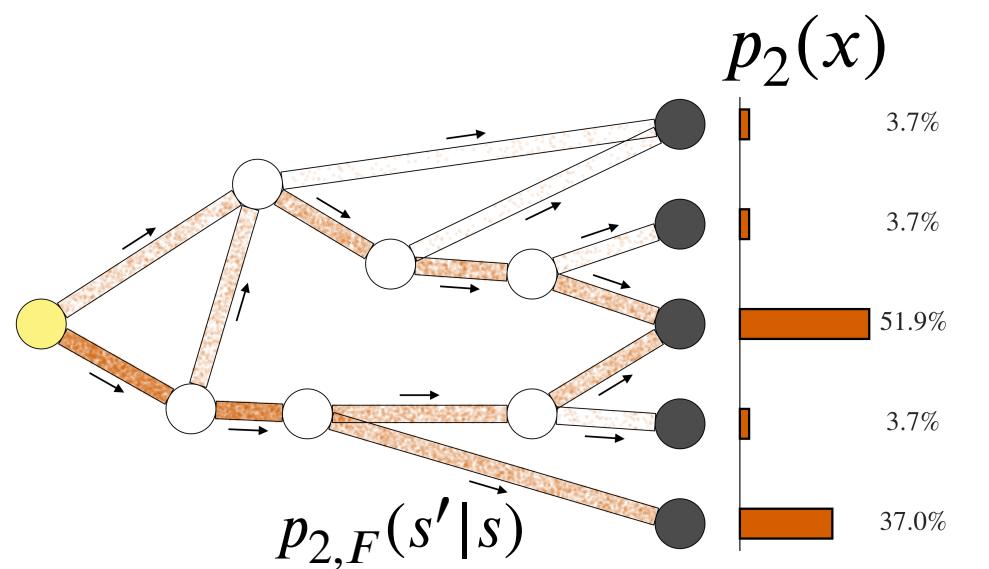
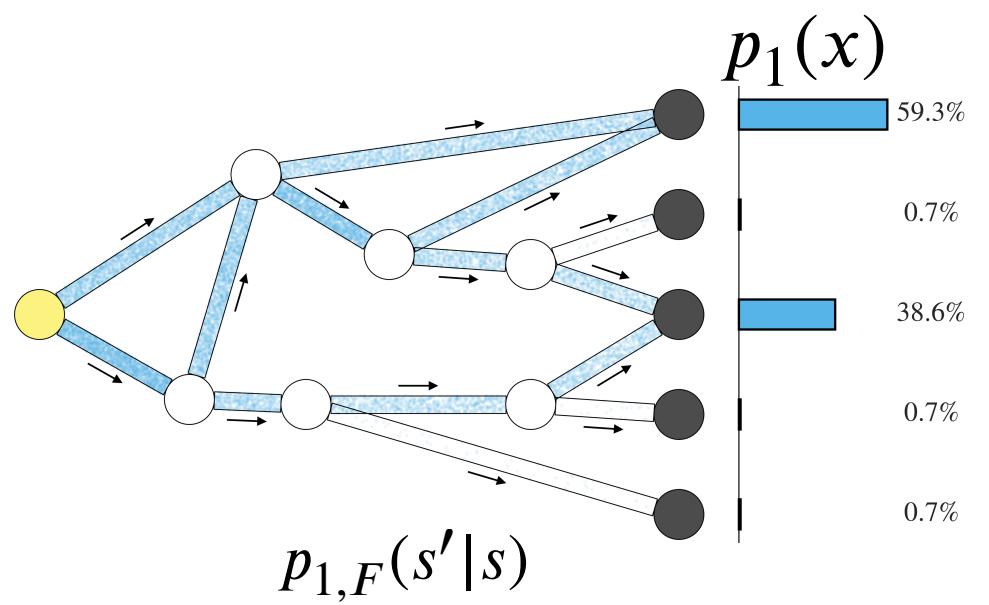
**“Harmonic Mean”**

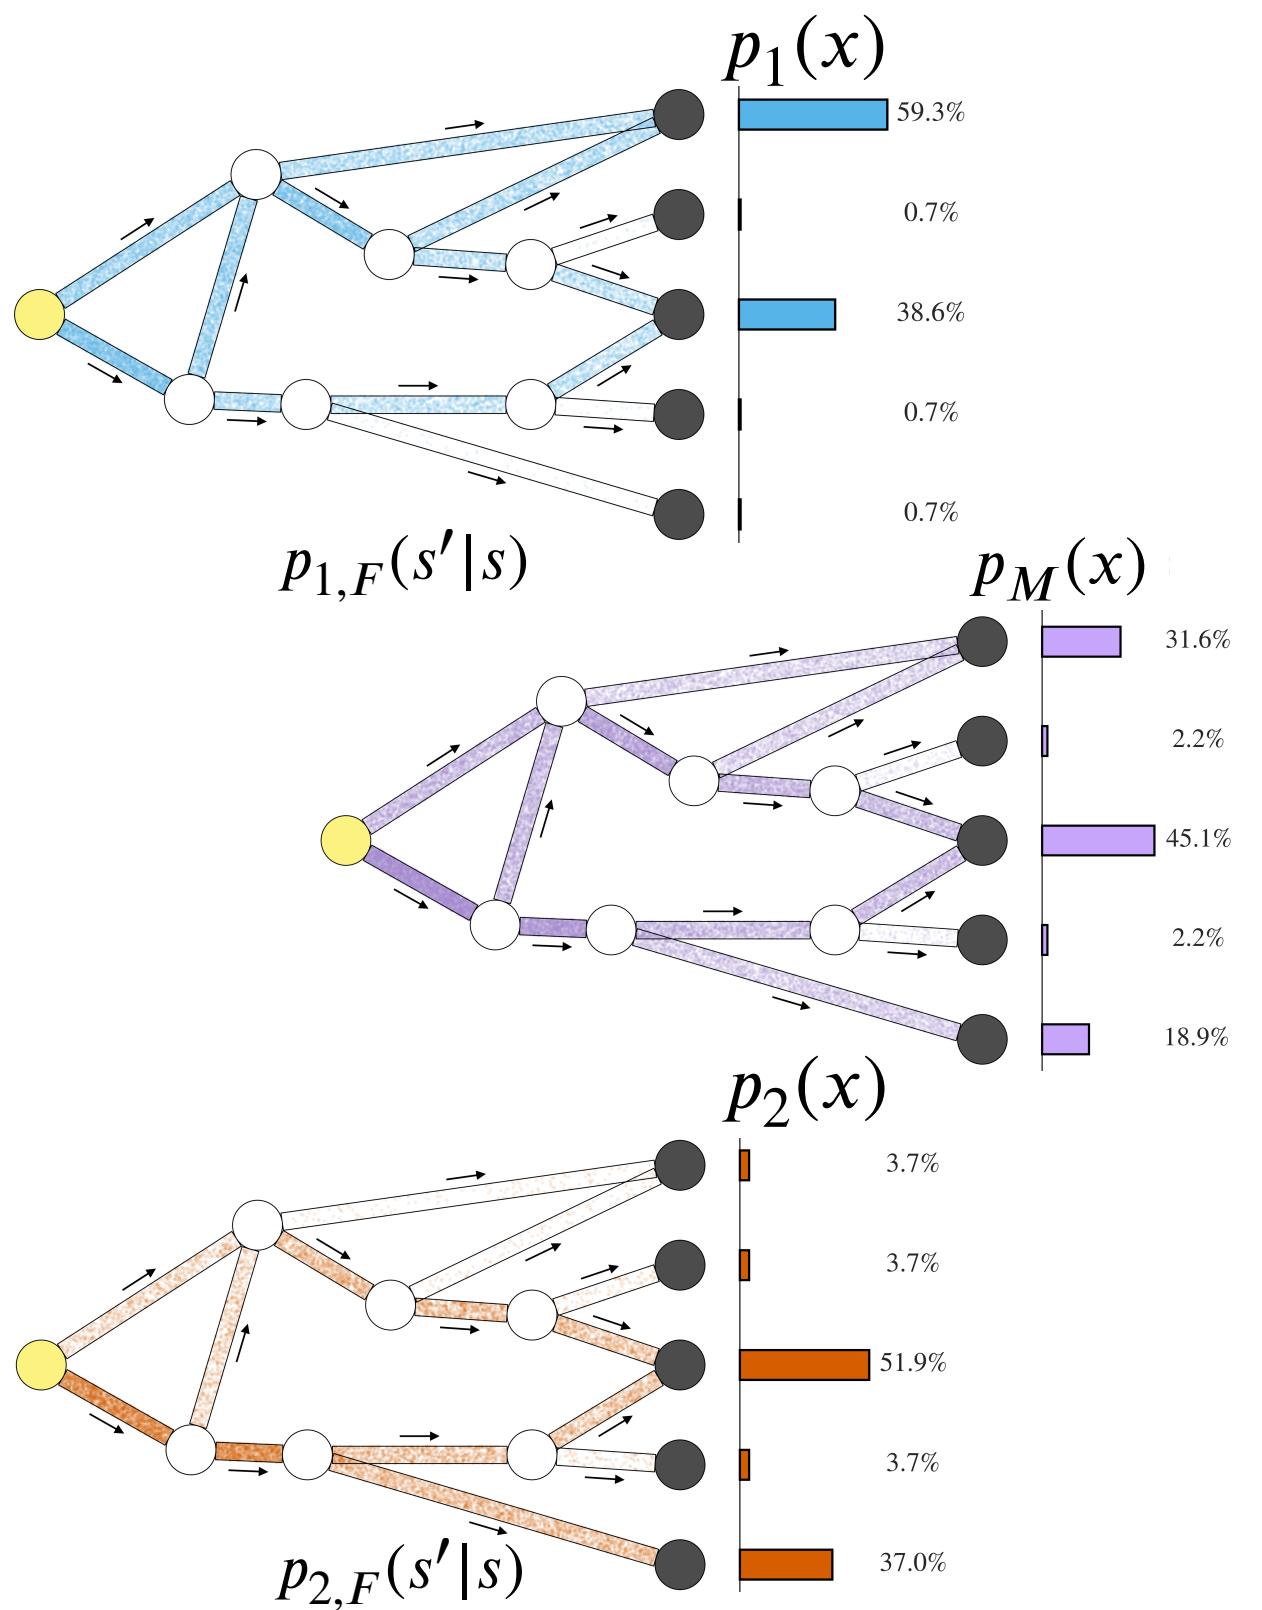
$$(p_1 \otimes p_2)(x) = \tilde{p}(x | y_1=1, y_2=2) \propto \frac{p_1(x)p_2(x)}{p_1(x) + p_2(x)}$$

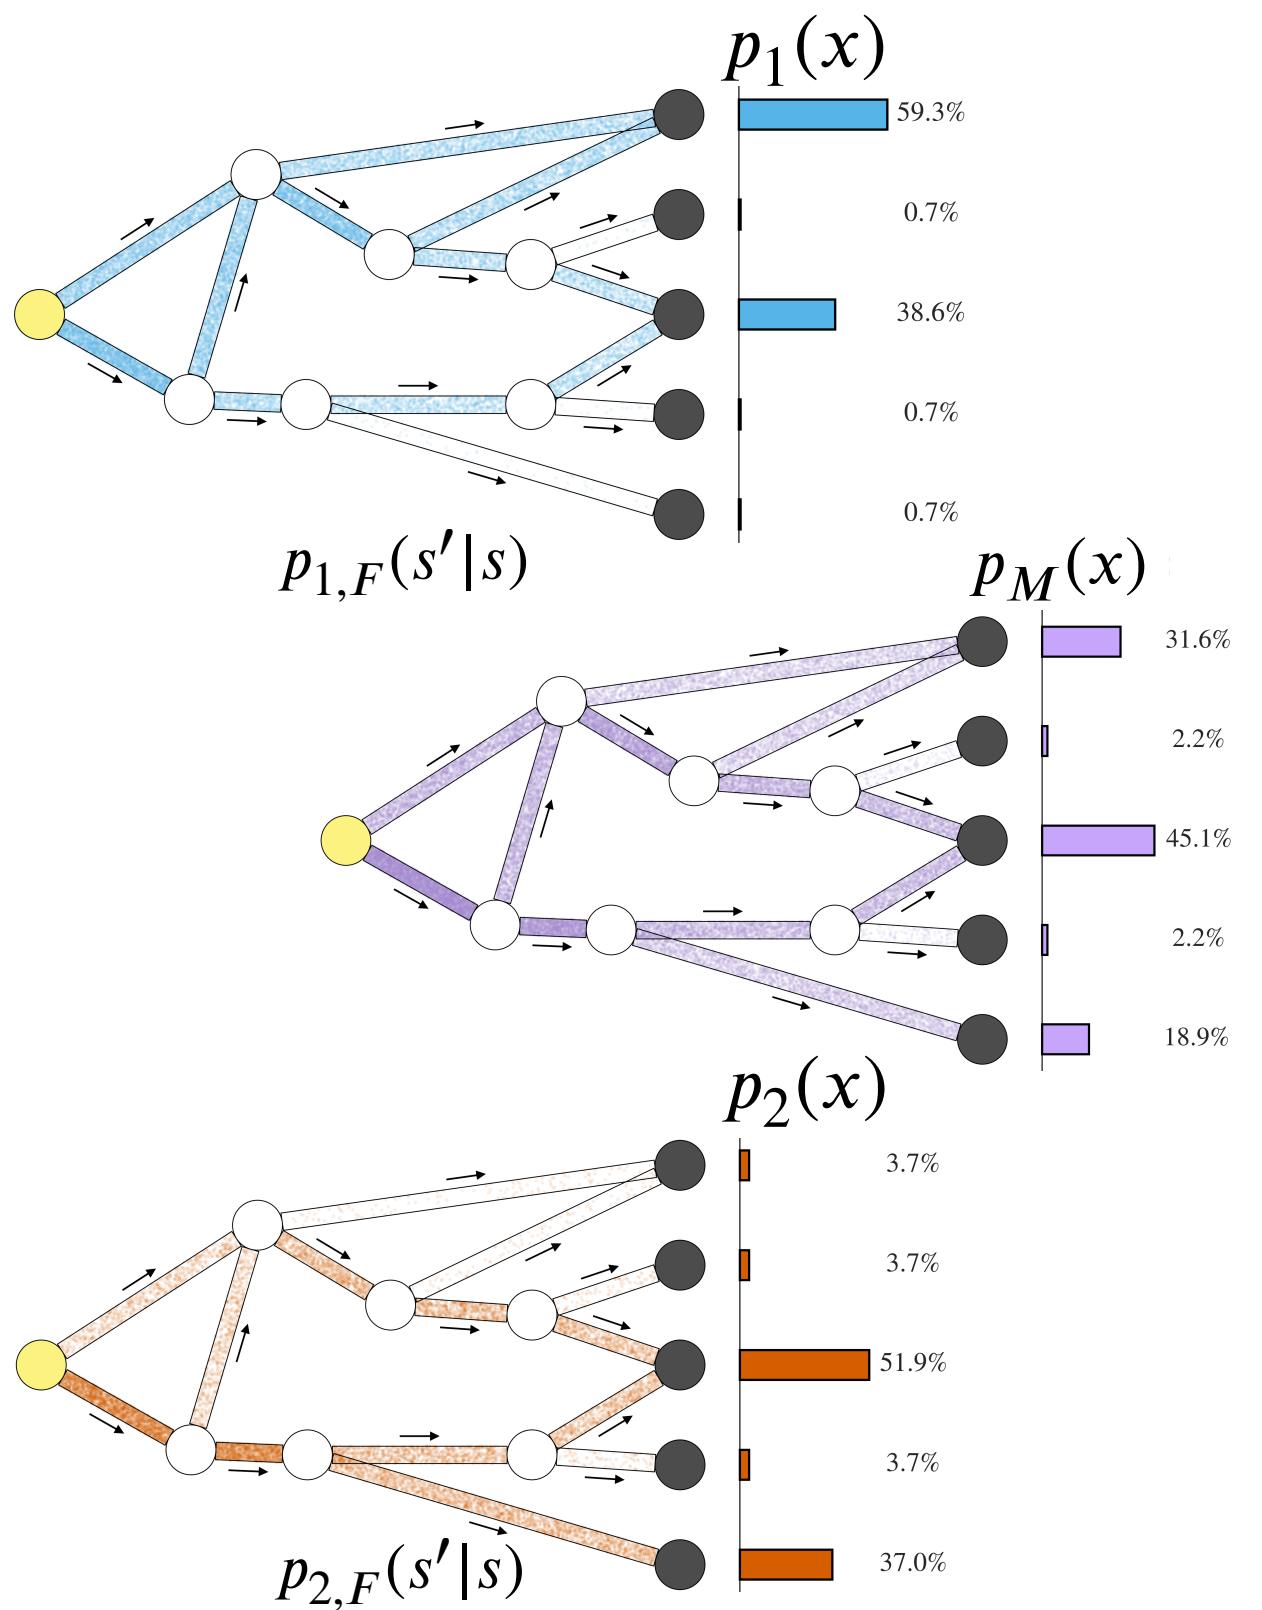
**“Contrast”**

$$(p_1 \odot p_2)(x) = \tilde{p}(x | y_1=1, y_2=1) \propto \frac{(p_1(x))^2}{p_1(x) + p_2(x)}$$

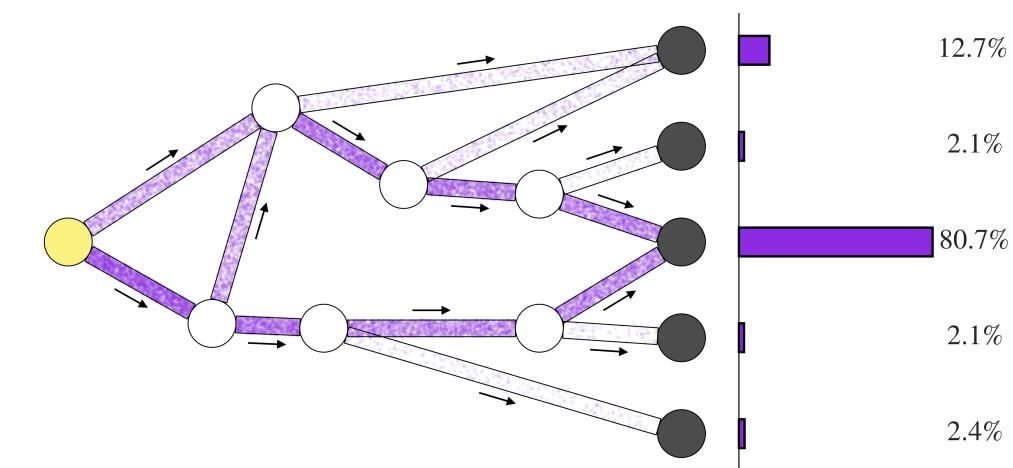
$$(p_2 \odot p_1)(x) = \tilde{p}(x | y_1=2, y_2=2) \propto \frac{(p_2(x))^2}{p_1(x) + p_2(x)}$$



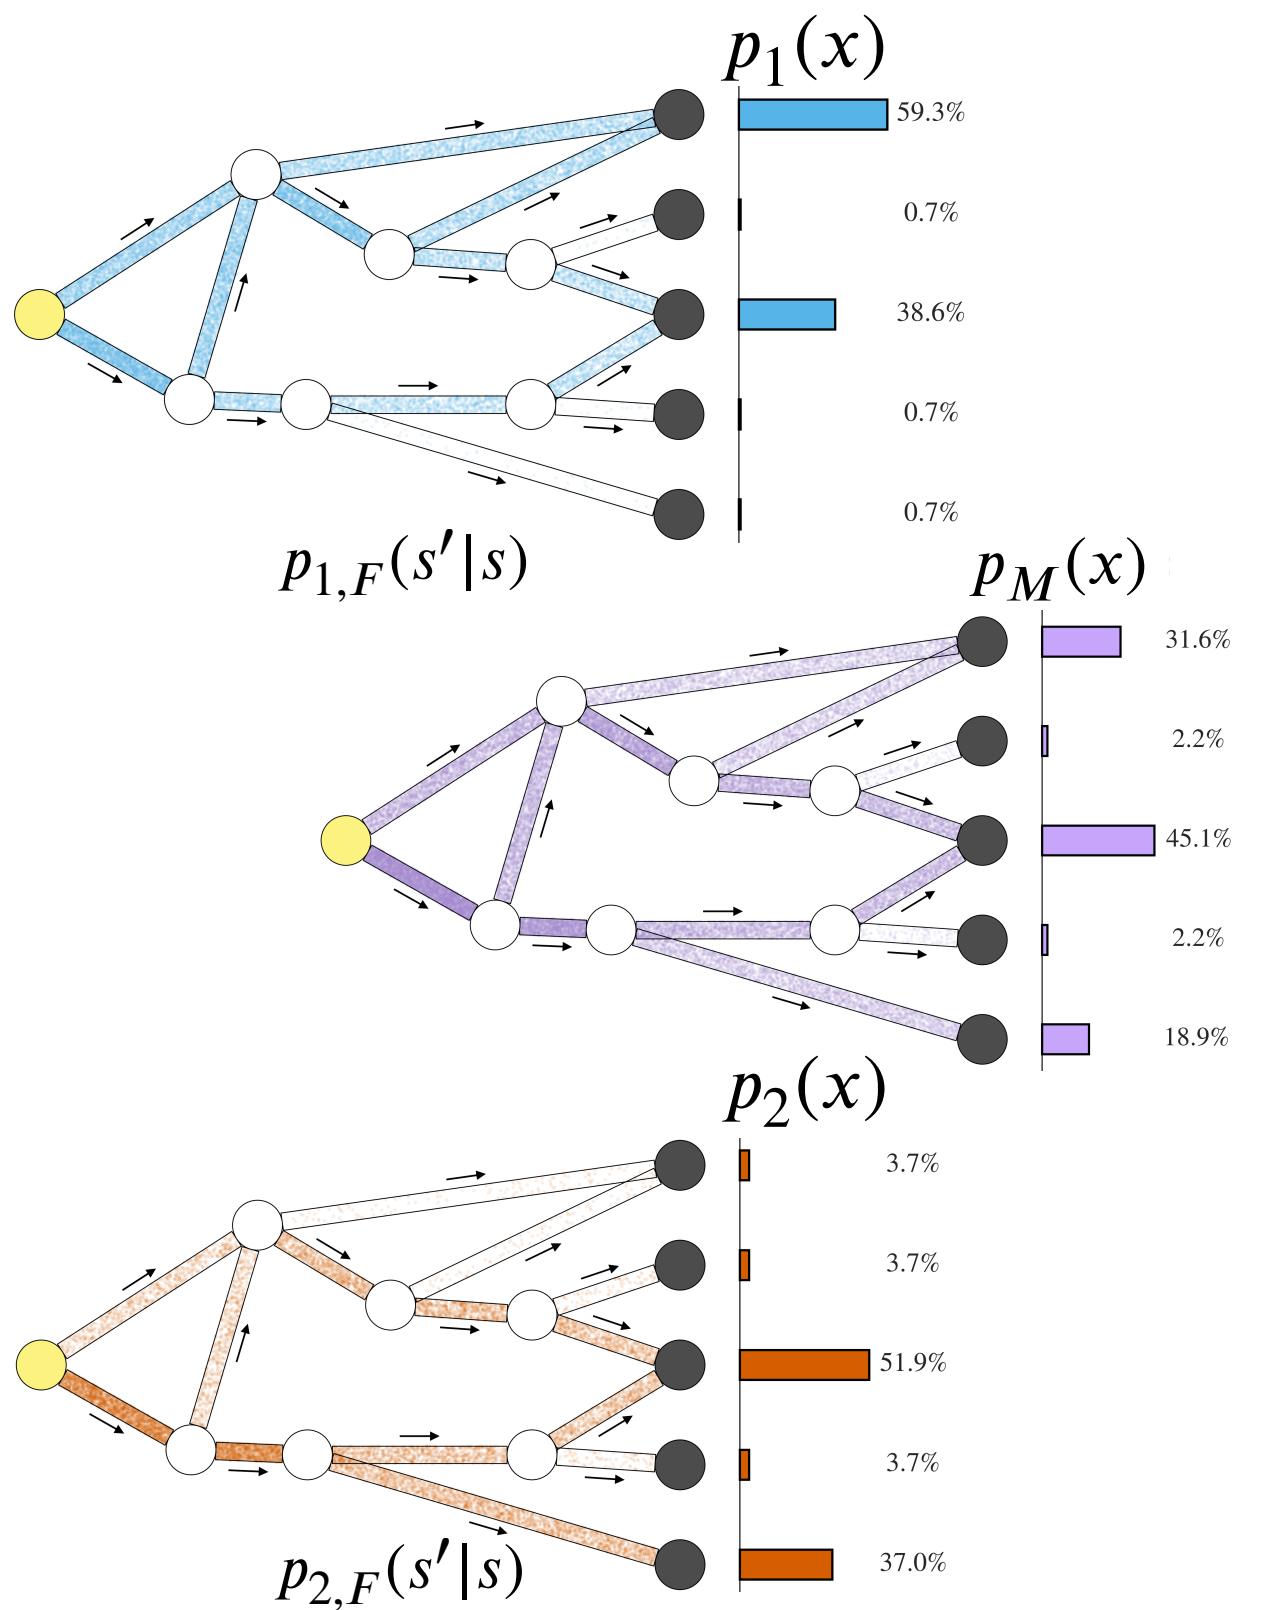




$$(p_1 \otimes p_2)(x)$$

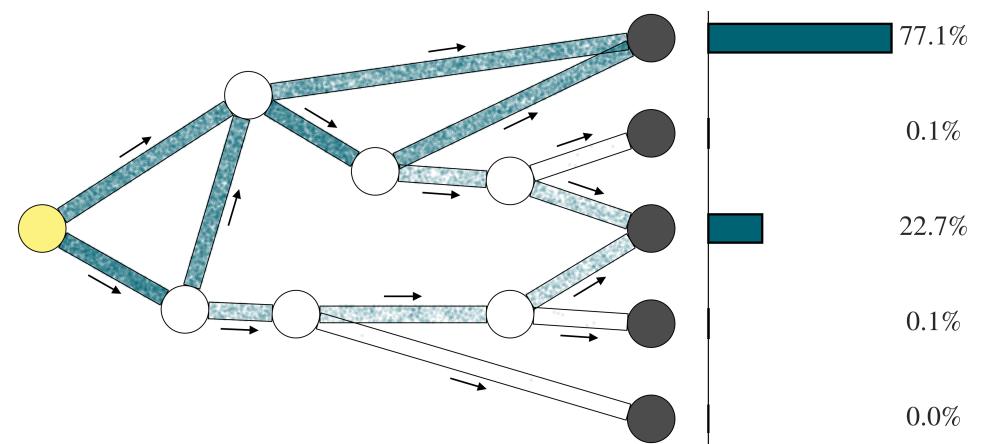


$$p_F(s'|s, y_1=1, y_2=2)$$



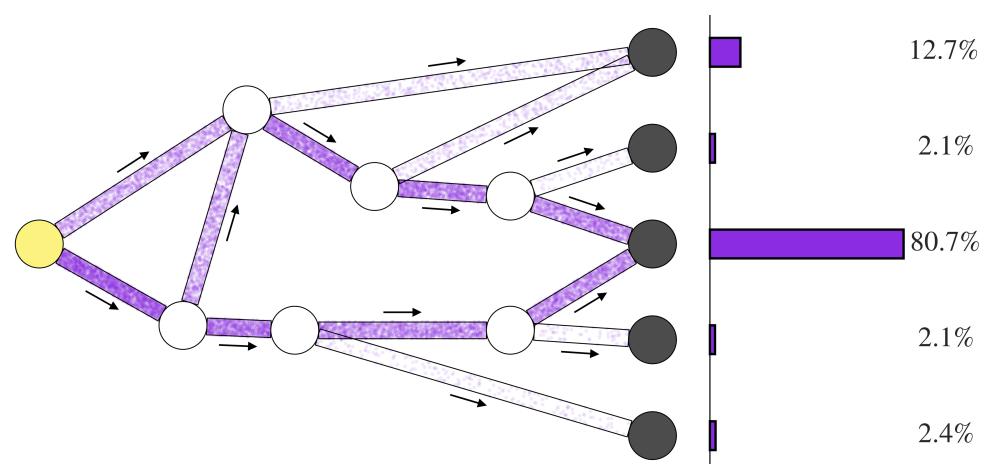
$(p_1 \bullet p_2)(x)$

$p_F(s'|s, y_1=1, y_2=1)$



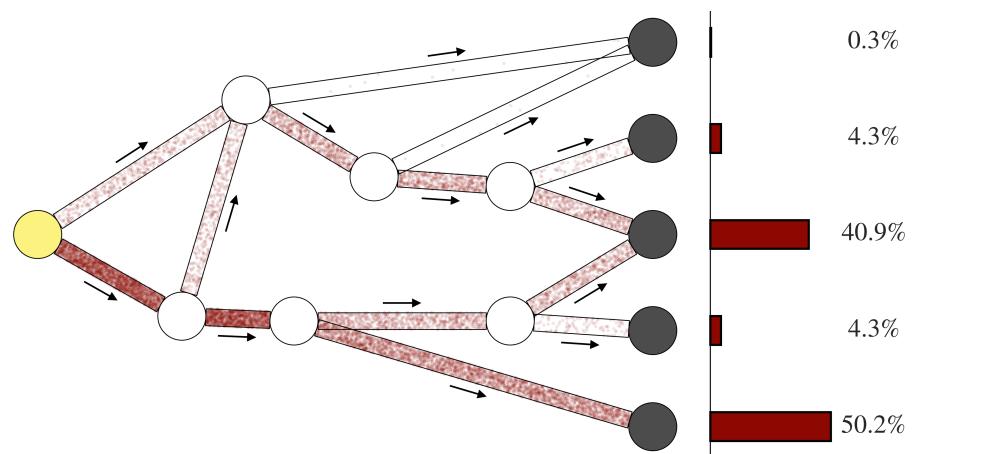
$(p_1 \otimes p_2)(x)$

$p_F(s'|s, y_1=1, y_2=2)$

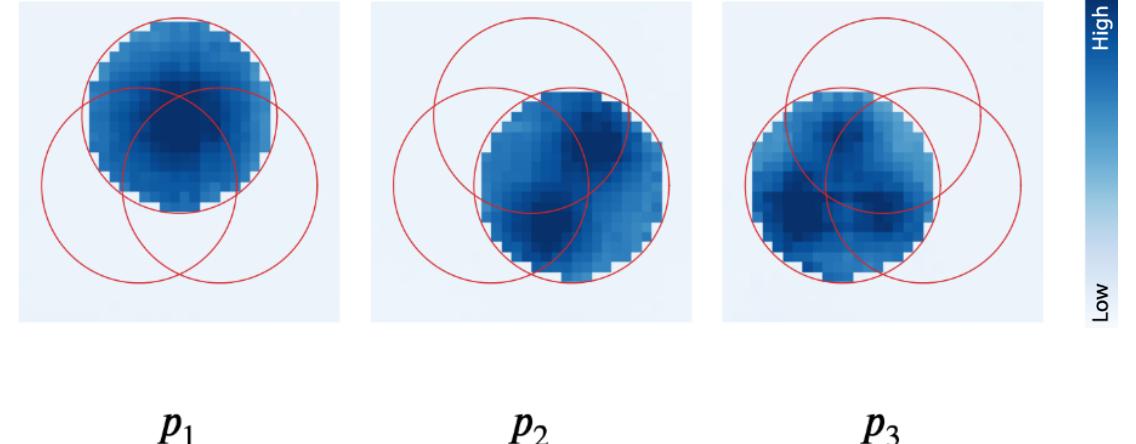


$(p_2 \bullet p_1)(x)$

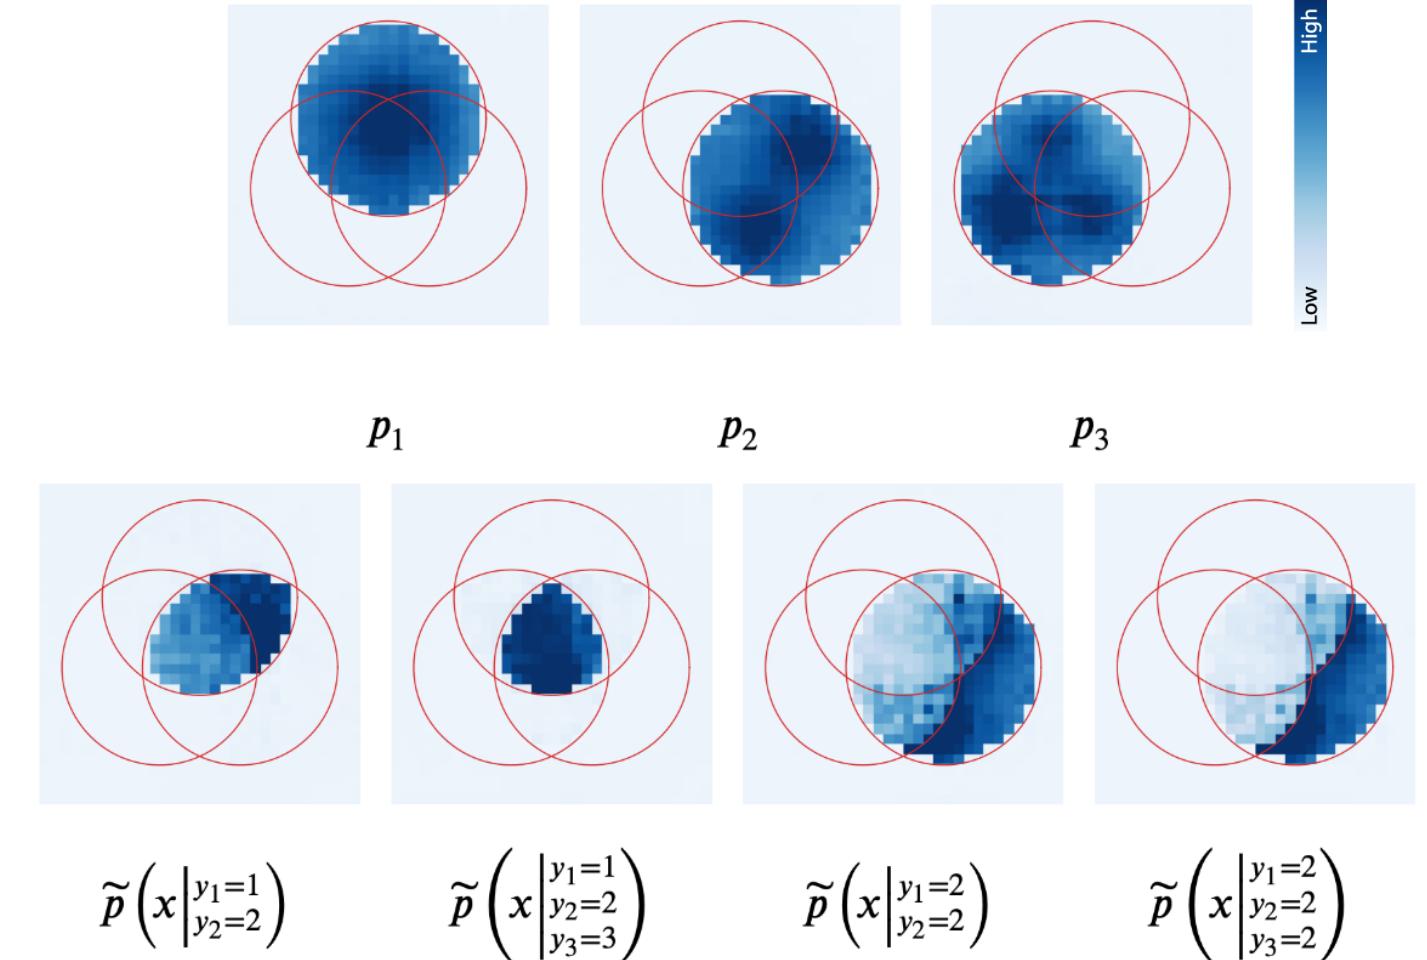
$p_F(s'|s, y_1=2, y_2=2)$



# Composition Operations



# Composition Operations

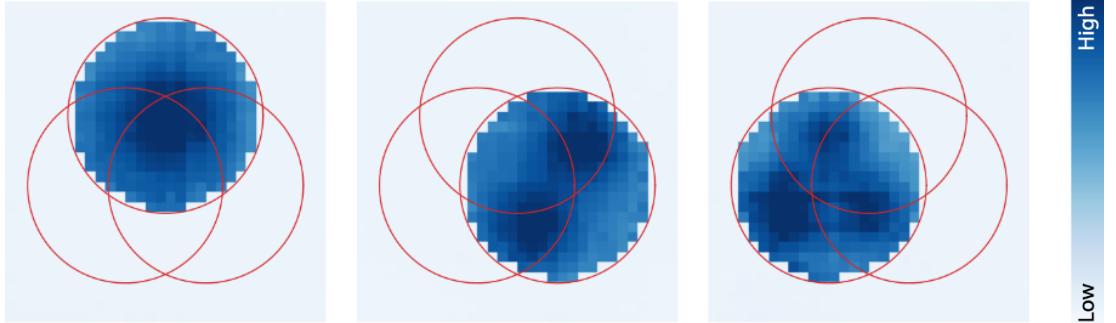


# Composition Operations

Given:  $m$  distributions  $p_1(x), \dots, p_m(x)$

**Prior**

$$\tilde{p}(x) = \frac{1}{m} \sum_{j=1}^m p_j(x)$$



**Observations**

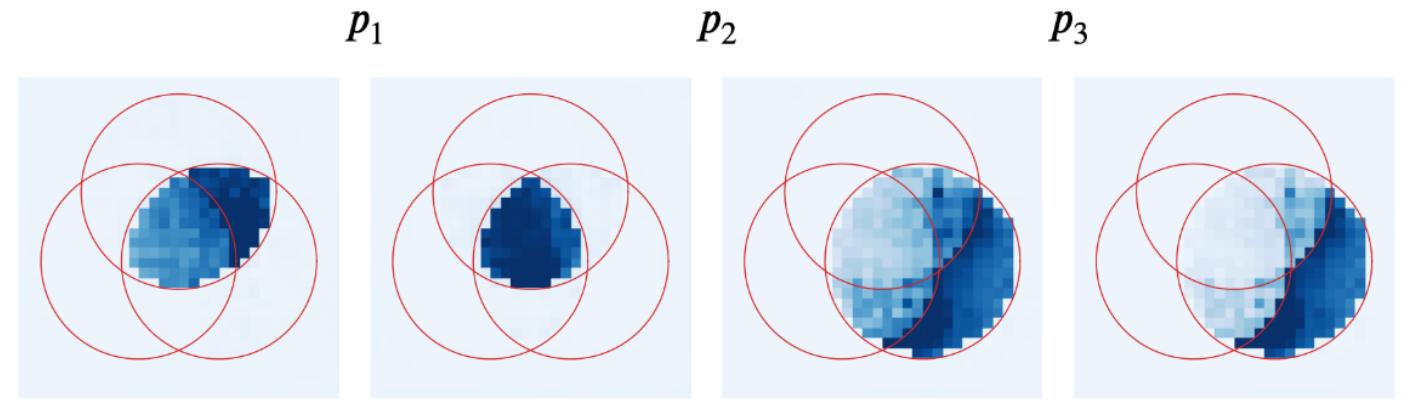
$$\tilde{p}(y_k=i \mid x) = \frac{p_i(x)}{\sum_{j=1}^m p_j(x)}, \quad i \in \{1, \dots, m\}$$

**Joint**

$$\tilde{p}(x, y_1, \dots, y_n) = \tilde{p}(x) \prod_{k=1}^n \tilde{p}(y_k \mid x)$$

**Posterior (composition)**

$$\tilde{p}(x \mid y_1=i_1, \dots, y_n=i_n) \propto \frac{\prod_{k=1}^n p_{i_k}(x)}{\left( \sum_{j=1}^m p_j(x) \right)^{n-1}}$$



$$\tilde{p}\left(x \middle| \begin{matrix} y_1=1 \\ y_2=2 \end{matrix}\right)$$

$$\tilde{p}\left(x \middle| \begin{matrix} y_1=1 \\ y_2=2 \\ y_3=3 \end{matrix}\right)$$

$$\tilde{p}\left(x \middle| \begin{matrix} y_1=2 \\ y_2=2 \end{matrix}\right)$$

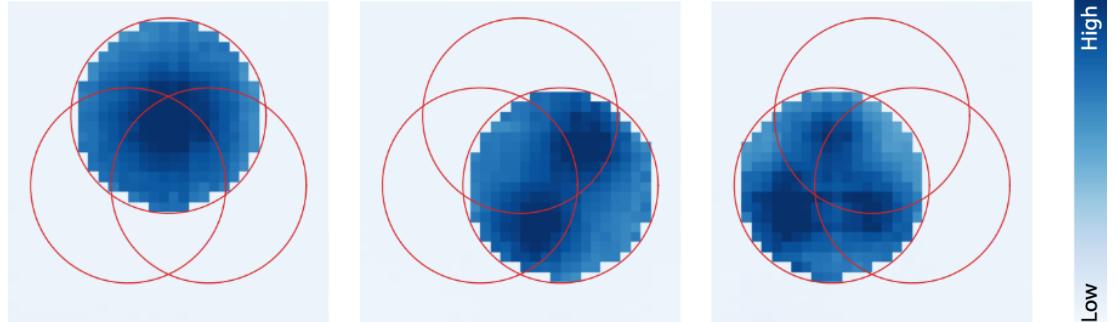
$$\tilde{p}\left(x \middle| \begin{matrix} y_1=2 \\ y_2=2 \\ y_3=2 \end{matrix}\right)$$

# Composition Operations

Given:  $m$  distributions  $p_1(x), \dots, p_m(x)$

Prior

$$\tilde{p}(x) = \frac{1}{m} \sum_{j=1}^m p_j(x)$$



Observations

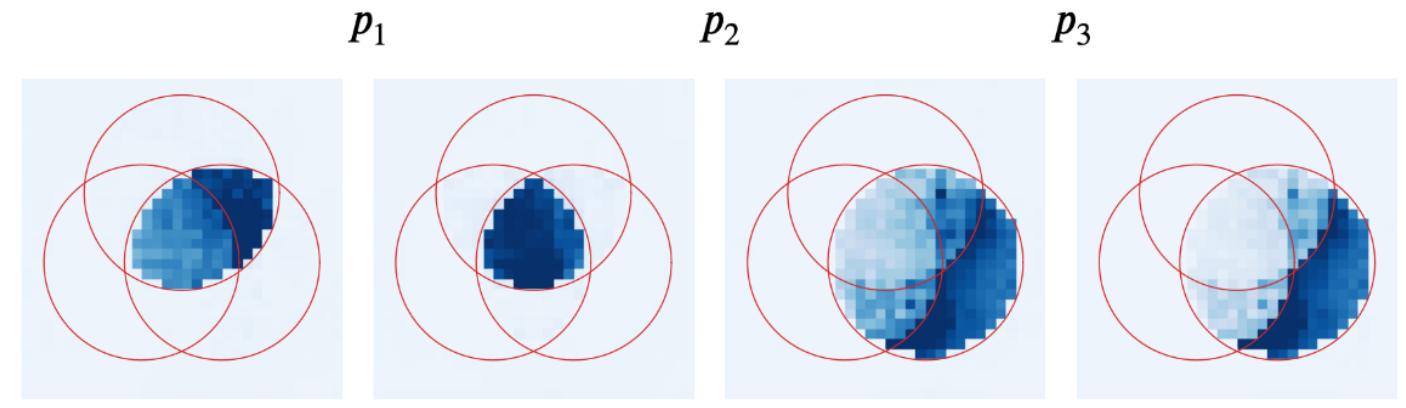
$$\tilde{p}(y_k=i \mid x) = \frac{p_i(x)}{\sum_{j=1}^m p_j(x)}, \quad i \in \{1, \dots, m\}$$

Joint

$$\tilde{p}(x, y_1, \dots, y_n) = \tilde{p}(x) \prod_{k=1}^n \tilde{p}(y_k \mid x)$$

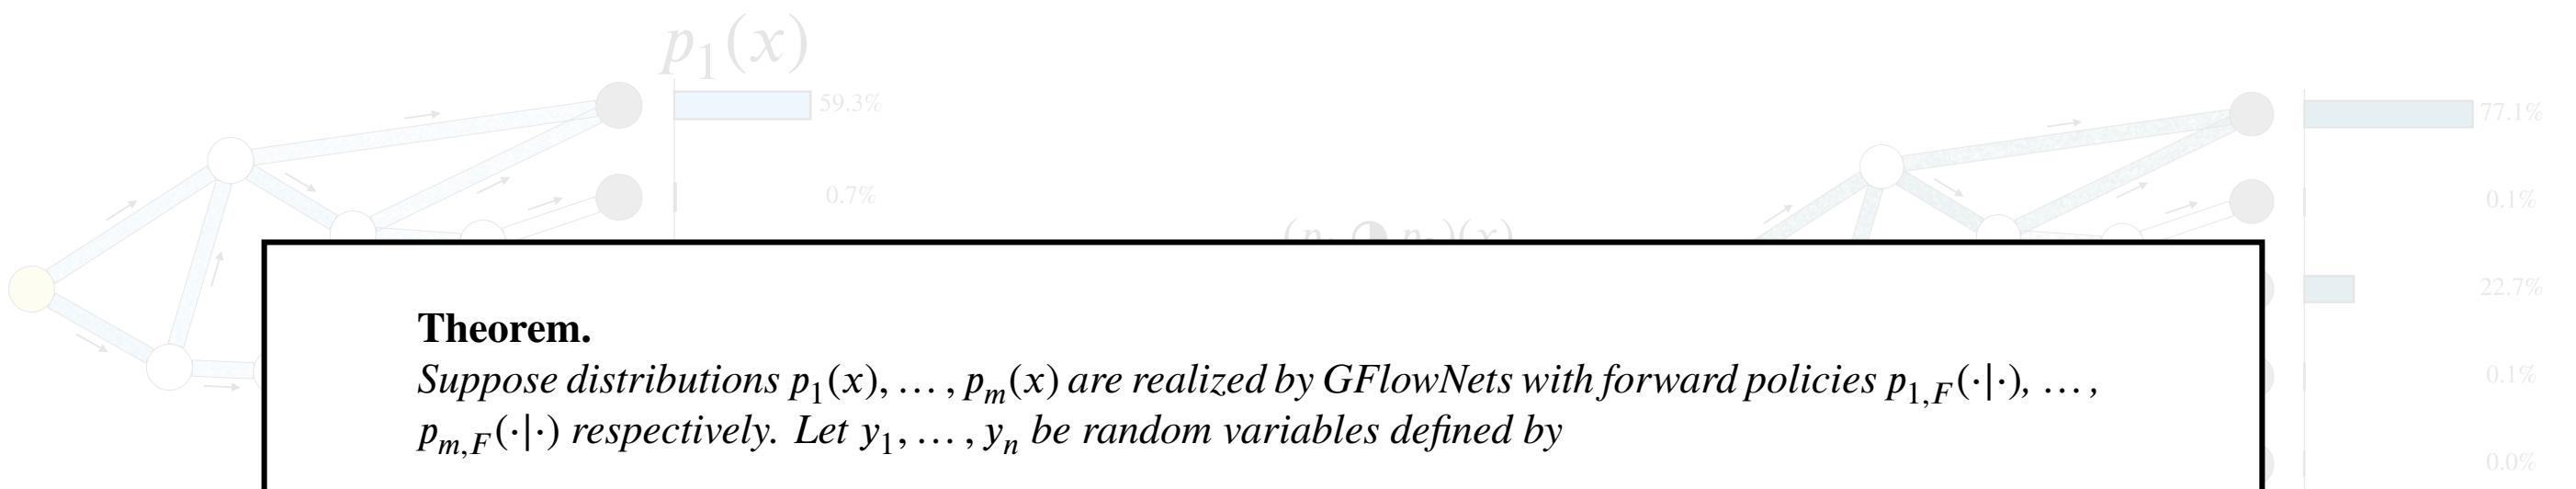
Posterior (composition)

$$\tilde{p}(x \mid y_1=i_1, \dots, y_n=i_n) \propto \frac{\prod_{k=1}^n p_{i_k}(x)}{\left(\sum_{j=1}^m p_j(x)\right)^{n-1}}$$



$$\begin{aligned} &\tilde{p}(x \mid y_1=1, y_2=2) \\ &\tilde{p}(x \mid y_1=1, y_2=2, y_3=3) \\ &\tilde{p}(x \mid y_1=2, y_2=2) \\ &\tilde{p}(x \mid y_1=2, y_2=2, y_3=2) \end{aligned}$$

Control of resulting distribution via observations  
+ more tools for control in the thesis



**Theorem.**

Suppose distributions  $p_1(x), \dots, p_m(x)$  are realized by GFlowNets with forward policies  $p_{1,F}(\cdot| \cdot), \dots, p_{m,F}(\cdot| \cdot)$  respectively. Let  $y_1, \dots, y_n$  be random variables defined by

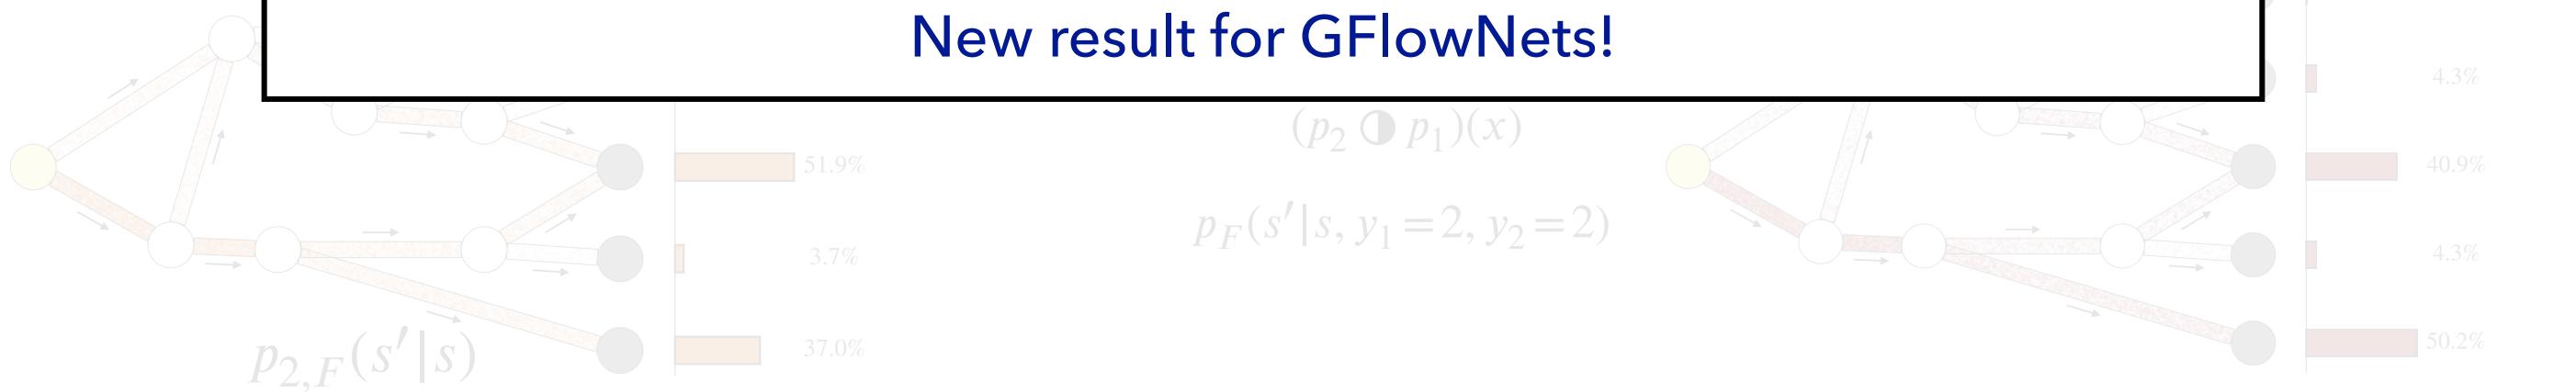
$$\tilde{p}(x, y_1, \dots, y_n) = \tilde{p}(x) \prod_{k=1}^n \tilde{p}(y_k|x), \quad \tilde{p}(x) = \frac{1}{m} \sum_{i=1}^m p_i(x),$$

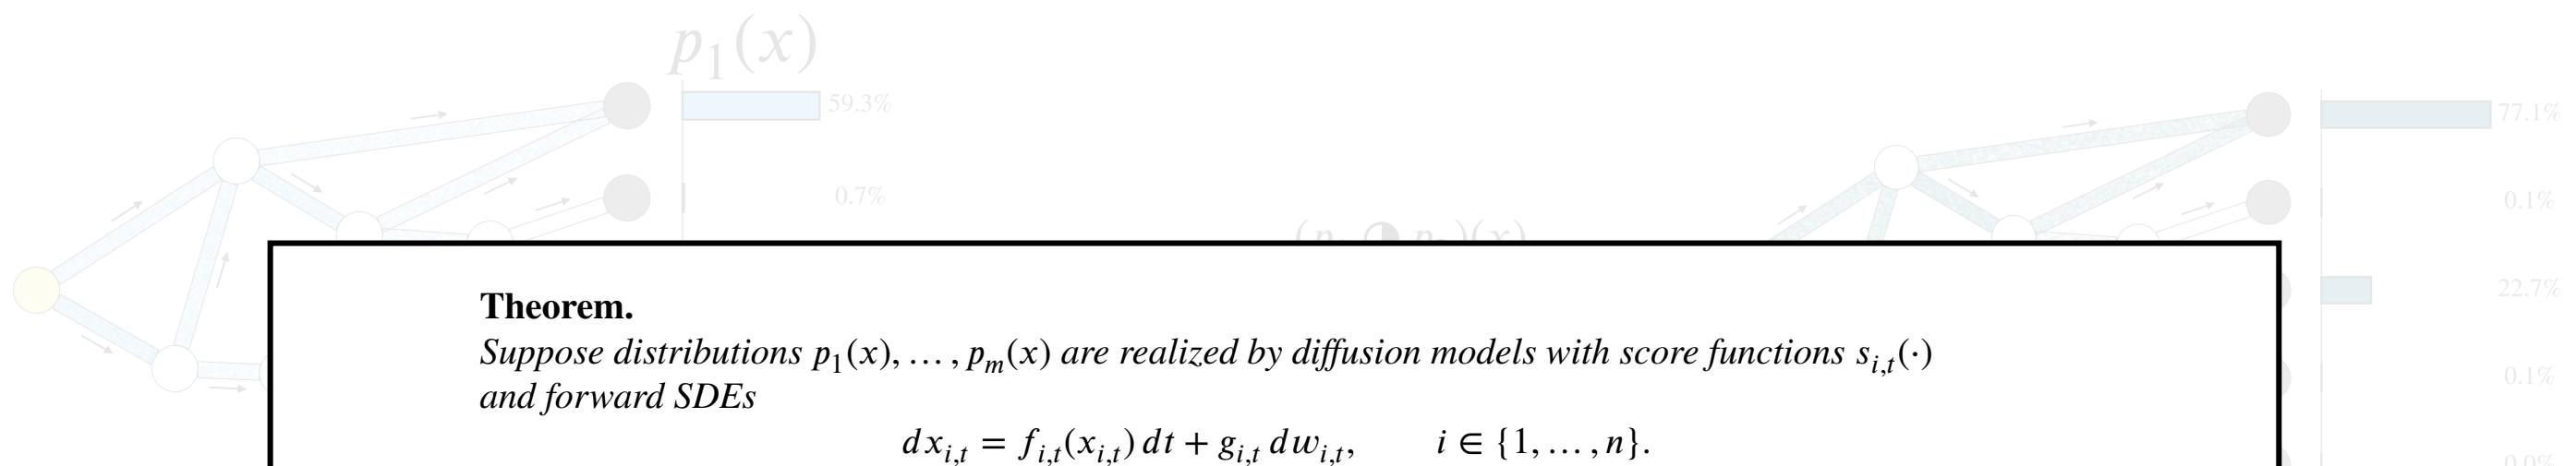
$$\tilde{p}(y_k=i) = \frac{1}{m} \quad \forall k \in \{1, \dots, n\}, \quad \tilde{p}(y_k=i|x) = \frac{p_i(x)}{\sum_{j=1}^m p_j(x)} \quad \forall k \in \{1, \dots, n\}.$$

Then, the conditional  $\tilde{p}(x|y_1, \dots, y_n)$  is realized by the forward policy

$$p_F(s'|s, y_1, \dots, y_n) = \frac{\tilde{p}(y_1, \dots, y_n|s')}{\tilde{p}(y_1, \dots, y_n|s)} \sum_{i=1}^m p_{i,F}(s'|s) \tilde{p}(y=i|s)$$

New result for GFlowNets!





**Theorem.**

Suppose distributions  $p_1(x), \dots, p_m(x)$  are realized by diffusion models with score functions  $s_{i,t}(\cdot)$  and forward SDEs

$$dx_{i,t} = f_{i,t}(x_{i,t}) dt + g_{i,t} dw_{i,t}, \quad i \in \{1, \dots, n\}.$$

Let  $y_1, \dots, y_n$  be random variables defined by

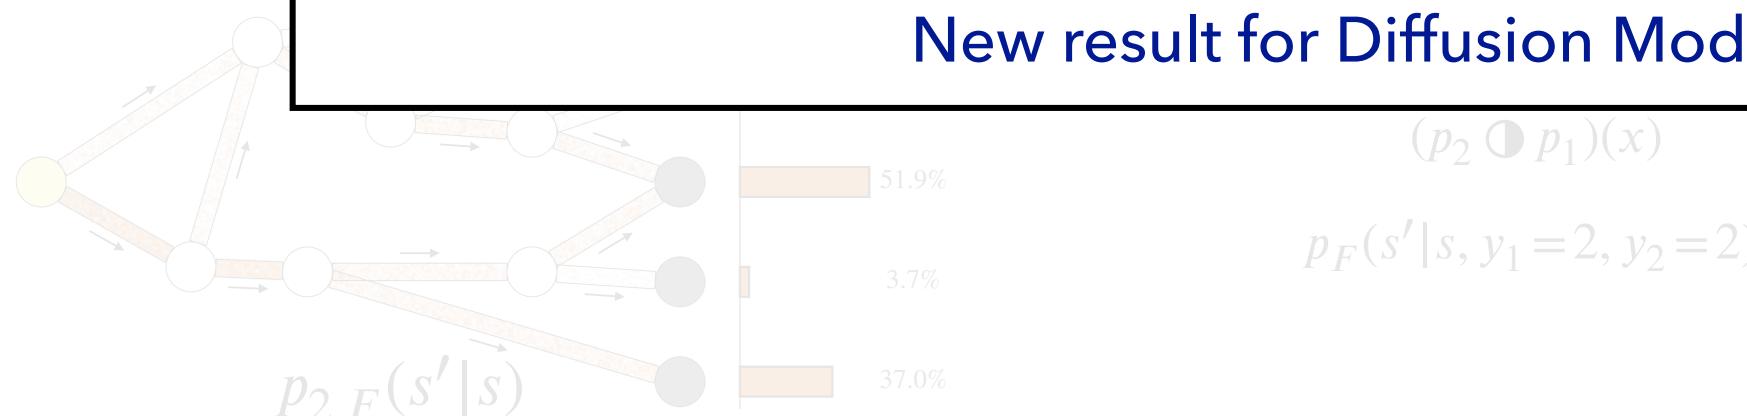
$$\tilde{p}(x, y_1, \dots, y_n) = \tilde{p}(x) \prod_{k=1}^n \tilde{p}(y_k|x), \quad \tilde{p}(x) = \frac{1}{m} \sum_{i=1}^m p_i(x),$$

$$\tilde{p}(y_k=i) = \frac{1}{m} \quad \forall k \in \{1, \dots, n\}, \quad \tilde{p}(y_k=i|x) = \frac{p_i(x)}{\sum_{j=1}^m p_j(x)} \quad \forall k \in \{1, \dots, n\}.$$

Then, the conditional  $\tilde{p}(x|y_1, \dots, y_n)$  is realized by a classifier-guided diffusion with backward SDE

$$dx_t = \left[ \sum_{i=1}^m \tilde{p}(y=i|x_t) \left( f_{i,t}(x_t) - g_{i,t}^2 \left( s_{i,t}(x_t) + \nabla_{x_t} \log \tilde{p}(y_1, \dots, y_n|x_t) \right) \right) \right] dt + \sqrt{\sum_{i=1}^m \tilde{p}(y=i|x_t) g_{i,t}^2} d\bar{w}_t.$$

New result for Diffusion Models!

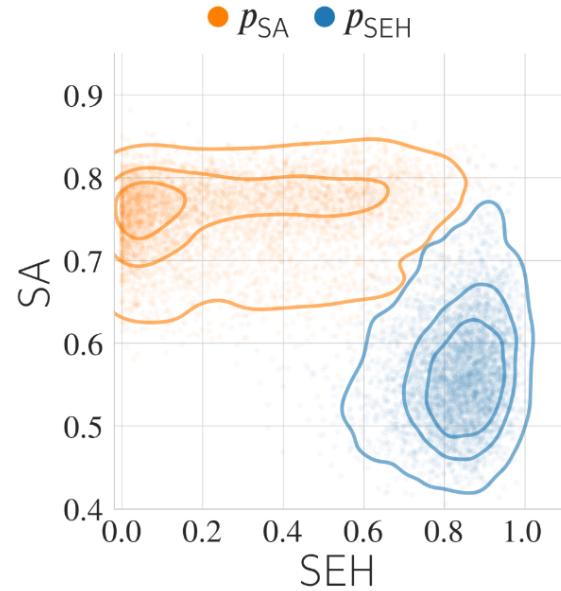


	<b>Models</b>	<b>Composition Operations</b>	<b>Sampling Algorithm</b>
<a href="#">[Hinton, Neural Computation 2002]</a> <a href="#">[Du et al, NeurIPS 2020]</a>	Energy-based models (EBMs) $p_i(x) \propto \exp(-E_i(x; \theta))$	Principle: energy-function arithmetic Product: $\frac{1}{Z} p_1(x) p_2(x)$ Negation: $\frac{1}{Z} \frac{p_1(x)}{(p_2(x))^\gamma}$	MCMC Langevin dynamics
<a href="#">[Liu et al, ECCV 2022]</a> <a href="#">[Du et al, ICML 2023]</a>	Diffusion models $p_i(x): s_{i,t}(x_t; \theta) \approx \nabla_{x_t} \log p_{i,t}(x_t)$	Principle: score-function arithmetic Product: $\frac{1}{Z} p_1(x) p_2(x)$ Negation: $\frac{1}{Z} \frac{p_1(x)}{(p_2(x))^\gamma}$	Diffusion sampling + annealed MCMC

**Challenge:** iterative generative processes (Diffusion models & GFlowNets) impose delicate balance conditions

	<b>Models</b>	<b>Composition Operations</b>	<b>Sampling Algorithm</b>
[Hinton, Neural Computation 2002] [Du et al, NeurIPS 2020]	Energy-based models (EBMs) $p_i(x) \propto \exp(-E_i(x; \theta))$	Principle: energy-function arithmetic Product: $\frac{1}{Z} p_1(x) p_2(x)$ Negation: $\frac{1}{Z} \frac{p_1(x)}{(p_2(x))^\gamma}$	MCMC Langevin dynamics
[Liu et al, ECCV 2022] [Du et al, ICML 2023]	Diffusion models $p_i(x) : s_{i,t}(x_t; \theta) \approx \nabla_{x_t} \log p_{i,t}(x_t)$	Principle: score-function arithmetic Product: $\frac{1}{Z} p_1(x) p_2(x)$ Negation: $\frac{1}{Z} \frac{p_1(x)}{(p_2(x))^\gamma}$	Diffusion sampling + annealed MCMC
<b>Challenge:</b> iterative generative processes (Diffusion models & GFlowNets) impose delicate balance conditions			
<b>Compositional Sculpting (ours)</b>	Diffusion models $p_i(x) : s_{i,t}(x_t; \theta) \approx \nabla_{x_t} \log p_{i,t}(x_t)$ GFlowNets $p_i(x) : p_{i,F}(s_{t+1} s_t; \theta)$	Principle: mixture & conditional generative processes <b>Harmonic Mean</b> : $\frac{1}{Z} \frac{p_1(x) p_2(x)}{p_1(x) + p_2(x)}$ <b>Contrast</b> : $\frac{1}{Z} \frac{(p_1(x))^2}{p_1(x) + p_2(x)}$ (+) other operations	Diffusion mixture + classifier guidance  GFlowNet mixture + classifier guidance

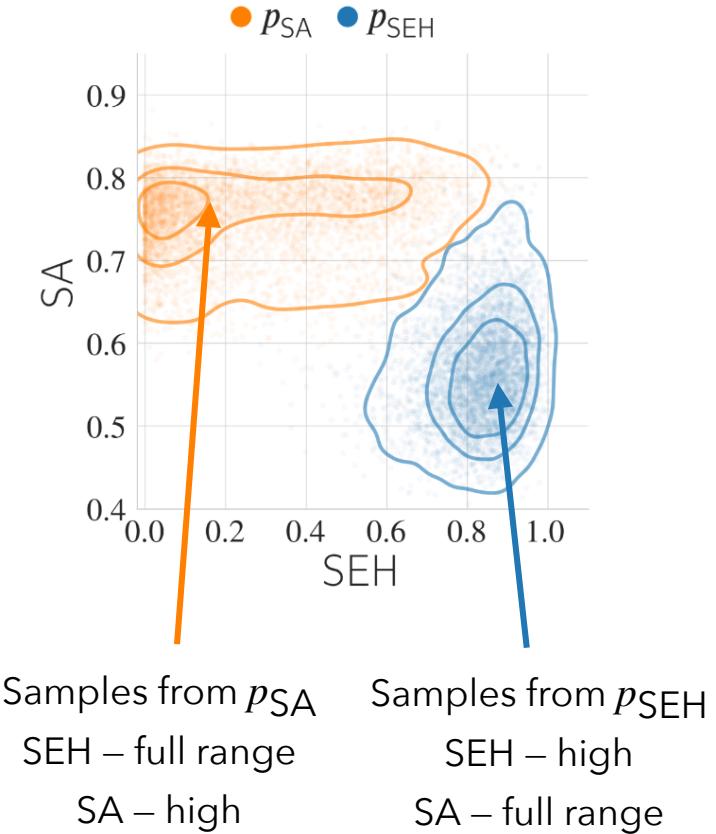
# Results: GFlowNet Composition For Molecule Generation



Molecular property scores ("Rewards")

- ▶ SEH – learned proxy of a protein binding score (soluble epoxide hydrolase)
- ▶ SA – synthetic availability score

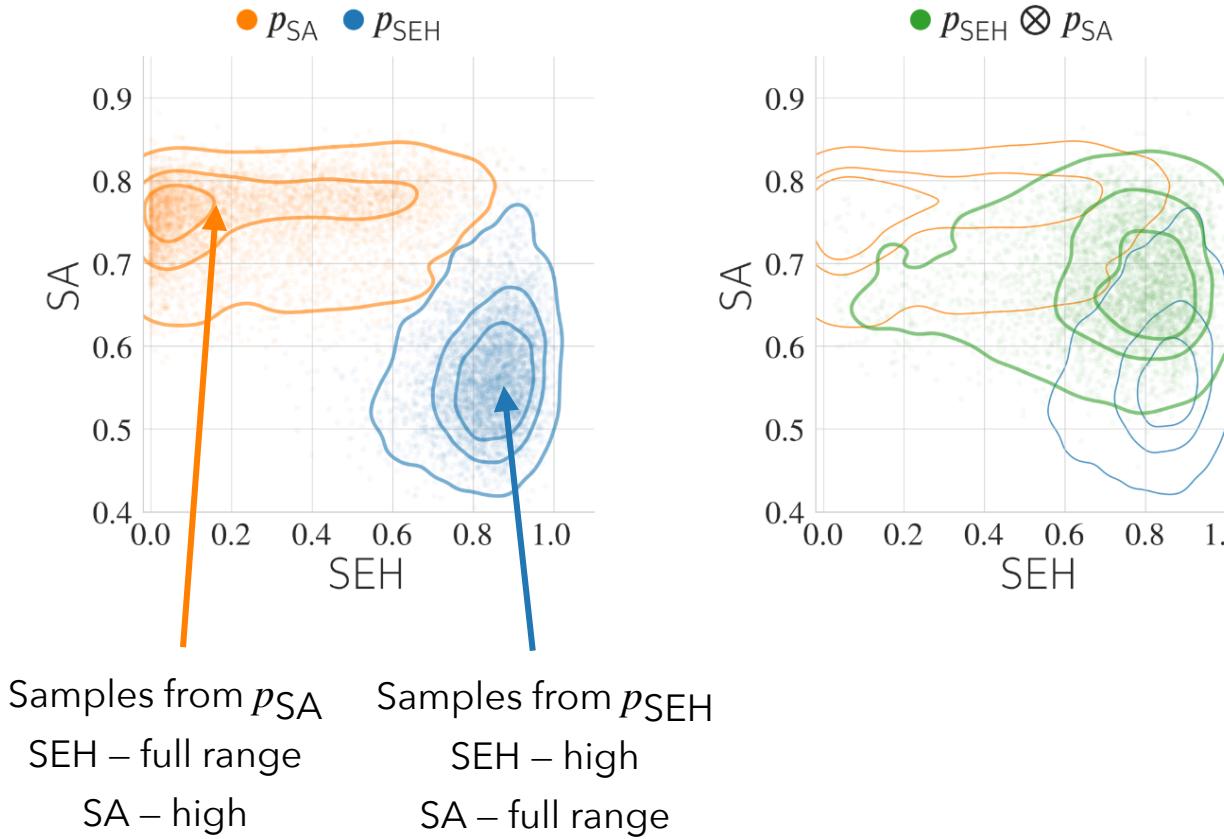
# Results: GFlowNet Composition For Molecule Generation



Molecular property scores ("Rewards")

- ▶ SEH – learned proxy of a protein binding score (soluble epoxide hydrolase)
- ▶ SA – synthetic availability score

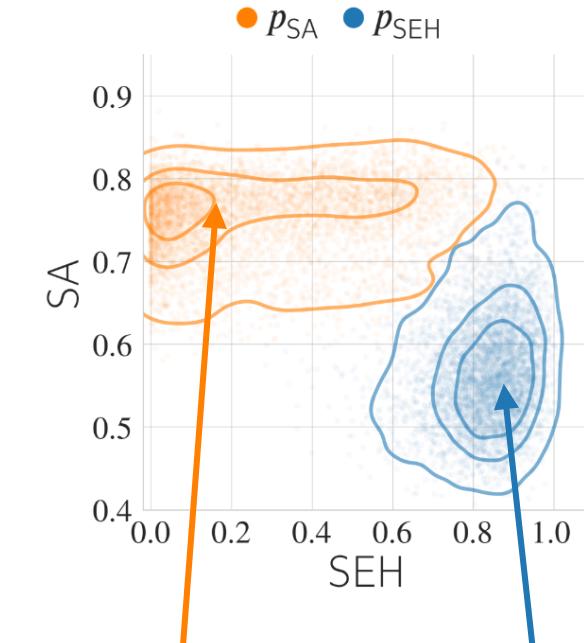
# Results: GFlowNet Composition For Molecule Generation



Molecular property scores ("Rewards")

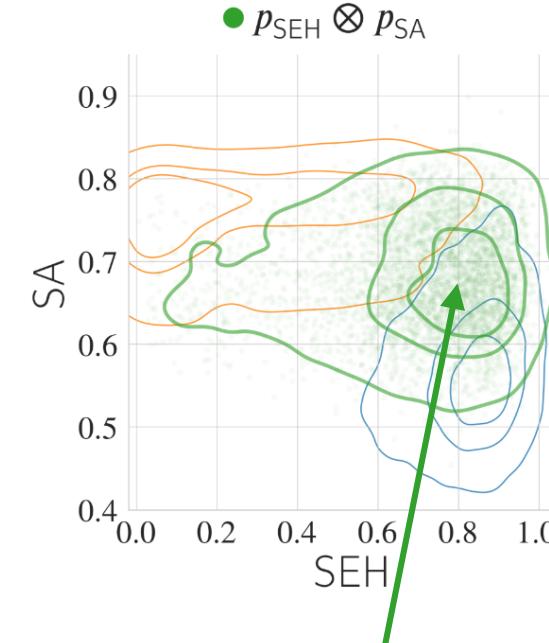
- ▶ SEH – learned proxy of a protein binding score (soluble epoxide hydrolase)
- ▶ SA – synthetic availability score

# Results: GFlowNet Composition For Molecule Generation



Samples from  $p_{SA}$   
SEH – full range  
SA – high

Samples from  $p_{SEH}$   
SEH – high  
SA – full range

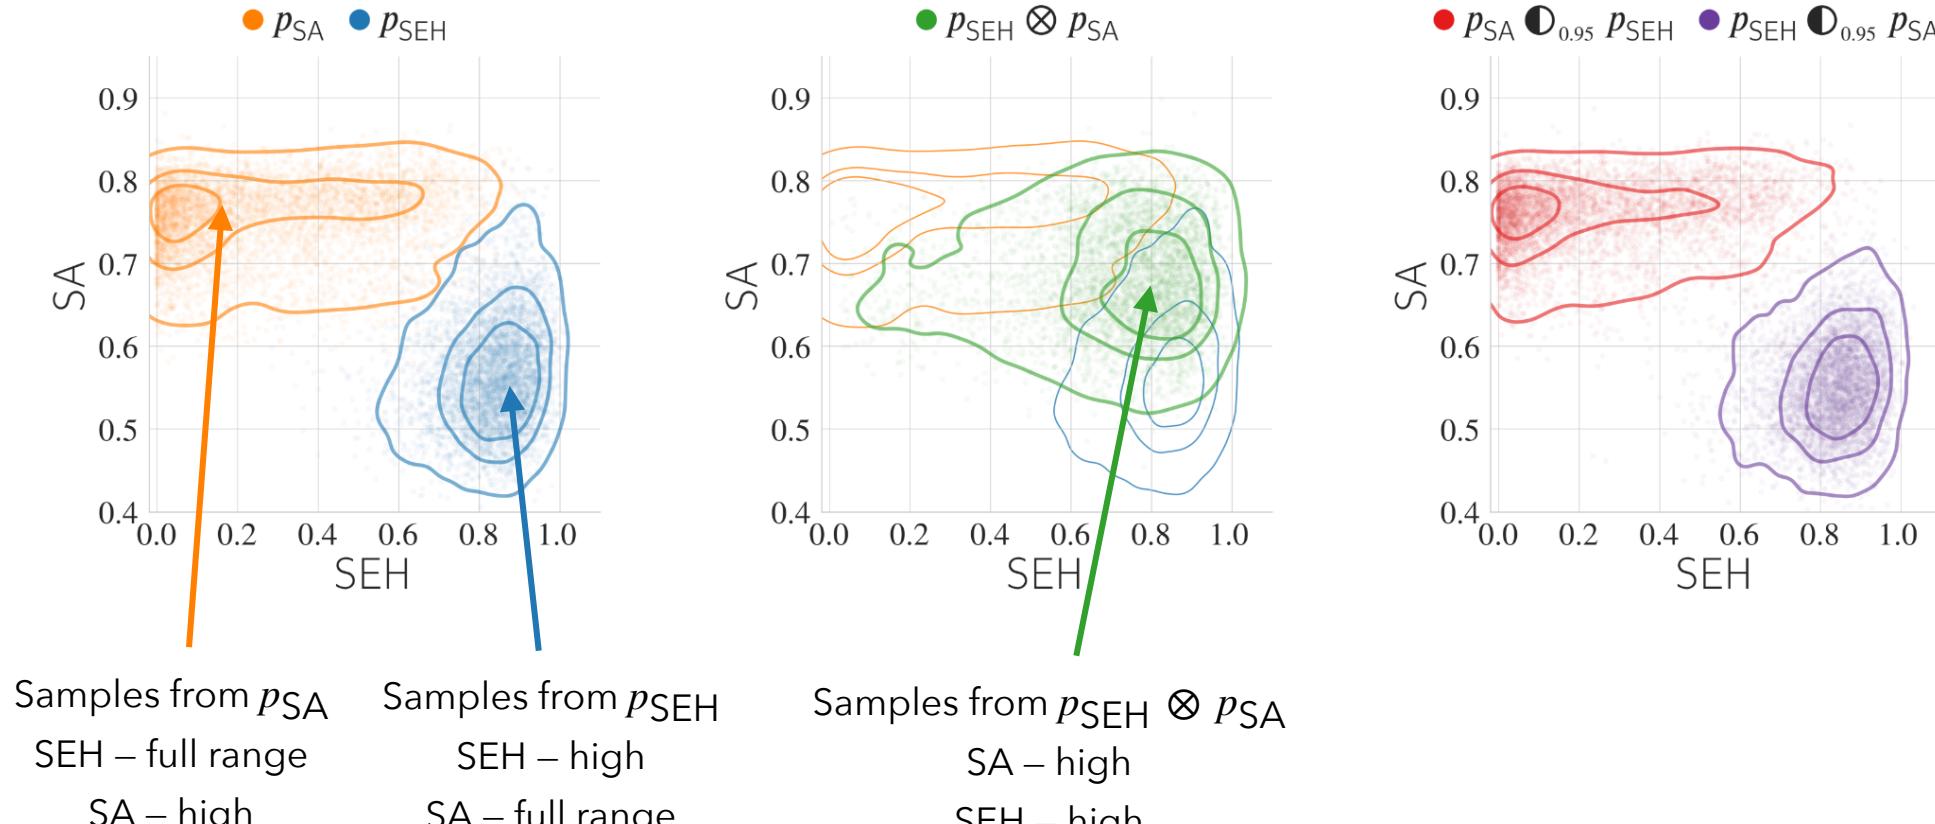


Samples from  $p_{SEH} \otimes p_{SA}$   
SA – high  
SEH – high

## Molecular property scores ("Rewards")

- ▶ SEH – learned proxy of a protein binding score (soluble epoxide hydrolase)
- ▶ SA – synthetic availability score

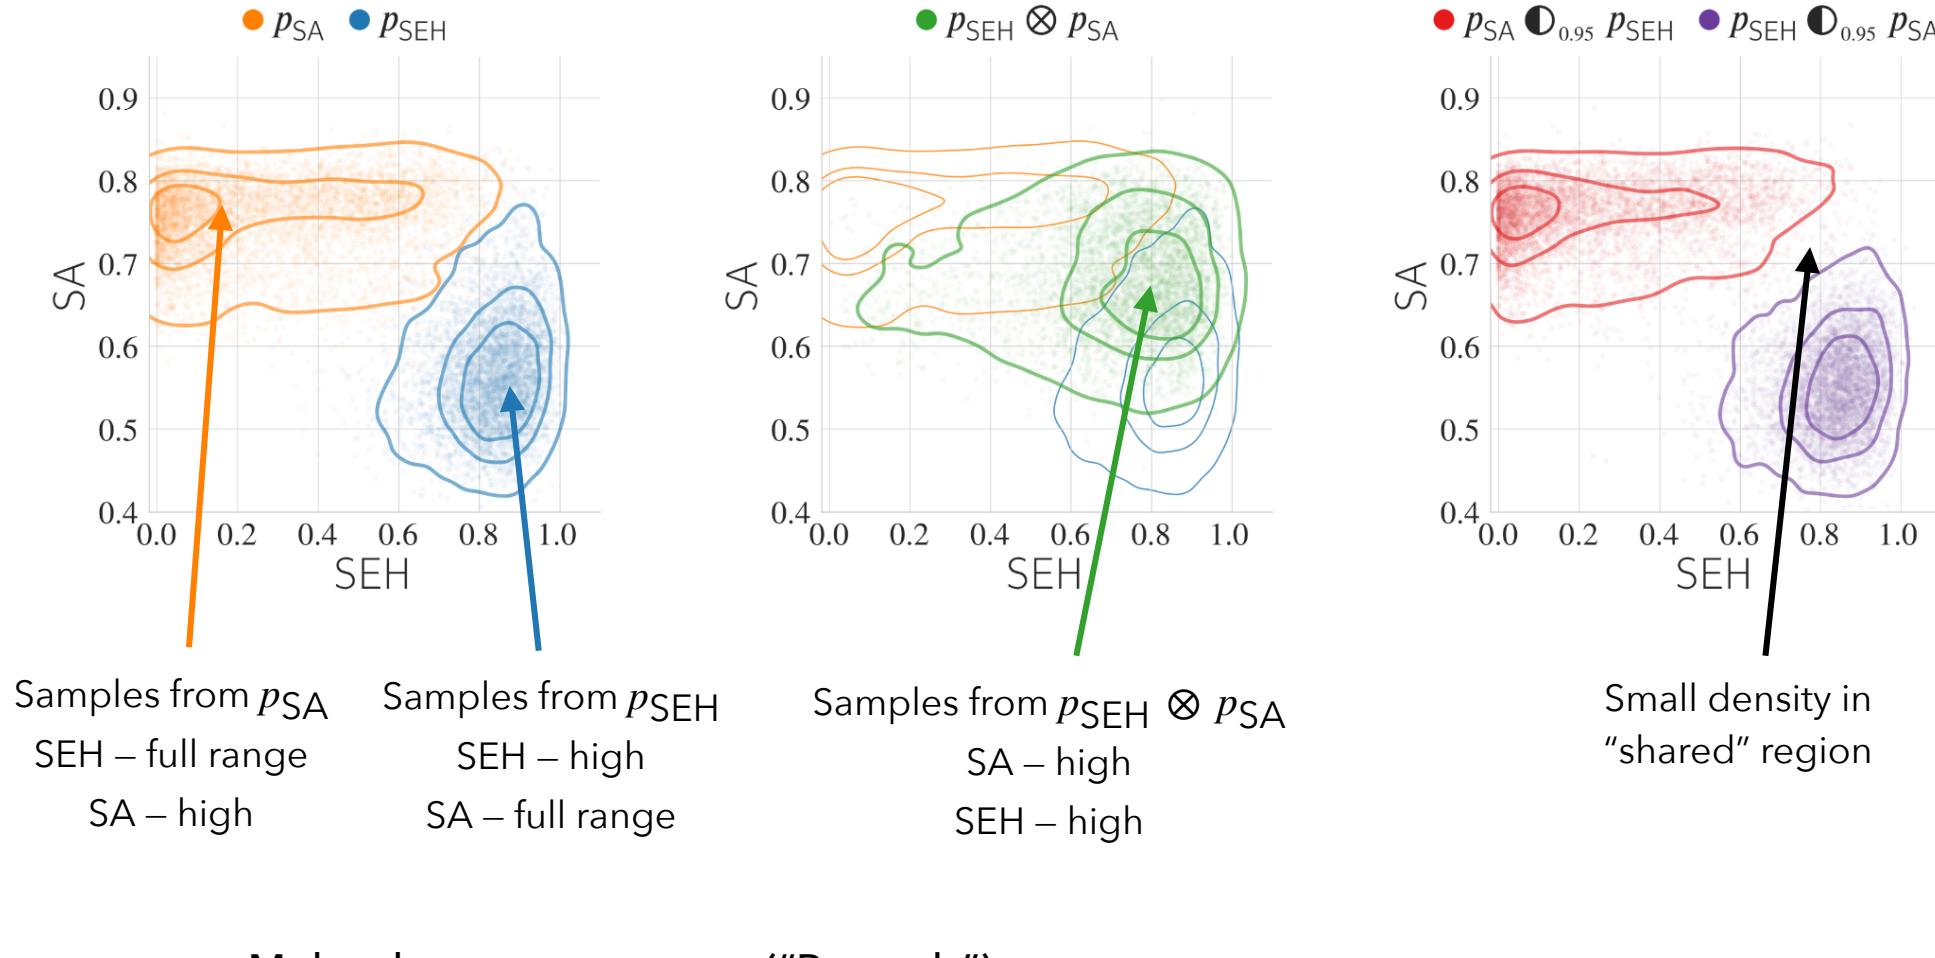
# Results: GFlowNet Composition For Molecule Generation



Molecular property scores ("Rewards")

- ▶ SEH – learned proxy of a protein binding score (soluble epoxide hydrolase)
- ▶ SA – synthetic availability score

# Results: GFlowNet Composition For Molecule Generation



- ▶ SEH – learned proxy of a protein binding score (soluble epoxide hydrolase)
- ▶ SA – synthetic availability score

# Results: Toy Diffusion Composition

1 0 8 1 0 3 1 1  
3 2 2 3 2 2 2 0  
3 3 3 3 3 0 2 1  
1 2 3 0 1 1 0 0  
1 3 1 1 3 0 0 0  
1 0 3 2 1 3 1 2  
0 1 3 2 2 3 8 0  
3 1 3 0 3 2 0 1

1 1 1 0 0 1 0 0  
0 1 1 0 1 1 0 1  
0 0 1 1 0 1 1 1  
1 1 1 0 0 1 1 0  
0 1 0 0 0 0 0 1  
1 1 0 1 0 1 1 0  
1 1 0 1 1 0 1 1  
1 1 0 1 0 1 1 0

2 2 2 2 2 0 0 0  
0 2 0 2 2 2 2 0  
2 2 0 0 2 2 2 0  
0 2 0 2 2 0 0 0  
2 2 2 2 2 0 2 2  
2 2 0 2 2 0 0 0  
2 2 0 2 2 0 0 0  
2 0 2 0 2 2 0 0

$p_1$

$p_2$

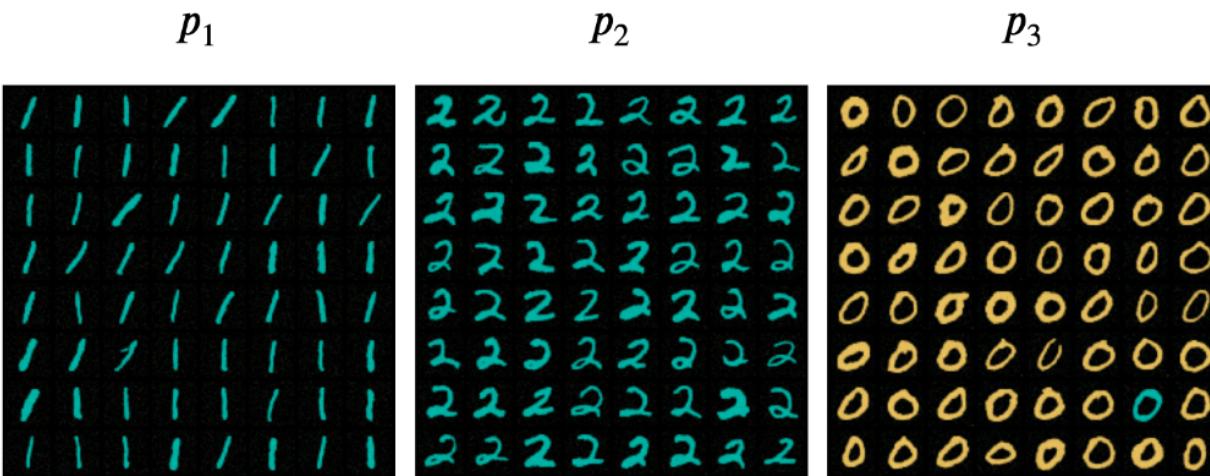
$p_3$

- 1) cyan digits
- 2) digits {"0", "1"} (cyan or beige)
- 3) digits {"0", "2"} (cyan or beige)

# Results: Toy Diffusion Composition



- 1) cyan digits
- 2) digits {"0", "1"} (cyan or beige)
- 3) digits {"0", "2"} (cyan or beige)



$$\tilde{p}\left(x \middle| \begin{matrix} y_1=1 \\ y_2=2 \end{matrix}\right)$$

$$\tilde{p}\left(x \middle| \begin{matrix} y_1=1 \\ y_2=3 \end{matrix}\right)$$

$$\tilde{p}\left(x \middle| \begin{matrix} y_1=2 \\ y_2=3 \end{matrix}\right)$$

$$\tilde{p}(x \mid y_1=i_1, \dots, y_n=i_n) \propto \frac{\prod_{k=1}^n p_{i_k}(x)}{\left(\sum_{j=1}^m p_j(x)\right)^{n-1}}$$

# Results: Toy Diffusion Composition

1 0 8 1 0 3 1 1  
3 2 2 3 2 2 2 0  
3 3 3 3 3 0 2 1  
1 2 3 0 1 1 0 0  
1 3 1 1 3 0 0 0  
1 0 3 2 1 3 1 2  
0 1 3 2 2 3 8 0  
3 1 3 0 3 2 0 1

1 1 1 0 0 1 0 0  
0 1 1 0 1 1 0 1  
0 0 1 1 0 1 1 1  
1 1 1 0 0 1 1 0  
0 1 0 0 0 0 0 1  
1 1 0 1 0 1 1 0  
1 1 0 1 1 0 1 1  
1 1 0 1 0 1 1 0

2 2 2 2 2 0 0 0 0  
0 2 0 2 2 2 2 0 0  
2 2 0 0 2 2 2 0 0  
0 2 0 2 2 0 0 0 0  
2 2 2 2 2 2 0 2 2  
0 2 0 0 2 0 0 2 2  
2 2 0 2 2 0 0 0 0  
2 0 2 0 2 2 2 0 0

3 3 3 3 3 3 3 3 3  
3 3 3 3 3 3 3 3 3  
3 3 3 3 3 3 3 3 3  
3 3 3 3 3 3 3 3 3  
3 3 3 3 3 3 3 3 3  
3 3 3 3 3 3 3 3 3  
3 3 3 3 3 3 3 3 3  
3 3 3 3 3 3 3 3 3

1 1 1 1 1 1 1 1 1  
1 1 1 1 1 1 1 1 1  
1 1 1 1 1 1 1 1 1  
1 1 1 1 1 1 1 1 1  
1 1 1 1 1 1 1 1 1  
1 1 1 1 1 1 1 1 1  
1 1 1 1 1 1 1 1 1  
1 1 1 1 1 1 1 1 1

2 2 2 2 2 2 2 2 2  
2 2 2 2 2 2 2 2 2  
2 2 2 2 2 2 2 2 2  
2 2 2 2 2 2 2 2 2  
2 2 2 2 2 2 2 2 2  
2 2 2 2 2 2 2 2 2  
2 2 2 2 2 2 2 2 2  
2 2 2 2 2 2 2 2 2

$p_1$

$p_2$

$p_3$

$\tilde{p}\left(x \middle| y_1=1\right)$

$\tilde{p}\left(x \middle| y_1=2\right)$

$\tilde{p}\left(x \middle| y_1=3\right)$

1 1 1 1 1 1 1 1  
1 1 1 1 1 1 1 1  
1 1 1 1 1 1 1 1  
1 1 1 1 1 1 1 1  
1 1 1 1 1 1 1 1  
1 1 1 1 1 1 1 1  
1 1 1 1 1 1 1 1  
1 1 1 1 1 1 1 1

2 2 2 2 2 2 2 2 2  
2 2 2 2 2 2 2 2 2  
2 2 2 2 2 2 2 2 2  
2 2 2 2 2 2 2 2 2  
2 2 2 2 2 2 2 2 2  
2 2 2 2 2 2 2 2 2  
2 2 2 2 2 2 2 2 2  
2 2 2 2 2 2 2 2 2

0 0 0 0 0 0 0 0 0  
0 0 0 0 0 0 0 0 0  
0 0 0 0 0 0 0 0 0  
0 0 0 0 0 0 0 0 0  
0 0 0 0 0 0 0 0 0  
0 0 0 0 0 0 0 0 0  
0 0 0 0 0 0 0 0 0  
0 0 0 0 0 0 0 0 0

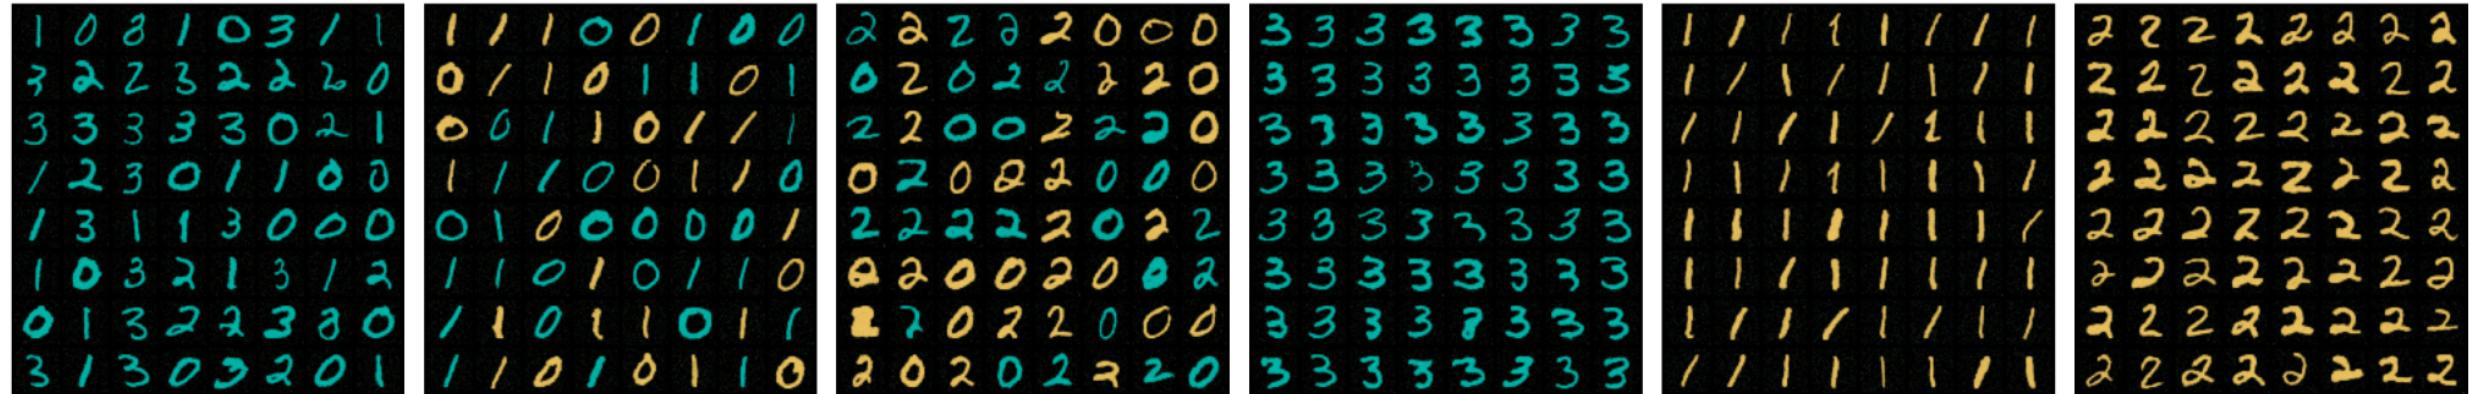
$$\tilde{p}(x \mid y_1=i_1, \dots, y_n=i_n) \propto \frac{\prod_{k=1}^n p_{i_k}(x)}{\left(\sum_{j=1}^m p_j(x)\right)^{n-1}}$$

$\tilde{p}\left(x \middle| y_1=1\right)$

$\tilde{p}\left(x \middle| y_1=2\right)$

$\tilde{p}\left(x \middle| y_1=3\right)$

# Results: Toy Diffusion Composition



$p_1$

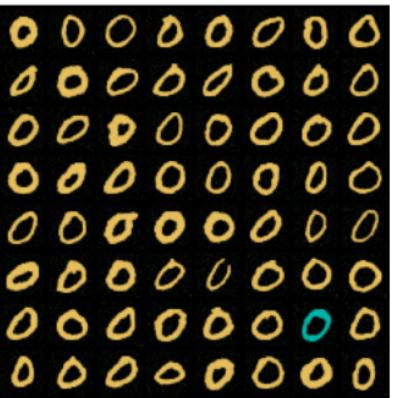
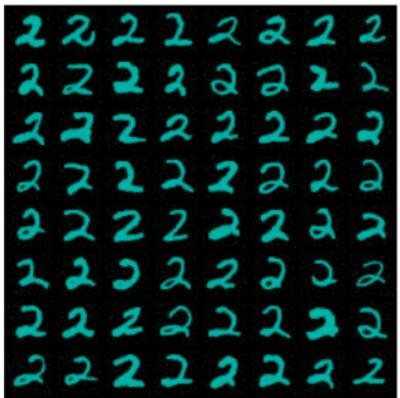
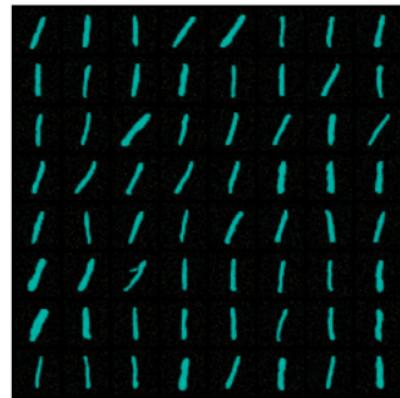
$p_2$

$p_3$

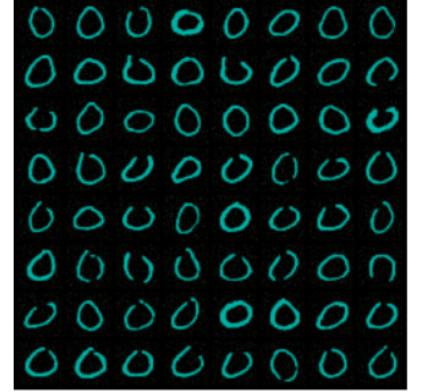
$$\tilde{p}(x \mid y_1=1)$$

$$\tilde{p}(x \mid y_1=2)$$

$$\tilde{p}(x \mid y_1=3)$$



$$\tilde{p}(x \mid y_1=i_1, \dots, y_n=i_n) \propto \frac{\prod_{k=1}^n p_{i_k}(x)}{\left(\sum_{j=1}^m p_j(x)\right)^{n-1}}$$



$$\tilde{p}(x \mid y_1=1)$$

$$\tilde{p}(x \mid y_1=2)$$

$$\tilde{p}(x \mid y_1=3)$$

$$\tilde{p}(x \mid y_2=1)$$

# Compositional Sculpting: Summary

New method for composition of diffusion models or GFlowNets

- ▶ **Controllable composition operations** through inference:

Base models → (prior, observations) → posterior

- ▶ Binary: “**Harmonic Mean**” and “**Contrast**”
- ▶ Generalized: N-ary, parameterized, chained

- ▶ **Tractable sampling algorithm**: classifier guidance on mixture of base models

Experimental validation

- ▶ GFlowNets, controllable synthetic domain (2D grid)
- ▶ GFlowNets, molecule generation
- ▶ Diffusion models, small image generation

# Thesis Overview

## Contributions

- ▶ Novel principled algorithms for training and inference in deep probabilistic models
- ▶ Guarantees
  - ▶ Training: optimality of the desired target configurations
  - ▶ Inference: sampling from target distributions

## Approach

- ▶ Guide complex models using signal derived from a simple auxiliary model (discriminator / coordinator)

## Applications

- ▶ Image generation
- ▶ Unsupervised domain adaptation under multi-faceted distribution shift
- ▶ Multi-objective drug-like molecule generation

## Models and analysis tools

- ▶ Generative adversarial networks, domain adversarial neural networks, diffusion models, GFlowNets
- ▶ Game theoretic training algorithms, optimal transport, dynamical systems, stochastic processes